HARVARD BUSINESS REVIEW

Why Employees *Override* AI

The problem isn't the technology. It's your organization.

PAGE 42

November–
December 2026

BEGIN YOUR
OWN TRADITION

PATEK PHILIPPE
GENEVE

THE CUBITUS COLLECTION

YOU NEVER ACTUALLY OWN A PATEK PHILIPPE.

YOU MERELY TAKE CARE OF IT FOR THE NEXT GENERATION.

PATEK.COM

The AI era is a human one.

We believe data and AI should be used to improve lives, strengthen communities and help solve meaningful problems. As AI becomes more capable, the need for responsibility and deep industry understanding only grows.

Domain expertise matters more than ever in the age of generative and agentic AI, grounding innovation in real-world context and judgment. Our commitment is to keep people in charge of the outcomes, and to build technology that supports transparency and trust.

This promise reflects the kind of company we set out to be in 1976 and the company we choose to be 50 years later. As we look to the next 50 years, we remain committed to keeping people at the center, innovating responsibly and serving as a trusted partner – combining advanced AI with proven industry expertise to help power the world's decisions.

Jim Goodnight, CEO, SAS

sas DATA & AI | The Power to Know® for 50 Years

© 2026 SAS Institute Inc. All rights reserved.

AI-modified

GAGGENAU

FOR THOSE WHO KNOW

Creators of luxury kitchen appliances.
Established 1683.

Experience it

The difference is Gaggenau

We help Carnival turn 5.6 billion transactions a year into endless fun.

Unisys powers the tech platform behind every booking and online transaction before anyone ever steps on board, helping Carnival connect millions of guests to moments of fun. We deliver the reliability and security Carnival has counted on — from their first ship to a global fleet — for almost 50 years and counting.
Discover how at unisys.com/carnival

We help Carnival
Keep Breaking Through

unisys | Carnival

Contents

November–December 2026

41 Features

42 AI & MACHINE LEARNING

Why Even Your Best Employees Override AI Systems That Work

It's not that they don't trust the technology. It's that they don't trust your organization.
Das Narayandas

50 OPERATIONS & SUPPLY CHAIN MANAGEMENT

Design Your Supply Chain as a Portfolio of Options

You need to be able to defer investments, quit unfavorable positions, expand in promising regions, and shift operations as policies evolve.
Luiz Mesquita et al.

58 LEADERSHIP

How CEOs Set the Tempo

One of the most underappreciated aspects of the role is deciding how fast to move and when.
Nitin Nohria

COVER PHOTOGRAPH
Jamie Chung

Photograph by **JAMIE CHUNG**

November–December 2026

"As AI makes relationships designable, brand architecture becomes, more than ever, relationship architecture."

"YOUR CHATBOT WILL DEFINE YOUR BRAND," PAGE 80

68 CUSTOMER STRATEGY
Do You Know What Your Customers Value?
Not according to the data. Here's how to fix that.
Vikas Mittal et al.

80 AI & MACHINE LEARNING
Your Chatbot Will Define Your Brand
As conversational interfaces begin to dominate the customer experience, relationship design becomes a strategic priority.
Julian De Freitas

90 ENTREPRENEURSHIP
The Risks of Scaling Too Fast
Startups that chase growth and ignore unit economics are more likely to fail.
Benjamin Hallen, Ron Tidhar, and Kathleen Eisenhardt

102 BOARDS
Building a Board That Fits
A good one is designed for purpose, people, process, and leadership in a specific context.
Timothy J. Rowley and Stanislav Shekshnia

114 STRATEGY
The New Platform Playbook
Success is no longer about winner-take-all. Here's a guide to making smarter strategic choices.
Chiara Farronato

Photograph by **HELGA STENTZEL**

HENRY CAVILL

Elegance is an attitude

LONGINES

THE LONGINES
MASTER COLLECTION

November–
December
2026

Our Commitment to Sustainability
We're proud that the paper we use in our print magazine is certified under the Sustainable Forestry Initiative® program, meaning that it comes from responsibly managed sources and is a renewable resource.

SUSTAINABLE FORESTRY INITIATIVE
Certified Sourcing
www.sfiprogram.org
SFI-00993

Vlad Tenev, CEO of Robinhood

17
Idea Watch
New Research and Emerging Insights

17 IN THEORY
For Real Performance Gains with AI, Start by Mapping Your Workflows
That's how to reveal the most promising opportunities. PLUS Why it pays to be a cooperative executive, the penalty faced by referred hires, why people resist reskilling, and more.

27
In Focus
Collaborating Amid Complexity

28
Four Hidden Traps of Team Dynamics
Susan MacKenty Brady, Stuart D. Kliman, and Leslie C. Smith

31
Prevent Team Friction from Turning into Dysfunction
Leonard A. Schlesinger, Joseph Fuller, and Robert V. Toomey

34
The Office Amenity That Actually Improves Teamwork
Ron Friedman

36
How Highly Effective Global Teams Collaborate Across Cultures
Zak Dychtwald

126
Experience
Advice and Inspiration

126 MANAGING YOURSELF
How Leaders Counter Pessimism in Trying Times
They model constructive behavior, reframe pressure, and highlight progress.
Michelle Gielan

130 HOW WE DID IT
Robinhood's CEO on Refounding the Business Post-IPO
The hardest leadership decisions begin after yesterday's strategy stops working.
Vlad Tenev

134 CASE STUDY
Are We Using the Right Customer Service Metrics?
The manager of a call center wonders whether her company is employing the right incentives.
Kristin Behfar, Martin N. Davidson, and Gerry Yemen

144 LIFE'S WORK
Michael Douglas

Departments
12 FROM THE EDITOR
13 CONTRIBUTORS
140 EXECUTIVE SUMMARIES

Photograph by **CHRISTIE HEMM KLOK**

Whittier Trust generally requires a minimum of $15 million in marketable securities for new client relationships. Investment and Wealth Management Services are provided by Whittier Trust Company and The Whittier Trust Company of Nevada, Inc. (referred to herein individually and collectively as "Whittier Trust"), state-chartered trust companies wholly owned by Whittier Holdings, Inc. ("WHI"), a closely held holding company. This document is provided for informational purposes only and is not intended, and should not be construed, as investment, tax or legal advice. Past performance is no guarantee of future results and no investment or financial planning strategy can guarantee profit or protection against losses. All names, characters, and incidents, except for certain incidental references, are fictitious. Any resemblance to real persons, living or dead, is entirely coincidental.

Your Organization's Real Trust Issue with AI

Amy Bernstein

CHANCES ARE GOOD that your organization has already installed at least one AI system that can make smart decisions. Maybe it does demand forecasting, maybe application processing. Whatever the case, chances are just as good that, even if it's a reliable system, humans are rejecting its recommendations.

This may look like a technology problem or a change management problem. Harvard Business School's Das Narayandas argues that it is usually something more basic: a problem of authority and accountability. In this issue's cover story, "Why Even Your Best Employees Override AI Systems That Work," (page 42) he describes organizations in which the AI performs well but humans continually overrule it.

Why? Because the person asked to trust the system is often the person held responsible when its decision goes awry. Under those conditions, choosing to override AI is not necessarily resistance to change. It can be an act of rational self-protection.

That distinction matters. If companies want AI to make binding decisions, Narayandas contends, they must reassign both the authority and the accountability for those decisions. That means redesigning roles, decision rights, incentives, and the rules governing when people should intervene.

AI adoption, then, depends not only on what the technology can do and whether employees trust it. It also depends on whether a business is willing to reorganize around what AI can do—and whether employees trust the business to follow through.

Amy Bernstein
Editor in chief

Connect with HBR

JOIN US ONLINE
WWW.HBR.ORG
FACEBOOK @HBR
INSTAGRAM
@harvard_business_review

LINKEDIN
Harvard Business Review
TIKTOK
@harvardbusinessreview

WHATSAPP
Harvard Business Review
X @HarvardBiz
YOUTUBE
@harvardbusinessreview

CONTACT HBR
PHONE
800.988.0886

EMAIL
editors@hbr.org
customerservice@hbr.org

Elie Honein

Contributors

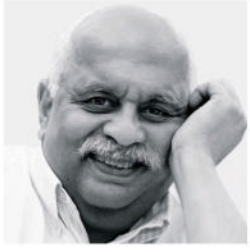

First as a seller of automation systems and then as a researcher, **Das Narayandas** observed that employees rarely use enterprise systems as their companies intend—and the gap has widened as AI has become more common in the workplace. Even when autonomous systems are designed to operate well independently, employees are often reluctant to let them do so. In his article, Narayandas, now a professor at Harvard Business School, examines that disconnect.

42 Why Even Your Best Employees Override AI Systems That Work

As a recreational tennis player, **Kathleen Eisenhardt** learned that speed alone isn't right for every shot. So when Silicon Valley embraced fast growth at almost any cost, the Stanford professor sensed an opportunity. Much of the rhetoric, she thought, focused on software businesses while ignoring the 90% of companies that ship products, hold inventory, and face major variable costs. In this issue she and her coauthors explain why most companies need to master unit economics before accelerating.

90 The Risks of Scaling Too Fast

Arizona State University professor **Thomas Choi** has come to see supply chains as complex adaptive systems: networks of interdependent firms making decisions to ensure their survival and prosperity. When tariffs became a pressing concern, he wondered if companies with more supply chain options would fare better because of their greater flexibility. He and his coauthors explore that question and offer a framework for managing supply chain decisions amid uncertainty.

50 Design Your Supply Chain as a Portfolio of Options

Chiara Farronato's interest in digital platforms began when she arrived at Stanford to pursue her PhD in 2009, just as companies such as Airbnb, Uber, and Lyft were taking off in Silicon Valley. Watching their rapid growth—and collaborating with some of them as a researcher—sparked her fascination with the way platforms evolve. In her article the associate professor at Harvard Business School describes how tighter regulation, greater capital discipline, and shifting customer preferences are reshaping the platform economy.

114 The New Platform Playbook

San Francisco–based photographer **Christie Hemm Klok** is drawn to people who work in science and technology because of their potential to have an outsize impact on the world. The Bay Area, she says, has provided rich subject matter, from AI researchers and robotics engineers to climate scientists and conservation biologists. In her portraits of Robinhood founder Vladimir Tenev, she leaned into dramatic lighting and contrast to reflect the scale, stakes, and reach of his work.

130 Robinhood's CEO on Refounding the Business Post-IPO

Meet the Interactive Issue

An intuitive way to experience HBR.

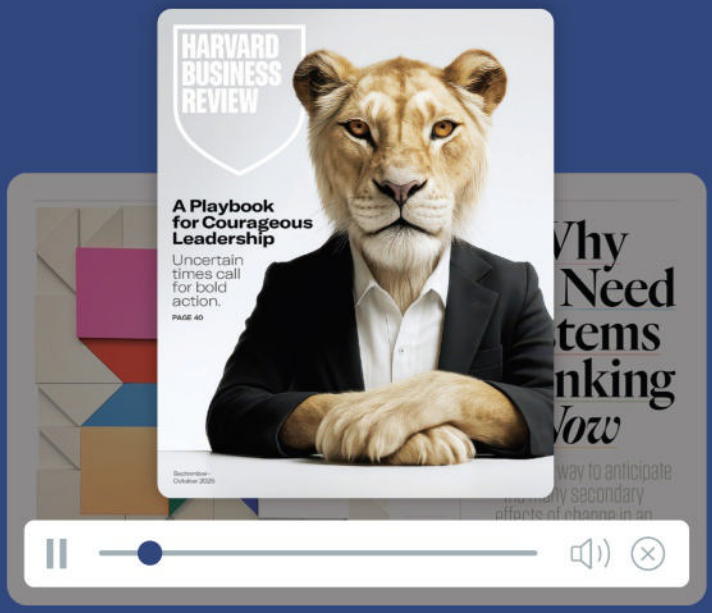

Swipe. Click. Listen. Search.

hbr.org/subscribe

HARVARD BUSINESS REVIEW

EDITOR IN CHIEF Amy Bernstein

VP, EDITORIAL ENGAGEMENT AND DIGITAL MEDIA
Maureen Hoch

EDITOR, HBR
Daniel McGinn

EDITORIAL DIRECTOR, HBR PRESS
Dave Lievens

EXECUTIVE CREATIVE DIRECTOR
John Korpics

EXECUTIVE EDITORS
Alison Beard
Ania W. Masinter
Sarah Moughty

EDITOR AT LARGE
Adi Ignatius

EDITORIAL

MANAGING EDITOR, HBR.ORG Gretchen Gavett

SENIOR EDITORS Scott Berinato, Courtney Cashman, David Champion, Paige Cohen, Susan Donovan, Kevin Evers, Alana Grambush, Eben Harrell, Jeff Kehoe, Scott LaPierre, Toby Lester, Juan Martinez, Steven Prokesch, Thomas Stackpole, Erica Truxler

SENIOR PRODUCER, AUDIO & LIVE EXPERIENCES Mary Dooe

PRODUCTION DIRECTOR & MANAGING EDITOR, HBR PRESS Allison Peter

SENIOR ASSOCIATE EDITORS Stefanie Fernández, Kelsey Hansen, Ramsey Khabbaz

EDITORIAL COORDINATORS Rachel DuRose, Emily Lang

EXECUTIVE ASSISTANT Christine C. Jack

CONTRIBUTING EDITORS Eileen Ahuja, Karen Dillon, Sarabeth Fields, Amy Gallo, Kelly Messier, JM Olejarz, Dana Rousmaniere

DESIGN

DESIGN DIRECTORS Stephani Finks *HBR Press,* Susannah Haesche *HBR,* Marta Kusztra *Product & UX,* Karen Player *Special Projects*

SENIOR GRAPHICS EDITOR Macaulay Campbell

SENIOR DESIGNERS Carey Bass *HBR,* Kim Lapan *UX,* So Young Park *Product*

SENIOR PHOTO EDITOR Sasha Patkin

PRODUCTION

MANAGER, EDITORIAL PRODUCTION Christine Wilder

SENIOR PRODUCTION EDITORS Victoria Desmond, Kristin Murphy Romano, Jennifer Waring

PRODUCTION DESIGNER Emily Peterson

PRODUCTION EDITOR Anne Starr

EDITORIAL OFFICES
60 Harvard Way, Boston, MA 02163
617-783-7410 | fax 617-783-7493
HBR.org

VOLUME 104, NUMBER 6
NOVEMBER–DECEMBER 2026
Printed in the U.S.A.

Copyright 2026 Harvard Business School Publishing Corporation. All rights reserved.

SENIOR VICE PRESIDENT, GROUP PUBLISHER Erika Heilman

VICE PRESIDENT, NATIONAL ADVERTISING & MARKETING
Craig Catalano

MANAGING DIRECTOR, DIGITAL PRODUCT STRATEGY
Dave Lefort

ADVERTISING
DIRECTOR, U.S. SALES & PRODUCT DEVELOPMENT Daniel Cohen
DIRECTOR OF RESEARCH, ANALYTIC SERVICES Beth Tracton-Bishop
SENIOR EDITOR, RESEARCH & SPECIAL PROJECTS Todd Pruzan
ASSISTANT EDITOR, ANALYTIC SERVICES Amanda Yao
ASSOCIATE DIRECTOR, MARKETING Samantha Holloway
DIRECTOR OF PROJECT MANAGEMENT, ANALYTIC SERVICES Samantha Barry
SENIOR RESEARCHER, ANALYTIC SERVICES Mary Callaghan
SENIOR PROJECT MANAGER, ANALYTIC SERVICES Audrey White
CONTRIBUTING CONSULTANT MaryAlice Holmes

MARKETING & COMMUNICATIONS
DIRECTOR, CONSUMER MARKETING Nini Diana
DIRECTOR, COMMUNICATIONS Amy Poftak
ASSOCIATE DIRECTOR, MARKETING FOR CUSTOMER JOURNEYS Caty Jackson
SENIOR MANAGER, LEAD GENERATION & CONVERSION Josh Getman
MANAGER, DIGITAL MARKETING Stephanie Kazemi
SALES ASSOCIATE, B2B & CUSTOM SALES Wendy Mkhize

PRODUCT MANAGEMENT
DIRECTOR, PRODUCT MANAGEMENT Anne Bartholomew
SENIOR PRODUCT MANAGER Jon Goldberg
SENIOR MANAGER, AI & MOBILE Gerrie Diaz
SENIOR TECHNICAL PRODUCT MANAGER Frederic Lalande
SENIOR PRODUCT MANAGER Marqui Mapp-Taylor
AGILE LEADS Kara Berrini, James Monia

BUSINESS & CUSTOMER ANALYTICS
SENIOR DIRECTOR, BUSINESS ANALYTICS & INSIGHTS Jessica Avery
DIRECTOR, CUSTOMER & MARKETING ANALYTICS Emily Freeman Kroger
ASSOCIATE DIRECTOR, EMAIL AUTOMATION & OPERATIONS Jerry Knott
ASSOCIATE DIRECTORS, EXPERIMENTATION & ANALYTICS Chris Chaky, Abigail Dawson
SENIOR MANAGER, EMAIL PRODUCTION Frances Lee
MANAGERS, EMAIL MARKETING Milo Goodman, Cara Vaccaro
SENIOR ANALYST, CUSTOMER & MARKETING ANALYTICS Sushane Sharma
ANALYST, EXPERIMENTATION & ANALYTICS Tony Manev

FINANCE & OPERATIONS
SENIOR DIRECTOR, FINANCE & OPERATIONS Greg St. Pierre
ASSOCIATE DIRECTOR, OPERATIONS & CUSTOMER EXPERIENCE Greg Daly
SENIOR MANAGER, FINANCE Krishna Dudhwewala
MANAGER, CUSTOMER EXPERIENCE Danielle Weber
SPECIALIST, DATA INTEGRITY & OPERATIONS Edward D. Domina IV

HBR PRESS
ASSOCIATE PUBLISHER Jon Shipley
DIRECTOR, MARKETING & COMMUNICATIONS Julie Devoll
DIRECTOR, B2B & CUSTOM SALES Lindsey Dietrich
ASSOCIATE DIRECTOR, MEDIA & COMMUNICATIONS Sally Ashworth
SENIOR MANAGER, OPERATIONS Rick Emanuel
SENIOR MANAGER, MARKETING Alexandra Kephart
SENIOR ASSOCIATE, SALES Jordan Concannon

TECHNOLOGY
CONTENT SYSTEMS ARCHITECT Matt Wagner
SOFTWARE ARCHITECTS Bill Munyan, Kenneth Pereira
SOFTWARE ARCHITECT, MOBILE Ahmad Khan
QUALITY ASSURANCE ARCHITECT Matthew Han
FRONT-END ENGINEER Caitlin Smith

WORLDWIDE ADVERTISING OFFICES

NEW YORK 973-769-3631
Maria A. Beacom *Sales Director, Technology & North American Research*

LONDON 44-7890-608-707
Alexander Bart *Sales Director, EMEA & APAC*

CALIFORNIA 415-986-7762
BEIJING 86-155-7661-4801
FRANCE 33-6-64-41-37-17
HONG KONG 852-9188-6157
INDIA 91-98201-32670
JAPAN 81-90-4202-9140
KOREA 82-10-8819-8156
SOUTH AMERICA 55-11-98160-0177
UAE 1-647-967-0153
UNITED KINGDOM 44-7890-608-707

For all other inquiries, please call 212-872-9280 or visit HBR.org/hbr-advertising-sales.

SUBSCRIPTION SERVICES
Harvard Business Review
20 Guest St, Suite 400
Brighton, MA 02135
customercare@hbr.org
US/Canada: 800-274-3214
International: 1-617-783-7500

RATES PER YEAR
United States $119
Canada US$139
Mexico US$139
International US$165

ARTICLE REPRINTS
To purchase reprints, go to HBR.org.

LIBRARY ACCESS
Libraries offer online access to current and back issues of *Harvard Business Review* through EBSCO host databases.

SUBMISSIONS
We encourage prospective authors to follow HBR's "Guidelines for Authors" before submitting manuscripts.

To obtain a copy, please visit: HBR.org; email: hbr_editorial@HBR.org; or write to: The Editor, *Harvard Business Review*, 60 Harvard Way, Boston, MA 02163.

A NOTE TO READERS
The views expressed in articles are the authors' and not necessarily those of *Harvard Business Review*, Harvard Business School, or Harvard University. Authors may have consulting or other business relationships with the companies they discuss.

Rotman Management has evolved. Introducing Rotman Intelligence.

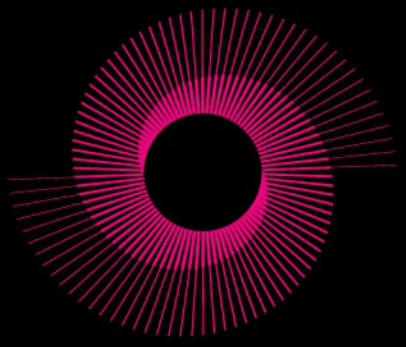

The same trusted source of perspectives on business, strategy, leadership and the accelerating pace of change.

Rotman Intelligence brings together the expertise, learning experiences and viewpoints of the Rotman community, offering thought-provoking insights and practical tools for leaders, innovators and entrepreneurs.

New name.
Same great thinking.

Subscribe today!
Uoft.me/Subscribe
RotmanIntelligence

Rotman INTELLIGENCE

HARVARD BUSINESS REVIEW

Join us at the World Economic Forum Annual Meeting

Davos, Switzerland

January 18–22, 2027

HBR Executive

Lead with insight. Act with impact. HBR.org/executive

In partnership with EgonZehnder

ideaWatch

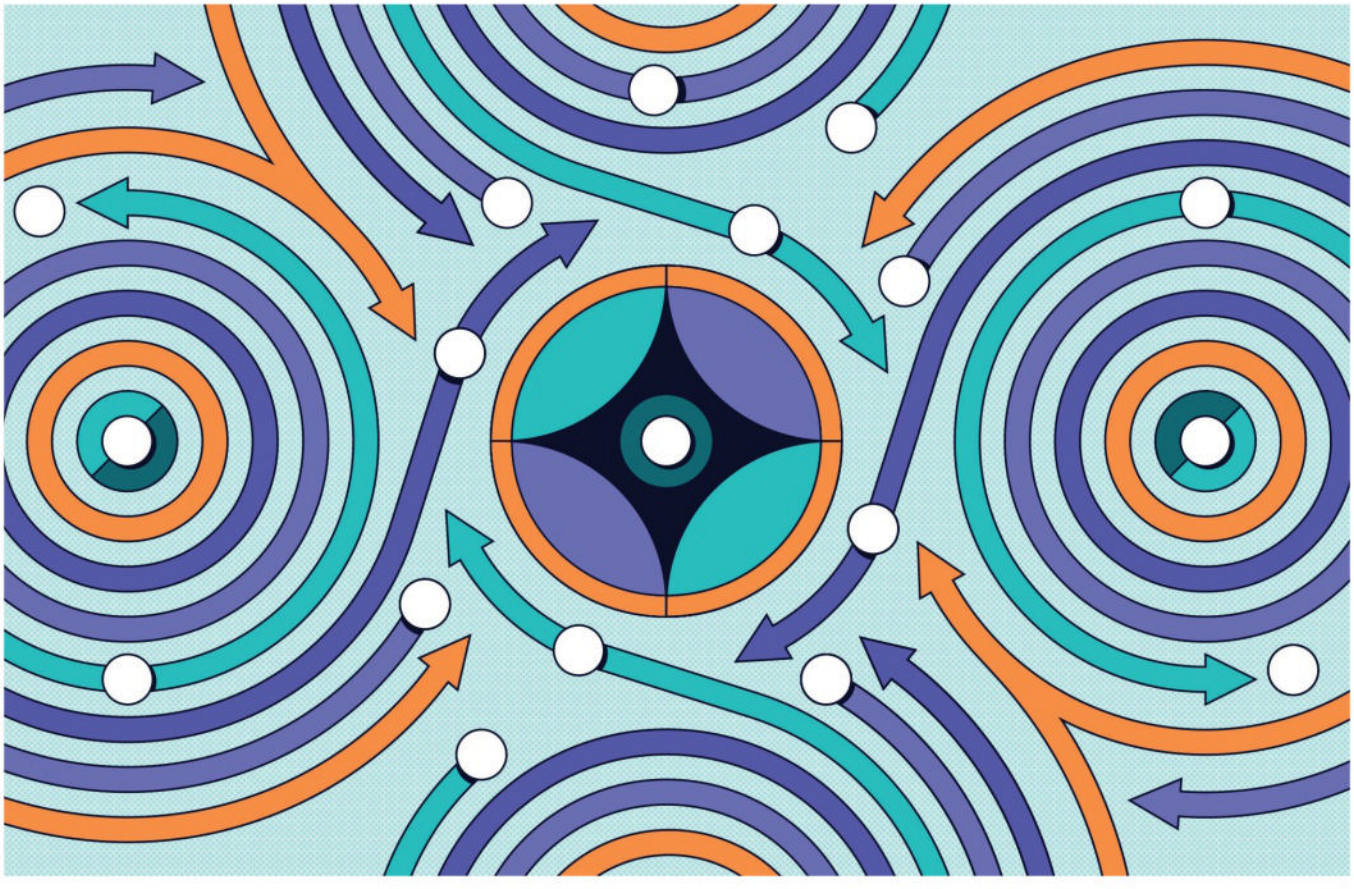

IN THEORY

For Real Performance Gains with AI, Start by Mapping Your Workflows

That's how to reveal the most promising opportunities.

NEARLY FOUR YEARS after the widespread release of large language models, most companies have made significant progress adopting them, but many have not successfully transformed that investment into productivity or performance gains. This disconnect has become a focal challenge that's mobilizing IT departments, affecting tech company stock prices, and raising the anxiety levels of senior corporate leaders.

New research from INSEAD's Hyunjin Kim and Dahyeon Kim and

Illustrations by **TIM BOELAARS**

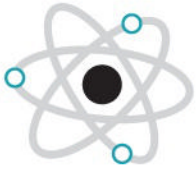

ideaWatch

Managers evaluating AI projects should measure success beyond just improved speed or task output.

Harvard Business School's Rembrand Koning sheds light on the challenge. Managers know that artificial intelligence can perform useful tasks, but identifying where and how it can reshape their businesses is difficult. The researchers call this the "mapping problem"—the challenge of identifying where AI can impact a company's production process and how surrounding activities must change to unlock its potential.

To better understand the mapping problem, the researchers conducted a randomized field experiment with 515 early-stage companies from around the world. Each received access to leading AI tools, technical training, mentorship, and API credits, but only half received examples of how other organizations had redesigned their workflows and business models around AI. The researchers then asked the companies' managers to rethink where to leverage AI across their own production processes.

The results were striking. Companies exposed to real-life examples identified 44% more AI use cases than the control group, completed 12% more tasks, were 18% more likely to acquire paying customers, generated nearly double the revenue, and required significantly less outside capital to grow. Learning about other companies' use cases had expanded managers' thinking about what was possible inside their own businesses, prompting experiments that produced new and better ways of working. That broader perspective led them to identify more, and ultimately more-valuable, applications for AI.

To help other companies get the most out of AI, the researchers suggest the following best practices:

Map your existing workflows. Marketing teams use AI to write copy. Developers rely on coding assistants. Those applications improve productivity but rarely produce the largest gains if the surrounding workflow remains unchanged.

Companies exposed to examples of AI-enabled workflow redesign looked beyond the obvious ways to automate individual tasks. They explored applications across product development, operations, strategy, administration, and other functions.

The lesson for managers is to examine how changes in one activity affect every step that follows. So map existing workflows from beginning to end. Examine where work changes hands, where approvals accumulate, where employees translate information from one system to another, and where decisions stall. Those transition points often offer greater opportunities than the tasks themselves.

It's critical to ask, "What's really slowing us down?" The answer usually isn't the most obvious task. More often it's a handoff or decision that holds up everything that comes after that task. Then drill down further, asking, "What is constraining us from producing 10 times the growth with the organization we have?" Hyunjin Kim says. "Finding the bottleneck that's closest to your customers or closest to the outcome you care about is often the first step."

Expect bottlenecks to reappear elsewhere. Many executives assume that AI removes constraints. Research suggests something more nuanced. When one activity accelerates, another can become an unexpected bottleneck.

Faster product development increases the volume of customer feedback, which means you need to process it more often. Automated reporting exposes weaknesses in data quality, which requires you to address problems you didn't know existed. Generating marketing content with AI is faster than doing it manually, but more content may increase time spent on governance, brand consistency, and legal review. Organizations often discover that yesterday's constraint has disappeared only to reveal another one farther downstream. That dynamic explains why many AI projects initially appear successful before business performance eventually plateaus. Organizations solve one problem but leave the next one untouched.

Treat AI implementation as an ongoing process rather than a onetime deployment. Every major improvement should trigger a fresh examination of where work now slows. The organizations that continuously identify and remove new constraints are more likely to capture value than those satisfied with their first round of automation.

Redesign the entire process. The highest-performing companies didn't stop at identifying bottlenecks. They redesigned the work surrounding them.

Consider product development. AI can generate software code, build prototypes, and create design alternatives in minutes. But those gains create little value if customer testing, approvals, or other downstream activities still move at the same pace.

AI changes the economics of a process only when organizations rethink how the activities fit together, the

research found. Managers evaluating AI projects should measure success beyond just improved speed or task output. "They should ask themselves whether the AI-enabled workflow allows them to reach more customers, produce better products or services, or require fewer resources than before," Hyunjin Kim says. "Otherwise, you're measuring activity rather than impact."

Allow people to experiment. The companies that performed best during the research study learned how others had reorganized around AI—but they didn't merely copy their ideas. Their managers also considered what was possible inside their own businesses. That led to experimentation, which led to new and improved processes.

Leaders can accelerate and scale this process by encouraging teams to explore different use cases, share redesigned workflows across departments, and regularly study how other organizations are reorganizing around the technology. Internal playbooks and cross-functional workshops may ultimately prove to be as valuable as investments in the next generation of technology.

"The biggest AI opportunities may not be obvious, and they don't always require the best, newest technology," Hyunjin Kim says. "You've got to put your people in a position to spot possibilities they never could have foreseen without AI." ■

HBR Reprint F2606A

ABOUT THE RESEARCH *"Mapping AI into Production: A Field Experiment on Firm Performance," by Hyunjin Kim, Dahyeon Kim, and Rembrand Koning (working paper, 2026)*

IN PRACTICE

"Don't Mistake Fancy Technology Demonstrations for Innovation"

Robert Oh is the chief digital and information officer at International Motors (formerly known as Navistar), a 124-year-old Illinois-based manufacturer of commercial vehicles, diesel engines, and buses. He spoke with HBR about why many artificial intelligence initiatives lose momentum after early successes, how his company has approached scaling AI across the business, and why redesigning work matters more than deploying new tools. Edited excerpts of the conversation follow.

How do you encourage employees to find new AI use cases?
I've worked across eight different industries, so I am still interested in what happens elsewhere. I take experiences from other businesses, and I turn them into six- to 10-page white papers to educate our company. The goal is to organize and share that information so that teams can apply it to their work and even feed it into their AI tools and agents to discover new uses. We also bring in external AI partners to work alongside our business and technology teams in six-week co-innovation pods. Those outside perspectives give people more context when they are thinking about our own sales, supply chain, manufacturing, and other operational processes.

Many companies see an initial burst of AI productivity before progress levels off. Did that happen for you?
Yes. Last year we encouraged employees to experiment with AI, identify opportunities in their own work, and build confidence with the technology. That generated dozens of experiments across the company and created real excitement. But we also realized we were improving individual tasks more than the business itself. The technology was helping but not as much as we had hoped.

Why did that happen?
It was our existing processes. They were slower than they needed to be. But technology alone can't fix that. Don't mistake fancy technology demonstrations for innovation. Instead, we started treating AI as an operating model change. We brought together our executive leadership team; mapped cross-functional workflows across sales, manufacturing, and the supply chain; and identified where AI could create measurable business value. Our employees generated a number of AI use cases across the company. That gave people confidence, but it wasn't enough to change the business. We then identified 21 cross-functional opportunities spanning sales, manufacturing, the supply chain, and service; attached an expected business value to each one; and then prioritized the list to the five we believed would have the biggest impact.

Can you offer any examples?
One area we focused on was after-sales parts. We're a truck manufacturer, and our customers make money only when their vehicles stay on the road, so parts availability and service are major competitive advantages. We worked with our after-sales organization to use AI to improve parts availability and to make more-consistent decisions.

Another team identified an opportunity to improve our accounts payable process while we were transitioning to a new enterprise-resource-planning system. Instead of relying on manual invoice reviews, we built a machine-learning system that flags duplicate and incorrect payments before they're made. The project went from idea to production in about 12 weeks and is expected to save millions of dollars annually while also protecting working capital. ■

Illustration by **JORI BOLTON**

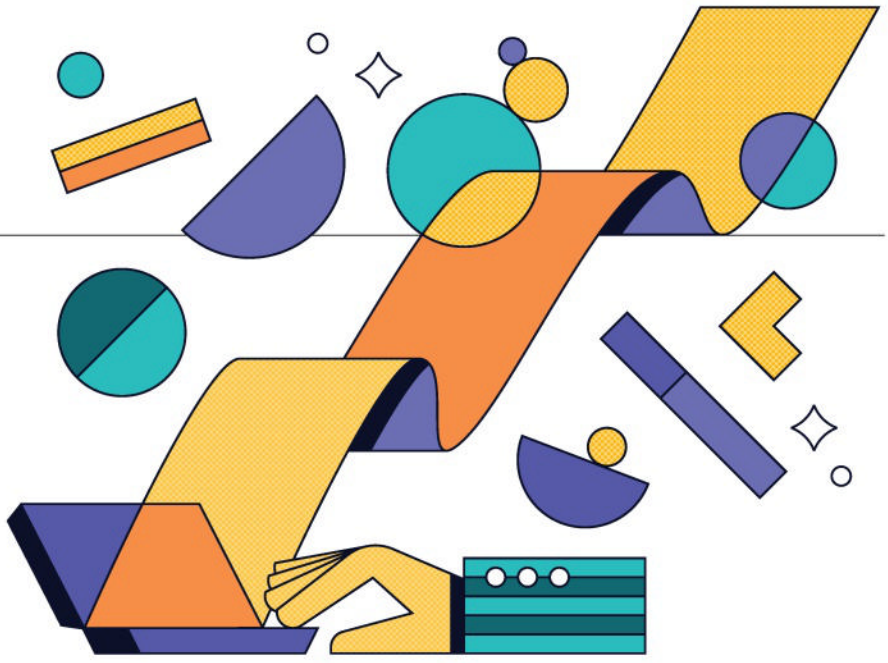

REMOTE WORK

When Jobs Go Online, Requirements Soar

Remote work is often viewed as a change in where work happens. But it can also change what a company looks for in potential hires.

Researchers analyzed more than 50 million online job postings across 28 European countries from 2018 to 2021, identifying which roles were fully remote and extracting the skills, education credentials, and work experience employers requested. To make an "apples to apples" comparison, the analyses compared postings within the same occupation, the same employer, the same country, and the same year whenever possible—so that the key difference was whether the job was remote or in person.

They found that remote postings demanded more—a lot more—from applicants. Compared with similar in-person roles, postings for remote work listed roughly 25% to 42% more required skills and asked for more prior experience. The remote requirement premium was largest in occupations that typically have more on-the-job training and in roles that had previously been hard to fill with qualified local applicants. It was smaller in regions that had high unemployment, consistent with the idea that when local labor markets are slack, employers can already be more selective, even without going remote.

An online experiment involving 1,250 participants with hiring experience echoed those results: People set the bar higher for both skills and experience when hiring for remote roles. Interviews with executives and HR managers and follow-up analysis suggest several reasons for that, including anticipated difficulty in training and supporting workers at a distance and the larger applicant pool that remote roles attract, making firms raise requirements to help winnow applications.

The researchers argue that their findings should prompt executives to view remote work not primarily as an issue of flexibility or cost savings but as a choice about the human-capital model they are building. Leveraging remote arrangements to tap a larger talent pool can yield highly skilled hires, but raising requirements may disadvantage less experienced, early-career applicants and make internal mobility more difficult, as current employees will need to compete against more-accomplished applicants.

ABOUT THE RESEARCH *"Remote Work and Hiring Requirements: Cross-Country Evidence from Job Postings," by Shinan Wang, Letian Zhang, and Zhenyu Liao* (Administrative Science Quarterly, *2026)*

NEGOTIATION SKILLS

People Prefer to Negotiate with Women

For decades, studies have found that women are less assertive negotiators than men are, which puts them at a disadvantage at the bargaining table. But new evidence suggests the opposite: Women negotiators may reap long-term economic gains because people prefer negotiating with them.

Across five studies with 2,401 U.S. participants, a team of researchers examined both objective outcomes (money, points) and subjective outcomes (trust, likability, satisfaction, and willingness to negotiate again) in negotiations that had a variety of gender makeups. In an MBA negotiations course that generated more than 2,000 evaluations, women were rated higher than men were on building trust, being fair, listening, communicating, and satisfying their counterparts' needs. Classmates were also more likely to say they wanted to negotiate with their women counterparts again. Overall, however,

men and women achieved equivalent economic results in their deals.

This pattern held in an online negotiation study, where partners were anonymous and communicated only via text. Third-party observers tried to guess the gender of negotiators based on the transcripts, and their accuracy was no better than chance. Yet the women participants had been liked more than the men had, and liking was correlated with greater satisfaction for their partners, again with no economic difference in their deal outcomes. In later experiments, third-party observers rated women negotiators as both warmer and more competent than men and expressed a stronger desire to work with them in future deals, whether as teammates or opponents in competitive or cooperative negotiations.

An analysis of thousands of negotiation messages suggested one reason: Women were more likely than men to accept offers, a form of yielding that makes partners feel heard and treated fairly, and doing so did not hurt women's economic outcomes. A simulation based on the data suggests that, over time, women could receive roughly 45% more opportunities to negotiate—advantages that would compound financially.

For organizations, the findings challenge prevailing narratives about women's negotiation skills. Rather than trying to "improve" how women employees negotiate, companies should recognize and reward the relational strengths that make people eager to work with them.

ABOUT THE RESEARCH *"People Prefer to Negotiate with Women, Even When Outcomes Are Identical and Gender Is Unknown," by Charlotte H. Townsend, Laura J. Kray, and Solène Delecourt* (Proceedings of the National Academy of Sciences of the United States of America, *2026*)

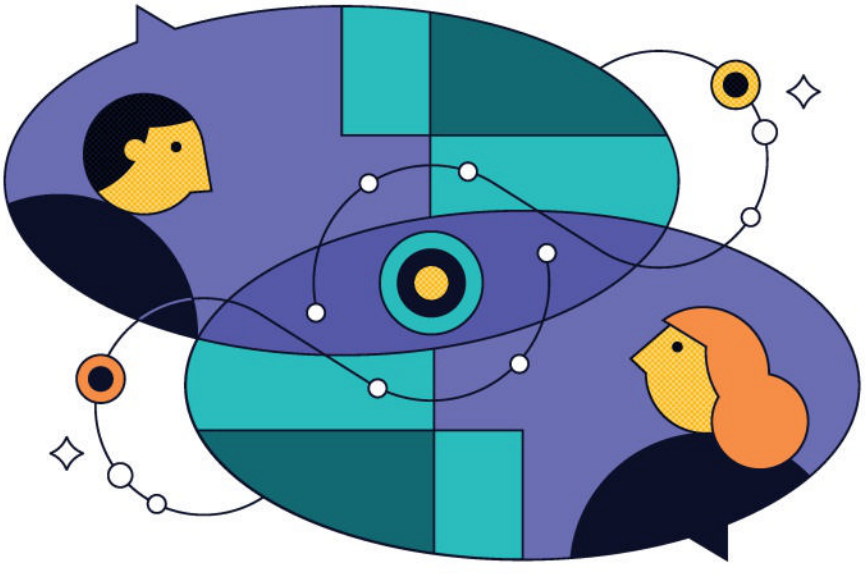

HIRING & RECRUITMENT

The Founder Qualities That Attract AI Talent

Many startups are racing to lock down top AI talent, but what grabs the attention of these professionals when founders reach out? To answer that, researchers used data from more than a million interactions on a talent-matching platform for startups. They wanted to know which founder attributes—such as having a law degree or previous work experience with a major tech company—made people more likely to accept a founder's invitation to connect and begin chatting. They discovered that founders who signaled that they were working full-time toward their idea and had already made some progress were more likely to have their invitations accepted. For the sake of comparison, researchers identified which founders had a background in AI and which did not.

Effect on invitation acceptance rates, by founder attribute and AI experience

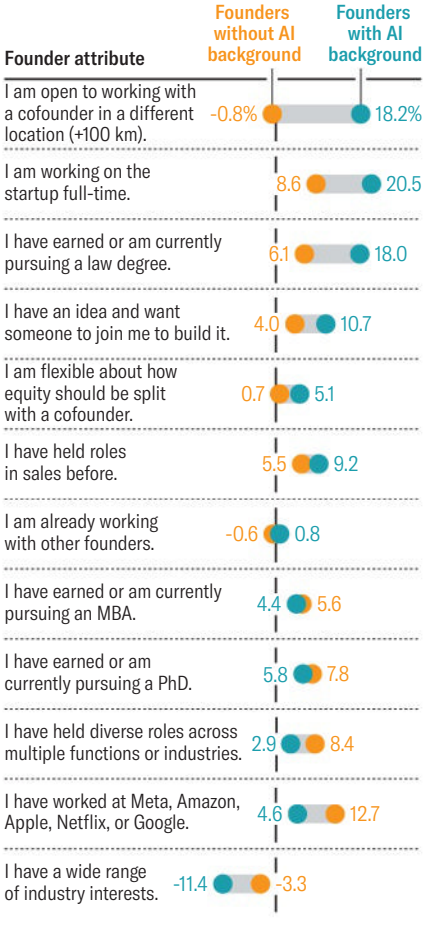

Source: CoffeeSpace, Fauzan Reza Maulana, Travis Howell

ideaWatch

ARE YOU BIASED TOWARD CANDIDATES WHO REPLY QUICKLY?

Employers may say they don't expect immediate responses from job applicants, but candidates may want to stay poised by their inboxes all the same. In a study of more than 11 million interactions on a global freelancing marketplace, researchers found that a one-hour lag cut applicants' chances of getting hired by about 46%, and a full day by roughly 90%.
"Speed Is a Signal: When Faster Replies Increase Hiring Likelihood," by Einav Hart et al.

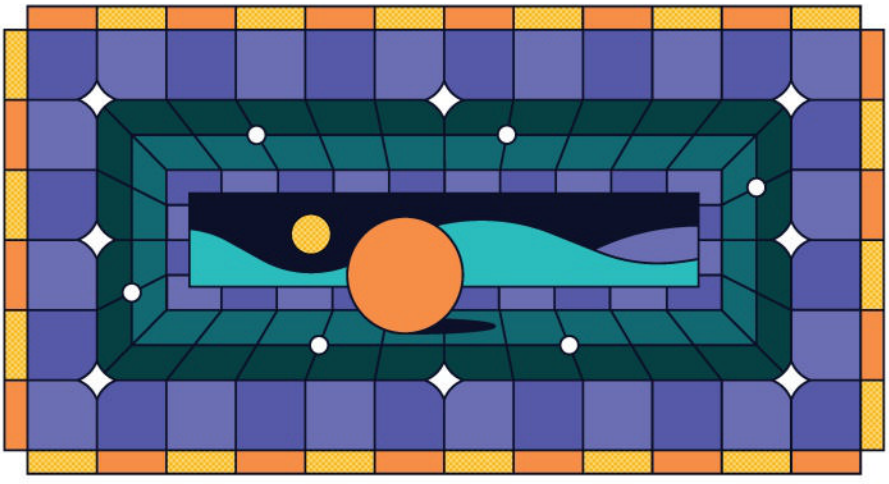

TECHNOLOGY & ANALYTICS

AI Agents Don't Just Speed Up Work—They Stretch What People Do

Many companies have rolled out AI tools to their workers. And they've seen efficiency gains when their employees use AI assistants to offload or speed up familiar tasks. But the biggest payoff comes when people start using AI agents. Because with them, employees don't just get faster at their work, they expand the scope of what they work on.

From February to May 2026, researchers analyzed hundreds of thousands of real-world AI sessions in two Perplexity tools: a conversational assistant and an autonomous agent. They found that users around the globe mostly used the assistant for short, bounded tasks, such as answering a question or polishing text. But because an agent can pull from outside sources (a user's email, logged-in accounts on websites, corporate databases, and more), its sessions unfolded as chains of autonomous actions: navigating websites, restructuring spreadsheets, and delivering a finished product. User behavior changed accordingly. Rather than manually doing the work themselves, they spent their time verifying or extending it and reported lower dissatisfaction with the agent than with the assistant.

That autonomy also led people to use the agent for a broader range of work. They were much more likely to make requests that demanded higher-order cognition: Seventy-six percent of agent requests involved analysis, creation, or abstraction, compared with 55% for assistants. Agents were asked to cross occupational boundaries more often, such as managers developing software or engineers building financial models. Most telling, about 23% of agent queries contained a request that never appeared in the same users' assistant logs at all. Those new tasks were, on average, more complex than the rest—the sort of work that only became practical once an agent could execute it end-to-end.

Because agents expand the type of work employees can take on, the researchers write, leaders need to redesign organizational roles to clarify accountability on consequential actions. What's more, firms should start measuring not only how much time is saved by agents but also what new, cross-functional work is made possible.

ABOUT THE RESEARCH *"How AI Agents Reshape Knowledge Work: Autonomy, Efficiency, and Scope," by Jeremy Yang et al. (working paper, 2026)*

COLLABORATION & TEAMS

Are Your Employees Penalizing Referred Hires?

Employee referral programs can bring in strong candidates quickly and cheaply. But they might also carry a long-term cost if your current employees view these new hires as undeserving—and treat them accordingly.

Across four studies in Europe and the United States, researchers examined how "regular" employees perceived and behaved toward coworkers who came in through a referral instead of through standard external channels. In two field studies, participants named a real colleague and reported both how that person was hired and how they treated that coworker day-to-day.

Stop reorganizing. Start adapting.

Becoming an Octopus Organization
Learn to break down hierarchy, get rid of silos, and create an adaptable, resilient, learning organization in this eight-week newsletter series.

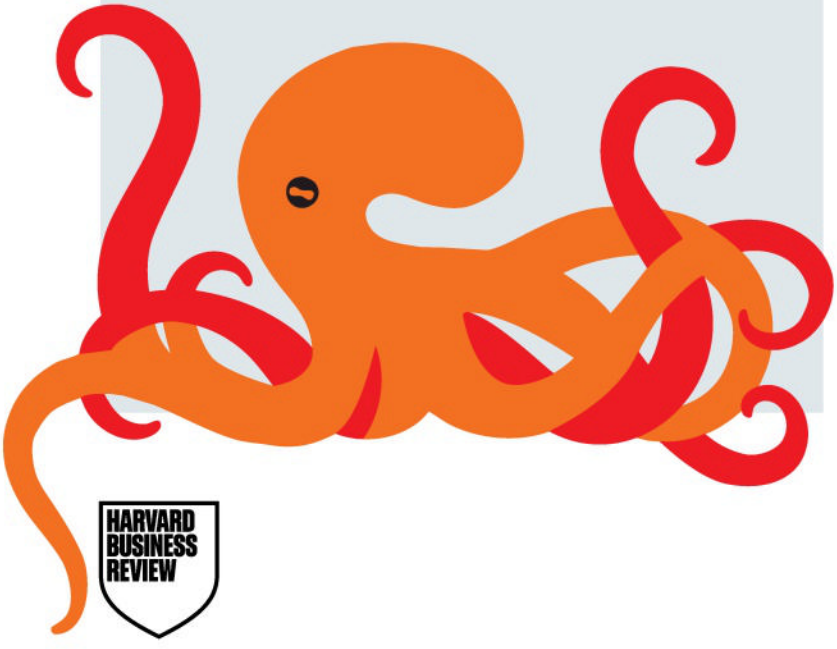

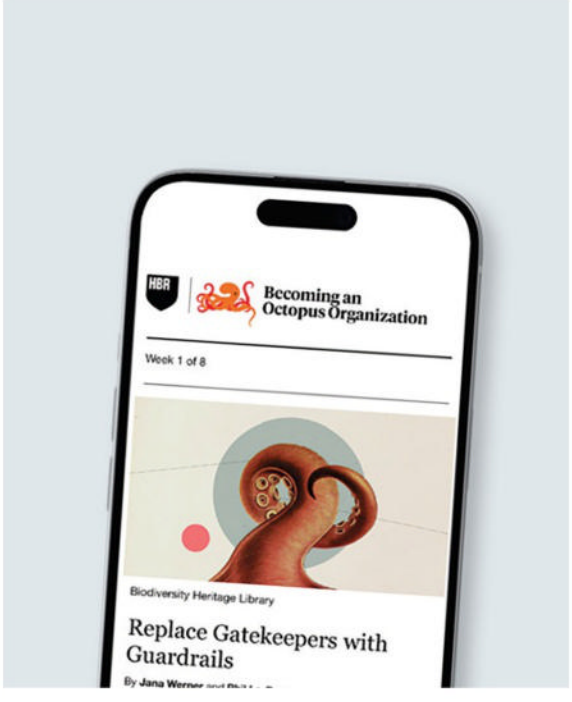

HARVARD BUSINESS REVIEW

SIGN UP TODAY
hbr.org/newsletters

Referred coworkers were consistently rated as less deserving of their job than nonreferred ones, even though they were performing similarly. People also reported that they were less likely to help referred colleagues (for example, by sharing information) and more willing to engage in subtle antisocial behaviors (such as withholding credit).

In another experiment, participants read vignettes about a new teammate whose qualifications and hiring process were clearly described. Even when the scenario stated that a referred newcomer had all the required credentials, went through the same vetting process as everyone else, and was performing above average, the person was still seen as less deserving than an equally strong nonreferred hire. And when people believe that a connection—not competence—has led to an applicant's success, they are less willing to help that person and more willing to punish or report that person for missteps. The perceived mismatch between "inputs" (merit) and "outputs" (getting the job) triggers equity-restoring behavior that undermines referred hires' socialization.

Firms must design referral programs that emphasize legitimacy and detail how fairness is upheld, the researchers say, so that companies can help ensure that the benefits of referrals aren't erased by peer backlash.

ABOUT THE RESEARCH *"The Referral Penalty: Decreased Perceptions of Merit Undermine Helping Behavior Toward Referred Employees," by Teodora K. Tomova Shakur and Rellie Derfler-Rozin* (Journal of Applied Psychology, *2026)*

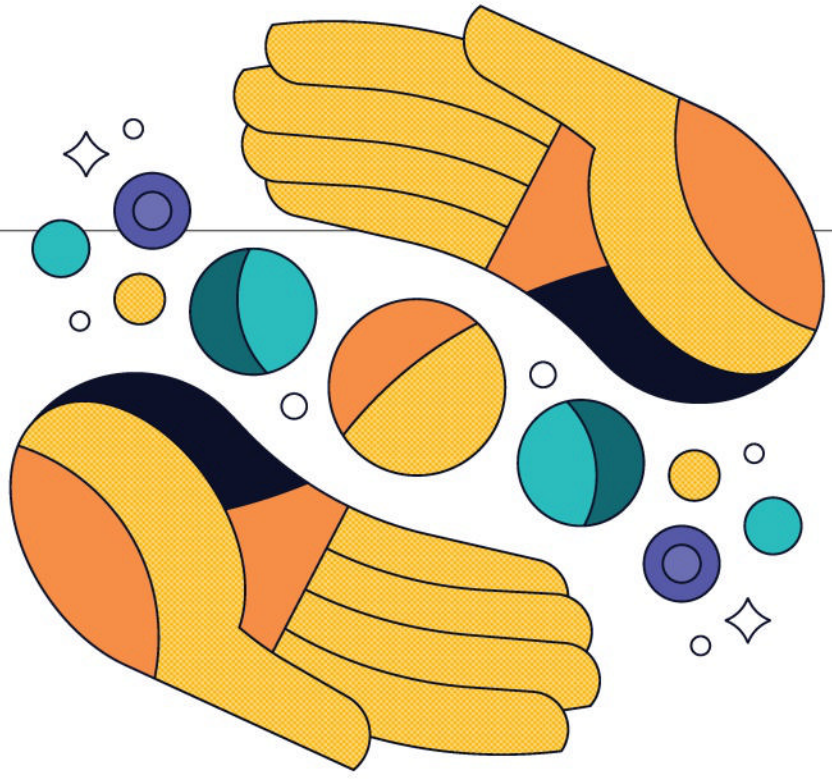

LEADERSHIP SKILLS

For Executives, It Pays to Be Cooperative

On earnings calls, some executives answer every analyst question themselves. Others routinely pull colleagues in: "Let me have our COO take that." New research suggests that the small act of sharing the mic signals something big—cooperativeness—and that it gives a boost to executives' careers.

The researchers analyzed 78,000 U.S. earnings calls from 2010 to 2020 and measured how often an executive explicitly invited a peer to answer an analyst's question. By adjusting for how frequently that executive spoke overall, they determined that person's cooperativeness score. The sample covered 10,864 executives at 2,486 public firms.

Executives who scored higher on cooperativeness, the researchers found, were more likely to serve on charitable boards, work in firms whose language and job ads emphasized teamwork, and be based in communities with higher social capital (measured by voter turnout, civic club density, and more). Older, more senior, and better-paid leaders were more cooperative, and cooperativeness was unrelated to short-term performance or negative analyst questioning, suggesting that sharing the mic wasn't simply a way to deflect blame. Cooperative behavior was persistent over time, though it tended to rise when a more cooperative CEO took over, indicating a spillover effect on the rest of the executive team.

When the researchers examined the career trajectories of executives, they found that among non-CEO members of the C-suite, executives whose cooperativeness scores were above average were almost twice as likely to get promoted in the next year. Those who engaged multiple colleagues—not just a single ally—had the highest odds of promotion.

HOW TO TELL WHICH CUSTOMERS WILL ESCALATE A PROBLEM

A new study that analyzed more than 160,000 service complaints filed with the U.S. Consumer Financial Protection Bureau found that customers who use the passive voice (for example, "A mistake was made") are more likely to dispute a company's proposed resolution than other complainers are. They also report stronger intentions to post negative reviews, bad-mouth the company to acquaintances, and pursue legal remedies. "Passive Voice in Consumer Complaints Reveals Fault Attribution and Predicts Escalation," by Amir Sepehri, Rod Duclos, and Jonah Berger

Investors seem to value this trait too. When firms appointed a CEO who was more cooperative than the preceding one, the stock gained about 0.6% in abnormal returns in the three days surrounding the announcement.

The researchers note that the cooperative advantage disappears in firms with very low teamwork cultures, where the trait clashes with how the organization operates. But in most settings, sharing the stage, not owning it, appears to be a hallmark of leaders on the rise.

ABOUT THE RESEARCH *"Executive Cooperativeness: Evidence from Conference Calls," by Wei Cai et al.* (Management Science, *2026*)

MANAGING UNCERTAINTY

Why Some People Resist Reskilling

Learning new skills can seem like the obvious answer for workers who are worried about potential layoffs or new technological change. But the biggest barrier to reskilling might not be time, money, or even confidence—it might be how workers see themselves.

In a study of about 1,100 job seekers in Italy from 2023 to 2024, researchers ran a "shopping-style" training experiment. Participants were given the choice to take a generic course in Microsoft Office skills, take an occupation-switching reskilling course, or take no course at all. The reskilling track was randomly assigned to one of two high-demand roles—IT support or construction technician (a white-collar coordination role)—and each option's features varied on factors such as out-of-pocket price per hour, course length, delivery mode, and the probability of finding a job after completion.

The researchers used those choices to determine how much people were effectively "willing to pay" for reskilling. They found that job seekers would pay about €3.5 an hour to learn the skills needed to land an IT support role, but they would need to be paid about €5.3 per hour to reskill into the construction technician role. A survey of the participants revealed that the strongest driver for choosing to reskill wasn't expected wages or skill overlap; it was "identity fit," or how closely the job felt aligned with participants' self-perceptions.

The researchers also tested two interventions. The first, a short informational video with wage and vacancy data, increased people's willingness to take a reskilling course. But the second, a video on cultivating a growth mindset, had no impact.

For a subset of the job seekers, the researchers also had administrative records that showed the effect of those videos over the following 12 to 17 months. The informational video increased real IT-related training enrollment and IT-sector employment—but not construction—for the same reason: Viewers didn't *see* themselves as people who worked in construction.

The researchers say that to encourage reskilling at scale, you can't market it like a generic benefit or a future-proof tactic. Instead you need to target each new role to people who can imagine themselves in it—and emphasize how the new skills will help them get hired. ■

ABOUT THE RESEARCH *"Unwilling to Reskill? Experimental Evidence from Real-World Jobseekers," by Alexia Delfino et al. (working paper, 2026)*

MARKET SEGMENTATION

Make Your Local Ads More Effective

Location-based advertising often relies on a simple strategy: Draw a radius around a store and target everyone living inside that area. But that approach is flawed, according to an analysis of television ad exposure and millions of U.S. store visits. What really matters for ad performance isn't distance alone—it's whether your store is closer to a customer than your competition is. The study found that relative proximity was the strongest predictor of subsequent in-store visits.

Advertising effectiveness, by customers' proximity to focal and rival stores

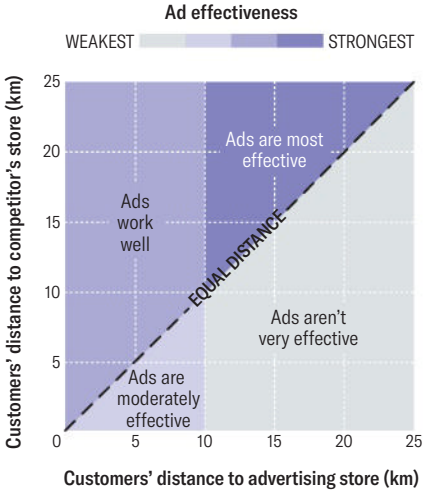

Source: Bowen Luo and Bhoomija Ranjan

HARVARD BUSINESS REVIEW

HBR Innovation Summit 2026

Turning breakthrough ideas into real results

LIVE VIRTUAL EVENT | NOVEMBER 12–13, 2026

Innovation requires new thinking in the era of AI. Where should you invest? What capabilities should you build? How do you create lasting value?

The HBR Innovation Summit features conversations with the CEOs of FedEx, IDEO, HCLTech, and Lumen Technologies, as well as research-based masterclasses from leading faculty. Explore how to integrate AI and emerging technologies, rethink business models, and scale innovation across your organization.

Leave with practical frameworks and ideas to turn innovation into measurable business results.

HARVARD BUSINESS REVIEW HOSTS:

Adi Ignatius
Editor at Large

Alison Beard
Executive Editor

VIEW THE AGENDA AND REGISTER TODAY
hbr.org/innovationsummit

One Challenge, Four Strategies

inFocus

akinbostanci/Getty Images

THE CHALLENGE

Collaborating Amid Complexity

1 Four Hidden Traps of Team Dynamics 28

2 Prevent Team Friction from Turning into Dysfunction 31

3 The Office Amenity That Actually Improves Teamwork 34

4 How Highly Effective Global Teams Collaborate Across Cultures 36

These articles were selected from **HBR.org**.

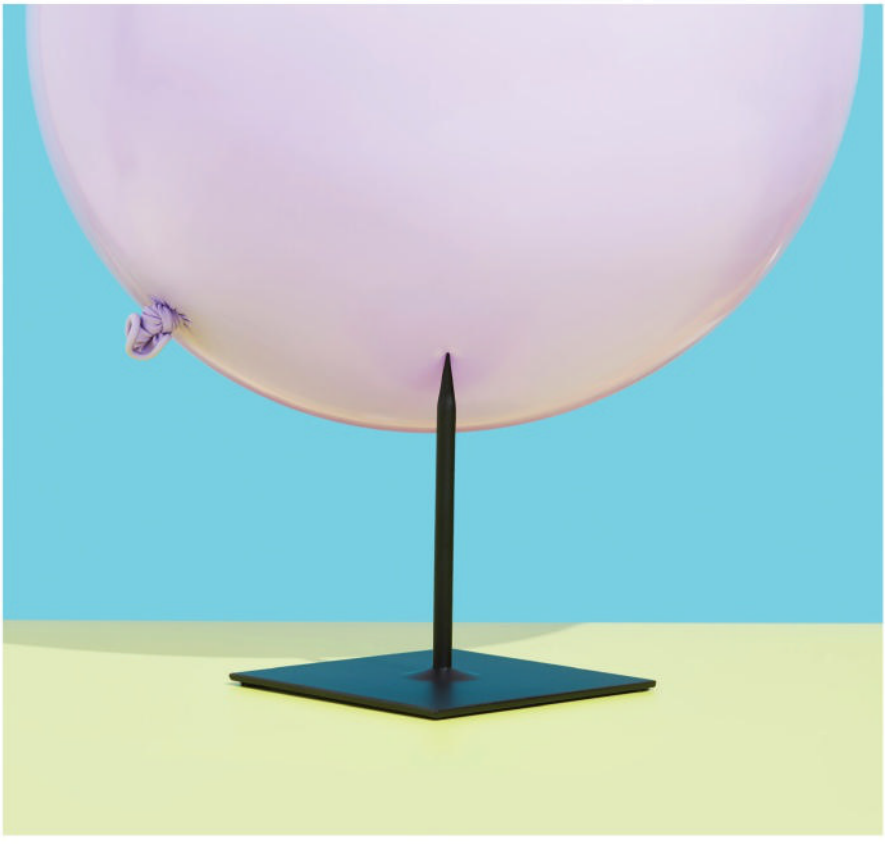

Four Hidden Traps of Team Dynamics

by SUSAN MACKENTY BRADY, STUART D. KLIMAN, and LESLIE C. SMITH

ORGANIZATIONS TODAY ARE incredibly diverse in every sense of the word. Teams span generations, cultures, functional specialties, educational backgrounds, political perspectives, and increasingly, relationships with technology and artificial intelligence. Work is more distributed, expectations are more personalized, and the pace of change is continually accelerating. Leaders are asked to unite people who see the world differently, solve increasingly complex problems, and move quickly despite uncertainty. Unfortunately, many organizations struggle to capture the full value of the differences present in their workforce.

The problem is not difference itself. Difference is often the source of adaptability, innovation, better decision-making, and stronger problem-solving. The challenge is that human beings are not naturally equipped to navigate difference without creating friction. When perspectives collide, assumptions are challenged, or tensions arise, leaders often fall into predictable behavioral traps, which erode trust, weaken collaboration, and diminish performance. We call these traps leadership "land mines"—hidden dynamics that can derail even the most well-intentioned leaders.

We identified these land mines from our work with all kinds of leaders and teams. But we also saw them in our own partnership. One of us has devoted her career to developing leaders as a CEO of a leadership institute and an executive coach. Another has spent his career helping organizations negotiate, collaborate, and solve high-stakes challenges. The third spent nearly four decades leading soldiers in some of the world's most complex environments. We have diverse professions, experiences, and ways of seeing the world and have witnessed firsthand how quickly well-intentioned people—including ourselves—can interpret the same situation in different ways. Those moments became the source of our best thinking and reminded us that difference is not what creates dysfunction. How we *respond* to difference is what determines whether value is created or diminished.

Understanding these land mines is the first step toward turning difference from a source of friction into a source of value.

Four Land Mines That Derail Leaders

Executives should be mindful of the following team dynamics traps:

1. Certainty. Perhaps the most common land mine is certainty: the conviction that our interpretation of a situation is correct and that opposing viewpoints are simply wrong.

Every leader has experienced this. In a strategic debate, one executive argues for aggressive investment; another insists on cost containment. One person advocates rapid action; another urges caution. Eventually the discussion stops being about exploration and starts becoming a contest over who is right.

In reality, people attend to different information, draw on different

Richard Drury/Getty Images

experiences, and interpret events through different mental lenses. Two intelligent, well-informed people can look at the same situation and reach different conclusions.

In one executive team we worked with, one member, Asher, felt a compelling need to focus planning discussions on making the next quarter's numbers. His colleague, Jen, wanted to develop a strategy that would better equip the business to grow over the long term. They both strongly believed they were right and the other was wrong.

We can all think of a conversation or issue we felt absolutely certain about. The evidence seemed clear. The conclusion felt obvious. Anyone who didn't see the situation similarly appeared to be missing something important. It is precisely in these confident moments that certainty becomes dangerous, because conviction quietly replaces curiosity.

Certainty feels efficient because it creates clarity. But it often closes off the very conversations that generate breakthrough thinking. Leaders who become overly attached to being right risk leaving valuable perspectives—and valuable solutions—on the table.

2. Inconsistency. The next land mine is inconsistency: the gap between what leaders say they value and how they actually behave.

Most leaders aspire to be collaborative, inclusive, respectful, and open-minded. Yet under pressure, their behavior doesn't always align with those intentions. Sometimes they know it; often they don't.

A manager may repeatedly emphasize the importance of teamwork while consistently rewarding individual achievement. A leader may genuinely value listening yet routinely interrupt others during meetings. Here are a couple of familiar examples:

- A vice president publicly commits to increasing gender diversity within his organization. Months later he misses a major internal event celebrating women in the company because of another professional obligation. The specific decision may have been entirely reasonable, but employees perceive a disconnect between the leader's stated values and his visible actions.
- An IT leader sends out a new comprehensive guide to using AI with integrity. She recommends that employees add a disclosure when a significant percentage of the output of a project or piece of writing is generated using AI. Yet she doesn't reveal that she used AI to create nearly all of the guide.

The challenge is not necessarily hypocrisy. It is humanity. Every leader has blind spots that can create a gap between intention and impact.

What makes inconsistency dangerous is that employees pay far more attention to what leaders do than to what they say. When words and actions diverge, trust begins to erode.

3. Emotional reactivity. Another land mine is emotional reactivity: the way that people respond differently, and often intensely, to the same event.

One employee shrugs off a blunt comment in a meeting; another spends days replaying it in his mind. One team member welcomes direct feedback; another experiences it as personal criticism. One person sees healthy debate; another sees disrespect.

The difficulty for leaders is that emotional reactions are shaped by factors that are often invisible: personality, life experiences, cultural norms, stress levels, previous interactions, and even physical well-being.

Imagine that a colleague brings a fresh idea to the weekly team meeting. Another colleague immediately starts critiquing it and offering other ideas. The first colleague reacts—but everyone else in the room sees, understands, and responds to that reaction differently. One teammate is visibly frustrated and doesn't participate in the discussion. Another says that the new perspectives are helpful and adds her own ideas. A third sees the entire conversation as off topic and moves to shut it down. All are reacting to the same stimulus but responding differently. That undermines the team's ability to execute—or to tap into all the thoughts, feelings, points of view, and experiences that sit underneath the varied reactions. When leaders assume that everyone shares their interpretation, misunderstandings multiply.

Emotional reactivity becomes especially dangerous when leaders dismiss others' reactions as irrational or oversensitive. Doing so often compounds the problem, causing people to feel unheard and invalidated.

The goal for leaders is not to eliminate emotional reactions—that would be impossible. They need to recognize that feelings are part of the landscape and to approach them with curiosity rather than judgment.

4. Justification. The final land mine is justification: our tendency to explain, defend, and rationalize our

Leaders must understand their own preferences, biases, triggers, assumptions, and default reactions.

own behavior rather than examine it critically.

When leaders are certain, they justify why they're right. When they behave inconsistently, they justify why the circumstances made it necessary. When they react emotionally, they justify why their response was appropriate.

A leader misses an important employee event and explains all the reasons that the decision made sense. A manager dismisses a colleague's concern because he was "just being realistic." An executive avoids a tough conversation because bringing it up would "only make things worse." In each case, the explanation may be reasonable, but the use of justification often prevents learning.

Justification of inaction may cause the most damage. Leaders frequently convince themselves that avoiding a difficult issue is prudent, when in reality it is merely making them more comfortable. The result? A conflict goes unaddressed. A problematic comment is ignored. A strategic disagreement remains unresolved. The issue doesn't disappear. It simply moves underground, where it quietly drains employee trust, engagement, and performance.

Justification transforms moments of potential growth into cycles of defensiveness. It turns curiosity into certainty, and accountability into blame.

Six Actions That Help Leaders Defuse the Land Mines

Fortunately, leaders are not powerless. They can develop the skills required to lead through the land mines by taking these six steps:

1. Know yourself. Everything starts with self-awareness. Leaders must understand their own preferences, biases, triggers, assumptions, and default reactions. They need to recognize when certainty is taking hold, when emotions are escalating, and when they are rationalizing behavior rather than examining it.

Here is a simple starting point: After a difficult interaction, ask yourself, "What was I feeling? What assumptions was I making? What might I have missed?"

The more leaders understand themselves, the better they can understand others.

2. Return to respect. Even the best executives get triggered. What distinguishes effective leaders is their ability to recover. When conversations become tense, they pause, reset, and reengage from a place of respect.

Suppose a meeting becomes heated. Rather than doubling down, a leader might say, "We're clearly seeing this differently. Let's slow down and make sure we understand each other before moving forward."

Respect is not agreement. It is the decision to treat others' perspectives as worthy of consideration.

3. Activate honesty. Trust grows when leaders tell the truth—even when the truth is incomplete or uncomfortable. Honesty means acknowledging uncertainty, naming tensions, and sharing what you genuinely think rather than hiding behind organizational scripts.

For example, a leader facing a difficult change initiative might say, "I know people have concerns. I don't have all the answers yet, but I want us to discuss your worries openly."

Such candor often creates more confidence than false certainty. Let's be clear, however. This kind of truth telling doesn't throw away empathy in the name of honesty. Leaders must be aware of the impact the message being delivered has on others.

4. See the full story. Most conflicts are fueled by incomplete information. Effective leaders resist the urge to rush to judgment. Instead, they seek additional perspectives and explore what may lie beneath the surface.

Here is a practical habit: Before forming a conclusion, ask, "What else could be true?" This simple question creates room for empathy, curiosity, and better decisions.

5. Ignite togetherness. Strong teams do not think alike. They think together.

Leaders create value when they shift the focus from "my solution" to "our solution." The best ones invite diverse perspectives early, encourage collective problem-solving, and help people feel ownership of the outcomes. They create a sense of "we're in this together" by knitting together disparate perspectives and views and tapping into those differences to create value.

A simple first step is to ask, "Whose perspective haven't we heard yet?" That question alone can transform the quality of a discussion.

6. Commit to action. Leading through difference is not a onetime event. It is an ongoing practice that starts with staying self-aware, respecting yourself and others, telling the truth even when it's hard, staying actively curious, and creating alignment.

Leaders must consistently reinforce expectations, address issues directly,

and hold themselves accountable for the culture they create.

After a difficult interaction, a leader might revisit the conversation, acknowledge any missteps, and clarify expectations moving forward. That accountability strengthens trust rather than weakens it. The goal is not perfection. It is progress.

Turning Difference into Advantage

Difference is unavoidable. In today's organizations, it is also one of the greatest sources of untapped value. The leaders who succeed at creating value from difference will not be those who eliminate tension or create perfect harmony. They will be those who recognize the hidden land mines that accompany human difference and develop the skills to navigate them with awareness, humility, and intention.

When leaders learn to do that, difference stops being a problem to manage and becomes an advantage to leverage. ■

HBR Reprint H0991A

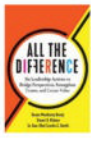

SUSAN MACKENTY BRADY *is a global leadership expert and the CEO of the Simmons University Institute for Inclusive Leadership.* **STUART D. KLIMAN** *is the chief strategy officer at Breez AI and the head of business development at PrescriberPoint.* **LESLIE C. SMITH**, *a retired U.S. Army Lieutenant General, is the vice president for leadership, education, and programs at the Association of the United States Army. This article is adapted from their book* All the Difference: Six Leadership Actions to Bridge Perspectives, Strengthen Teams, and Create Value *(Harvard Business Review Press, 2026).*

MirageC/Getty Images

Prevent Team Friction from Turning into Dysfunction

by LEONARD A. SCHLESINGER, JOSEPH FULLER, and ROBERT V. TOOMEY

IT WAS THE SECOND night of working on a global consulting project at a top-tier telecommunications company in Hanoi when a team of MBA students found itself stuck.

The team was at the tail end of a meeting in which one student had been assigning tasks, locking down timelines, and pushing for closure. Suddenly another student, who had been quiet throughout the conversation, introduced a new strategic perspective that reopened decisions the rest of the team thought were settled. Frustration grew. The discussion began to loop without resolution.

Similar scenes occurred during the 2024 session of Harvard Business School's signature FIELD Global Immersion program, in which nearly 1,000

MBA students spent 10 days working in small, nonhierarchical teams on real-time business challenges in 16 countries.

The experience mirrors the pressures many corporate teams face daily: ambiguous data, unfamiliar markets, cross-cultural dynamics, and tight timelines. But given the level of talent—students who'd been admitted to a highly competitive business school—most people would assume that the teams would function smoothly.

They didn't. By the end of the program, more than a quarter of the teams experienced communication breakdowns serious enough to require faculty intervention. It wasn't because the students lacked intelligence or motivation; it was because they didn't have a shared way to work through cognitive and communication differences.

Beyond Intelligence and Motivation

What we observed among these MBA teams is not unique. Variations of the same dynamic play out every day inside organizations. Modern teams often consist of highly capable people with very different ways of processing information, approaching decisions, and communicating under stress. Without a shared framework for navigating those differences, friction escalates, trust erodes, and execution slows.

Under normal conditions, differences may seem manageable. Under stress, however, they become amplified. One team member pushes for rapid closure, while another continues exploring alternatives. One interprets direct feedback as constructive advice, while another experiences it as dismissiveness. One needs time to process before responding, while another interprets that silence as disengagement.

In each case, the issue is not capability. It is interpretation. As we looked at the groups' dynamics, we concluded: Teams don't fail because their members disagree. They fail because people misread one another's intent.

Two of us (Joseph and Leonard) are Harvard Business School faculty members. As leaders of the FIELD program—Joseph is the current course head, and Leonard preceded him in that role—we were forced to confront a difficult reality: We had rigorously trained students to analyze markets, build strategies, and execute plans, but we had not fully prepared them to understand how individual people process information, respond to pressure, and work through conflict together. We hadn't sufficiently trained them to work well in a team.

The Rise of Interpersonal Competence

That realization led us to rethink part of the immersion experience and ultimately to partner with TypeCoach, an organization that combines training and online tools to help business teams adapt more effectively across differences in how people think, communicate, and work. (The third author of this article, Robert, is TypeCoach's cofounder and president; he taught its program to all 960 students in this year's MBA class.)

The ability to identify how others process information, anticipate where misalignments may occur, and modify one's approach in real time is what we call interpersonal competence. This capability helps teams work through differences before friction turns into dysfunction.

Most team members are willing to collaborate. They start out motivated, capable, and aligned around the objective. What they often don't recognize is how differently people interpret the same situation, how their own behavior is perceived by others, and how quickly those differences can compound when the stakes are high. Without a common language, performance eventually begins to deteriorate. And, as stress increases, people tend to return to their default style, the gap in communication widens, and progress can come to a halt.

Those differences, though, are often beneficial in surprisingly practical ways. Some people focus on concrete realities and immediate practicalities, while others naturally orient toward patterns, implications, and future possibilities. Effective teams need both perspectives. But without tools and training, one side may view the other as being impractical or vague, while the other may feel that its counterpart is imposing unnecessary constraints or resisting innovation. In stressful moments, those misunderstandings often create unnecessary tension, slow decision-making, and weaken trust across teams.

Interpersonal competence isn't a natural instinct or a nebulous talent to "read the room." It is a set of skills that can be learned, practiced, and strengthened over time.

Most people think of communication skills as the ability to factor in

Interpersonal competence isn't a natural instinct or a nebulous talent to "read the room." It is a set of skills that can be learned.

the context of a situation—roles, hierarchy, relationships, culture, and social expectations—and adjust accordingly. But many communication breakdowns occur even when people understand the context perfectly. The deeper issue is cognition: Two people can experience the same interaction in completely different ways. Without a shared framework for recognizing those differences, team members often misinterpret one another's behavior, intentions, and priorities. Interpersonal competence helps teams recognize patterns in real time and adapt before friction undermines execution.

Reducing Team Breakdowns from 45 to One

In 2025 Harvard Business School introduced a four-stage methodology developed in collaboration with TypeCoach to strengthen interpersonal competence across all FIELD Global Immersion teams. The methodology focused on four practical capabilities:

1. Understand one's own communication preferences and work style.
2. Recognize how others process information and engage based on observable behavioral clues.
3. Anticipate potential challenging dynamics before they escalate.
4. Modify one's approach in real time.

Before entering the immersion experience, students completed online tools and training designed to help them recognize the various ways people process information, communicate, and collaborate through observable behavior and adjust effectively in real-world situations involving feedback, decision-making, conflict, and time pressure.

This preparation had a clear payoff. In 2024, 45 of the 158 student teams required faculty intervention to manage conflict and collaborate effectively during the immersion experience. In 2025, after the students received interpersonal competence training, only one team needed assistance.

The difference was adaptation.

The 2025 teams entered the immersion experience with stronger interpersonal competence. They understood more clearly how different teammates were likely to communicate, respond under pressure, and approach decisions. They also had practical tools, including a team playbook, that helped them respond more effectively in real time rather than defaulting to defensive reactions or inaccurate assumptions.

That changed how the teams operated under pressure. Silence was less likely to be misread as disengagement. Direct feedback wasn't always interpreted as personal criticism. Team members tackled tension earlier and in neutral language, calibrated more effectively to one another's working styles, and maintained momentum in situations that previously might have stalled or required faculty intervention.

Most important, students stopped interpreting different approaches as incompetence or withdrawal. They began recognizing them as predictable differences in decision-making, communication, and responses to pressure.

These dynamics appear every day inside organizations. Meetings can feel chaotic or unproductive, decisions can drag, a few voices may dominate while others detach, and friction or misalignment can slow execution. Over time, this often results in burnout, attrition, and leadership fatigue, even with highly intelligent and motivated employees.

Interpersonal Competence in the Age of AI

As artificial intelligence makes information, analysis, and execution increasingly accessible, the competitive advantage is shifting elsewhere: Organizations whose leaders and teams develop the skill of adaptive communication will collaborate faster, execute more effectively, and sustain stronger performance under pressure.

Many organizations do not lack intelligence, expertise, or strategy. They struggle with the growing cost of misalignment—stalled execution, burnout, and talent loss caused by teams failing to communicate clearly, coordinate successfully, and maintain trust as complexity increases. Organizations that develop interpersonal competence reduce friction, solve complex situations more efficiently, and have more-productive teams. ■

HBR Reprint H097RY

LEONARD A. SCHLESINGER *is the Baker Foundation Professor at Harvard Business School, where he serves as chair of its practice-based faculty.* **JOSEPH FULLER** *is the MBA Class of 1960 Professor of Management Practice and a faculty cochair of the Managing the Future of Work project at Harvard Business School.* **ROBERT V. TOOMEY** *is a cofounder and the president of TypeCoach.*

The Office Amenity That Actually Improves Teamwork

by RON FRIEDMAN

WHAT KIND OF workplace helps teams perform at their best?

Organizations have grappled with that question for decades, answering it with one workplace reinvention after another: private offices, cubicles, open floor plans, Silicon Valley–style campuses, and more recently, hybrid-work arrangements.

Despite those efforts, surprisingly little is known about which workplace features truly foster high-performing teams.

To better understand what distinguishes exceptional teams from average ones, my team and I surveyed more than 6,000 knowledge workers across a wide range of industries and identified a small group of exceptionally effective teams. We then explored a variety of factors that could explain their success, including leadership practices, team habits, work arrangements, and workplace amenities.

Participants in our study indicated which of more than 20 workplace

Tanja Ivanova/Getty Images

Teams need individual members to think critically and solve problems. That work requires uninterrupted concentration.

amenities were available where they work, from collaborative workspaces and private meeting rooms to office gyms, Ping-Pong tables, and complimentary food and coffee.

Theoretically, any of those perks could elevate workplace performance by creating opportunities for group bonding, improving workers' well-being, or simply making it easier for people to do their best work.

But in our survey, only one offering consistently distinguished the workplaces of high-performing teams: quiet places where employees could do focused work. Members of superteams were 52% more likely than those of average teams to report having access to a place where they could concentrate without distractions.

What does the availability of quiet workspaces have to do with teamwork? Effective collaboration requires more than constant togetherness. Of course, teams need to discuss ideas, coordinate efforts, and make decisions. But they also need individual members to think critically, solve problems, and execute those decisions. That work requires uninterrupted concentration, and a quiet workspace makes that possible.

The Remote-Work Debate Misses the Point

The desire for more focus is a major reason why remote work is so appealing, especially for those whose offices aren't conducive to sustained concentration. Working from home eliminates many of the distractions that plague office life: background chatter, surprise drop-ins, the feeling of always being "on."

Previous studies on remote work have produced mixed results when it comes to individual productivity, with some showing a boost of 20% or higher and others showing modest declines. Therefore, in our study we decided to focus on team rather than individual productivity: Where are high-performing teams more likely to emerge? Does collaborating in person give teams an advantage? Does working from home lead to better results? Or is the best option a hybrid arrangement, which combines the structure of the office with the flexibility of home?

The answer is actually none of the above. We found that superteams were just as likely to work remotely as they were to have hybrid or fully in-person setups.

That finding seems counterintuitive until you consider that every work arrangement creates its own obstacles to focused work. In-person teams struggle with interruptions, overheard conversations, and the expectation of constant availability. Remote and hybrid teams face a different set of challenges. Research suggests that employees working in hybrid environments increase their use of messaging and video calls, even on days when they are physically in the office. The result can be the worst of both worlds: Employees commute into the office only to spend much of the day in virtual meetings and responding to messages, which leaves little time for focused work.

How Matters More Than *Where*

We found that superteams intentionally carve out opportunities for focused work, regardless of where their members do their jobs. They schedule dedicated focus blocks during which people are not expected to monitor messages or respond immediately to requests, and they establish meeting-free days, creating large stretches of uninterrupted time for meaningful work.

Just as important, they reduce the need for interruption in the first place. They document how routine tasks should be completed and use task-management systems that allow team members to check the status of projects without checking in with their colleagues.

The bottom line: Instead of relying on a steady stream of meetings and messages to stay aligned, superteams build systems that make information simple to find and progress easy to track. *Where* a team works matters less than *how* it works.

Create a Better Environment

No one work arrangement has a monopoly on high performance. Superteams distinguish themselves by protecting focused work, no matter where it happens.

Our findings suggest three ways leaders can create better conditions for teamwork.

1. Treat focused work as an essential part of collaboration. The best collaborations rarely emerge from constant interaction. Teams need space to think independently, refine ideas, solve problems, and execute decisions.

Research on brainstorming supports this point. Groups that brainstorm together often generate fewer and less original ideas than groups in which people

are encouraged to formulate ideas on their own before bringing them to the group for discussion and refinement—a practice known as "brainwriting."

Brainwriting illustrates an important principle: Effective collaboration isn't just about bringing people together. It's also about creating opportunities for people to think independently before contributing ideas to the group. The best leaders recognize that effective teamwork depends on both individual contribution and collective improvement.

2. Protect time for sustained concentration. Superteams are significantly more likely than average teams to establish focus blocks and meeting-free days that allow employees to work without the expectation of immediate responses. Those practices acknowledge an important reality: Every meeting and message imposes a cost. A single interruption may seem insignificant, but dozens of small interruptions throughout the day can fragment attention and make meaningful progress on complex work nearly impossible.

Rather than attempting a sweeping overhaul of how your team works, start small. Designate a 60-minute focus block once or twice a week during which team members are not expected to respond to messages or attend meetings. Then discuss the experience as a team. Did people make headway on important work? Did any problems emerge? Use those insights to refine the approach over time.

The highest-performing teams don't stumble upon better ways of working; they continually experiment, learn, and adapt.

3. Give employees greater control over where they do different kinds of work. Many organizations assume that everyone should work in the same environment throughout the day. But different tasks require different conditions. Writing, analyzing complex information, and strategic thinking often benefit from uninterrupted concentration. Aligning team priorities, gathering input, and refining decisions frequently benefit from discussion with others.

Leaders don't need to redesign entire offices to support this flexibility. They can reserve conference rooms for quiet work during certain hours, establish low-distraction zones, or explicitly encourage employees to choose workspaces that match the demands of the task at hand.

FOR YEARS, ORGANIZATIONS have debated workplace design and return-to-office policies. My team's findings suggest that leaders may benefit from asking a different question: Have we created conditions that allow people to do their best thinking?

The most effective leaders recognize that great teams need more than a lot of interaction. They need workplace conditions that help people contribute their best thinking to the group. ■

HBR Reprint H098TC

RON FRIEDMAN *is a social psychologist, a keynote speaker, and the founder of Superteams Inc. His research focuses on teamwork, collaboration, leadership, and high-performing teams. He is the author of* Superteams: The Science and Secrets of High-Performing Teams *(Simon & Schuster, 2026).*

How Highly Effective Global Teams Collaborate Across Cultures

by ZAK DYCHTWALD

A SENIOR STRATEGY executive at a U.S.-based *Fortune* 100 company recently asked me a question that I hear from global leaders all the time. His teams in Asia were performing well, but very little of his company's strategy originated in Asia, and very little innovation from the Asia-based teams fed back into the company's global strategy. With so much disruptive innovation originating from the region, that was a missed opportunity. Struck by the trends, he asked, "Do the teams in Asia just struggle with speaking up?"

It was an honest question, rooted in the idea that Asia-based teams often come from cultures of indirect communication. But when we compared the company's headquarters with its largest Asia-based team, he and I discovered that the data told a different story: Collaboration was breaking down at the systems level, not the culture level.

Culture was a factor, of course. Speaking up across hierarchy does not come naturally to many people from Asian countries. But culture was neither the whole cause nor the primary

The mistake most leaders make in solving the challenges of global collaboration is to pour their energy into "friction amplifiers."

turner890/Getty Images

cure. Instead, the problem stemmed from several system-related factors, which my colleagues and I at BridgeWorks Global have identified during years of cross-geography research, including surveys of the Global Collaboration Index, hundreds of hours of leadership interviews, and a series of "friction mapping" workshops with 10 global companies that ranged from *Fortune* 10 to *Fortune* 500 firms to Chinese tech giants that were expanding internationally.

Our research on global collaboration has centered on the United States and China—two of the most physically, culturally, and linguistically distant markets in the world—but it extends across India, Mexico, Southeast Asia, East Asia, and Western Europe. When we began our work, we expected culture to be the dominant reason for the breakdown of global collaboration. But it wasn't. Instead, we found that seven elements determine how well global teams work—and leaders have varying degrees of control over them (see the exhibit "The Seven Elements of Cross-Border Collaboration").

In this article I'll show that the mistake most leaders make in solving the challenges of global collaboration is to pour their energy into "friction amplifiers," which they have little or no control over. But they should be focusing on "trust-and-alignment multipliers," which they *can* control and which determine whether the frictions they contain become business problems.

Friction Amplifiers

These are the three elements of global collaboration that you can manage but never really change:

1. **Time zones.** Someone on the team always has to take the 5:00 AM or the midnight call.

2. **Language.** Colleagues who can argue nuance and risk in their first language always have an edge over those doing it in their second, even when their analysis is sharper.

3. **Country culture.** When speaking up across levels isn't the norm, advocacy genuinely suffers. That's more true when everything runs through a global, second-language channel.

All of these elements can be managed with market visits, relocations, and years of language study. "You don't really understand a market until you wake up in it," a *Fortune* 10 executive once told me. True, but 95% of his team will never be able or allowed to wake up there. (At one *Fortune* 500 company, senior leaders outscored managers and others in lower positions by 15 points on the Global Collaboration Index in market knowledge, enthusiasm, and alignment.) Travel fixes things only for the upper crust. The daily load falls on a front line that may never see the market or headquarters at all.

Running global teams without accounting for friction amplifiers is like running a marathon with a rock in your shoe: Everybody does fine for the first mile, most people are limping by

mile 10, and many of your best people won't finish.

Trust-and-Alignment Multipliers

These are the four elements of global collaboration that you can change through intentional design:

1. HQ-region power dynamics. HQ-region power dynamics are in the center between friction amplifiers and trust-and-alignment multipliers. They determine whether a regional leader experiences a global call as a real strategy conversation or a status update and whether "We can't do that in our market" is heard as intelligence or resistance. The clearest sign that things have gone wrong is if, when you ask HQ and regional teams to rate trust and influence separately, they report high levels of trust (they feel respected) but low levels of influence (they feel underpowered). Influence flows to those who advocate best for themselves, not those who read the market best. The loudest voice wins, not the wisest.

2. Company culture. Company culture, unlike country culture, is something you can change, but it rarely travels well. One CEO of a *Fortune* 100 company shared how astonished she was that the phrase "Do the right thing" could be interpreted so differently across her U.S., Indian, Chinese, and Latin American teams. Companies going global stumble here most: Many organizations have never translated their company culture, let alone made it explicit to a workforce that doesn't share their language or assumptions. Company culture tends to live in pithy sayings that carry unwritten rules, and across languages those rules vanish, leaving people to guess at what the company actually expects of them in challenging situations.

3. Market knowledge. Market knowledge matters at HQ, and investing in it pays the highest returns on investment: An HQ that knows the market decides faster, captures more value, and trusts its regions more. Raise this in a leadership workshop anywhere and everybody will nod in agreement. But ask participants what they do to build that knowledge systematically and the room will go quiet. The effort is almost always ad hoc, and leaders heavily rely on the exceptional "bridge people" who are making the extra effort. But when people who have taken two or three years to learn a market are assigned to a different region, what they've learned walks out the door with them.

4. Process. Process is where the drag collects. Most companies run on processes built for HQ, where people who share a building can hash out decision rights: who decides what, which calls should be handled by a global team rather than a local one. During a crisis colocated teams meet in a war room to battle through it, but that's impossible for teams that are 12 time zones apart, with decision rights undefined. Companies must create a global way of working that is distinct from HQ's.

The Global-Collaboration Tax

One regional leader we work with runs a multibillion-dollar business for a U.S. company. On paper she has one job: Win her market. In practice she works two shifts: 9:00 AM to 6:00 PM doing the job she was hired for, and 6:00 PM to midnight doing the one no one sees—explaining her market to an HQ that has rarely visited it, justifying decisions, and translating directives into something her team can follow. She is a perfect example of a bridge person.

Bridge people can be anybody, from frontline operators to senior leaders, but their real job is the hidden work of holding the seams of a global company

The Seven Elements of Cross-Border Collaboration

Most leaders focus on friction amplifiers, which they *can't* control, rather than trust-and-alignment multipliers, which they *can*.

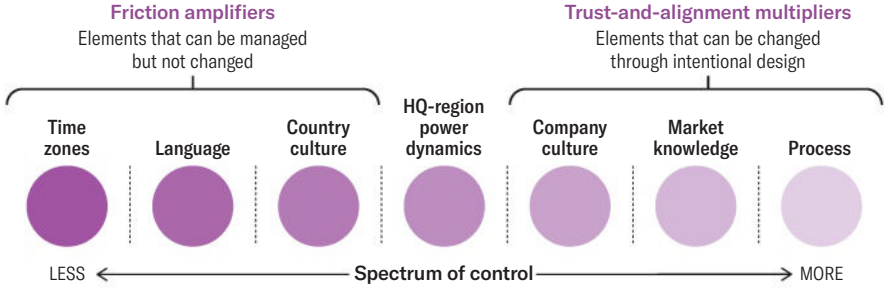

Translate company culture, then make its unwritten rules understandable to employees who don't share your language.

together by translating context between HQ and the market, absorbing the time-zone load, and moving decisions across language and hierarchy. Every global organization relies on them; few acknowledge, accommodate, or reward them. The sum of what they carry is what my colleagues and I call the *global-collaboration tax*.

Bridge people are often treated as employees who must toughen up and fit the system, but the system is what's broken, and bridge people absorb the friction it creates. It is why international work is often seen as a tour of duty rather than a worthwhile destination. The more capable bridge people are, the more they absorb, until the best of them leave. In the companies I study, attrition among bridge people runs roughly twice that of their domestic peers.

Bridge people should not be considered a cost to manage. They are among the most valuable talent a global company has and should be protected, developed, and rewarded, not bent to fit a structure that punishes them for doing its hardest work.

The Way Forward

Companies that are struggling with global collaboration can take a set of actions to improve their situations.

Map leadership team members on the seven elements. Have your leaders rate themselves honestly. The point isn't to get a tidy score; it's to find the gap between where they spend attention and where the leverage sits.

Carefully design decision rights. Don't aim for equality between the regions; aim for intentionality. Define, region by region, what input is advisory and what is determinative, and create an escalation plan before a crisis happens so that no one has to wait for HQ to wake up while the brand burns. That includes codifying innovation or ideation pipelines from high-impact regions rather than relying on ad hoc advocacy. Remember, if it's not in someone's remit, you're winging it.

Systematize market knowledge at HQ. Treat it as a shared asset, not a heroic individual trait. Conduct reverse briefings, in which regions inform HQ of what they're seeing, publish FAQs for recurring questions, and create a regional knowledge hub that captures what would otherwise be lost if an employee left.

Formalize opportunities for strategic input. Innovation disruption is increasingly global. Few HQs create formal opportunities for market-led teams to take part in upstream planning or strategy; they're incentivized on the year, not the decade. Creating strategy summits that are designed to surface, co-create, and incorporate ideas from the market is an excellent way to integrate strong ideas and deepen human ties with regional teams.

Make the company culture explicit. Translate it, then make its unwritten rules understandable to employees who don't share your language or context. You'll lose the poetry of the succinct phrases but gain a team that knows what the company stands for when faced with a difficult situation.

Reduce work and codify how employees accomplish it. This is the fastest win we see: Cut global meetings by 20%–25% while holding HQ-regional alignment steady. Schedule global meetings only two or three days a week; nobody wants meetings every night. Devise and distribute templates for asynchronous written communication so that the company has one standard, not 15. That may seem trivial, but a client recently estimated that following a single standard will save it millions of dollars annually in wasted time, rework, and turnover and attrition.

Protect and reward bridge people. Put the work of bridging into performance reviews, promotion criteria, and how you talk about talent. At a minimum, accommodate it (the person on an 11:00 PM call doesn't owe you a 9:00 AM start the following workday), but even better, reward it. That way you don't just keep your best global operators; you tell everyone doing this work that the second shift is seen, valued, and central to how the company wins.

GLOBAL WORK WILL always be harder than domestic work. But leaders who treat it as a system to be designed, not a mystery to be endured, will build organizations that decide faster, spend less to maintain stability, and succeed in the markets where competitors are moving faster. Country culture shapes how cross-geography collaboration feels. Operating design determines whether it wins. ■

HBR Reprint H098S2

ZAK DYCHTWALD *is the founder and CEO of BridgeWorks Global, a cross-geo collaboration lab, and the author of* Young China: How the Restless Generation Will Change Their Country and the World *(St. Martin's Press, 2018).*

Evidence. Insight. Action.

Get expert-backed intelligence—drawn from the HBR Advisory Council and global business leaders—and explore hundreds of white papers and research reports from HBR Analytic Services.

RECENTLY PUBLISHED

Building Strong Brands
Sponsored by Frontify

Driving Business Value with Agentic AI
Sponsored by Wipro

Leveraging Alternative Payment Models to Ensure Broad Access to Personalized Medicine
Sponsored by BU Questrom

Procurement at a Crossroads: Evolve or Be Left Behind
Sponsored by Zip

Accelerating Growth Through Revenue Enablement
Sponsored by Seismic

HARVARD BUSINESS REVIEW

View all the reports at
hbr.org/hbr-analytic-services

Register for the HBR Advisory Council at hbr.org/advisory-council.

November–December 2026

Features

"The more company executives and customers agreed on what drove customer value, the higher the company's revenue, gross margins, free cash flow, and market value were."

"DO YOU KNOW WHAT YOUR CUSTOMERS VALUE?" PAGE 68

Photograph by **GABRIELE GALIMBERTI**

Why Even Your Best Employees Override *AI Systems* That Work

AI & MACHINE LEARNING

It's not that they don't trust the technology. It's that they don't trust your organization.

AUTHOR **Das Narayandas** // *Professor, Harvard Business School*

PHOTOGRAPHER **JAMIE CHUNG**

AI & MACHINE LEARNING

In late 2022 a global bank approved an AI system designed to automate routine credit decisions. On paper it looked exactly like the kind of forward-thinking move experts had been urging companies to make. In pilots the system processed applications faster than humans did, reduced errors, and promised meaningful cost savings. The risk management team signed off on it, compliance approved it, and the CEO touted it publicly as proof that the bank was leaning into the future.

Six months later the system was still scoring applications and generating recommendations. Yet most files continued to flow through a human review. Exceptions had multiplied, and temporary safeguards like dual review thresholds and manual confirmation flags had hardened into permanent process steps. While the AI system was part of the bank's workflow, it never actually made most final decisions.

The explanations sounded familiar: The organization was not quite ready, the model needed refinement, employees needed more training, change is hard. But the technology worked, the economics were clear, and the system had been approved formally and enthusiastically. What it had never really been granted was authority—the right to make decisions within agreed-upon boundaries with the organization's full backing.

Versions of this story are playing out across many industries. A global manufacturer rolls out AI-based demand forecasting, yet planners still use their age-old playbooks and treat the AI's forecasts as a second opinion. A healthcare network introduces AI-assisted triage that reduces wait times and improves consistency, and even though clinicians agree the model is accurate, they continue to manually review a majority of the recommendations. The system is used more for documentation than decision-making.

The problem isn't failed pilots. It's that organizations' structures for accountability and legitimacy haven't evolved to support the use of AI systems. By legitimacy structures, I mean the shared understanding of who is entitled not only to make a decision but also to have that decision be trusted. Too many institutions are accumulating intelligence without reorganizing around it.

We Have Seen This Before

This pattern feels new because the technology is new, but structurally it is not. Earlier waves of decision automation—for example, enterprise resource planning (ERP) systems and credit scoring—exposed the same institutional habit: Organizations adopt systems that can make better decisions but preserve the human authority structures that prevent those systems' decisions from being accepted.

When ERP systems (such as those from Oracle and SAP) were introduced, in the 1990s, they promised centralized visibility and optimization across the enterprise. Nestlé USA's experience implementing ERP is instructive. The company's

IDEA IN BRIEF

THE PROBLEM

Many organizations deploy AI systems successfully but never allow them to make meaningful decisions. Employees keep overriding the AI because they are still accountable for outcomes, even when decision authority is supposedly delegated to the systems.

THE CAUSE

Organizations redesign technology without redesigning governance. Employees are expected to trust AI but remain personally exposed when its decisions produce bad outcomes. Under those conditions, overriding the system is an act of rational self-protection, not resistance to change.

THE SOLUTION

Treat AI adoption as an organizational redesign challenge, not simply a technology rollout. Before scaling up AI systems, explicitly reallocate decision rights, clarify who owns outcomes, redesign human roles around oversight and exceptions, and align accountability with delegated authority.

Companies must address the gap between authority and accountability. If they don't, people will continue to override AI systems.

divisions had operated autonomously, with each accountable for its own production, cost, and supplier contracts. When the ERP system required them to surrender that discretion to a common process, the employees who remained personally accountable for local outcomes had good reason to keep making the decisions themselves. The system produced a single integrated forecast; the accountability structure rewarded local overrides. As Tom Davenport of Babson College has noted, the company treated the system as a technology installation rather than a general-management redesign.

At around the same time, automated credit scoring was following the same pattern. Models could codify risk faster and more reliably than humans, but override cultures abounded. Studies of auto lending show that algorithmic underwriting produced higher loan profits and lower default rates than human underwriting, and that the biggest performance gap was where loan officers had the strongest incentives to preserve their own discretion. Other research found that loan officers manipulated inputs and ratings to push borderline applications over approval thresholds to earn volume incentives. Even though authority had been delegated to the models on paper, accountability remained with the people.

The airline industry's experience with revenue management systems provides an instructive contrast. These systems achieved lasting authority because the airlines supported them by reassigning accountability. The carriers clarified when human intervention was appropriate, structured the escalation process, changed performance evaluations, and made adhering to the system professionally defensible. Authority shifted not because the algorithms became perfect but because the airlines changed the governance surrounding them.

The patterns in those examples is consistent. When decision logic migrates faster than the structures surrounding accountability and legitimacy can evolve, authority gets frozen in place. When governance changes, authority can shift. Davenport, Erik Brynjolfsson, and others have noted this pattern before. What this article identifies is the precise reason for it: Authority fails to move when people remain personally liable for the system's outcomes.

Companies must address this gap between authority and accountability. If they don't, people will continue to override AI systems—not as an act of resistance but as one of self-protection.

No amount of training or communication can overcome this gap. An organization can manage change well and still allow the technology to settle into a permanent advisory role. It can cultivate psychological safety and still punish the individual who followed the model when the decision goes badly. When accountability remains personal, people will take steps to protect themselves.

Other explanations for why AI systems haven't been given decision-making authority point to the wrong remedies. For example, some have posited that structural inertia causes resistance, but with AI we often see apparent compliance: Dashboards are running, workflows include the system, and employees use it. But still the system isn't given real decision rights. Research on automation complacency warns that humans defer too readily to machines, but that's not the issue here, where they *don't* defer, because accountability structures make doing so professionally dangerous. The productivity paradox literature says that IT investment fails to show up in performance gains, but it mostly fails to specify the organizational mechanism that explains why. Research by Ethan Bernstein on transparency shows that workers who are observed conceal their practices to avoid scrutiny, but they use the visible apparatus of compliance, documented overrides, escalation steps, and review layers as vehicles for preserving control.

What's Really Going On

In most of these situations, an AI system is used and gets touted in board presentations, yet when the system and a human disagree, the human almost always wins quietly. Professionals reshape the system's boundaries for reasons of prudence—they establish more review steps, set slightly lower thresholds, create more documented overrides, and issue new flags for sensitive cases. These adjustments are individually defensible. Collectively, they are ruinous.

UPS's deployment of its ORION routing system illustrates this tension. When rollout began, in 2013, many drivers treated its computer-generated routes as suggestions, favoring local knowledge and customer relationships. Over time UPS changed its scorecard: Drivers were measured against

AI & MACHINE LEARNING

ORION's prescribed route rather than their own past performance, and their adherence to it was tracked. Following the algorithm thus became the professionally defensible act.

I saw this dynamic recently inside a different global logistics firm, one that handles high-value shipments for corporate clients. The company approved an AI routing system after pilot projects had significantly reduced late shipments and improved asset utilization. Dispatchers admitted privately that the system often surfaced better options than they would have considered on their own. Yet after its rollout, the same dispatchers continued to adjust routes throughout the day for customers they had long-standing relationships with, for shipments that had political importance inside the firm, and for more-trusted drivers. At first these adjustments were described as fine-tuning, but over time the AI system ceased being a decision engine; its outputs were merely the starting points for conversation.

When the firm's leadership pressed for fuller adoption, the dispatchers didn't argue that the system was wrong, nor did they refuse to use it. They added more structure: flags for sensitive shipments, additional escalation steps, and lower review thresholds. A dispatcher explained the logic to me in a disarmingly honest way. He said, "If the system makes the call and something goes wrong, it is my name on the report. If I make the call, at least I can explain it." The dispatcher wasn't confused, nor did he feel psychologically unsafe. He was reading the accountability structure accurately.

Research on professional identity threats shows that workers whose expertise is being automated can feel genuine anxiety about their status, relevance, and career trajectories, and that this anxiety can shape their behavior in ways that are neither irrational nor trivial. But that doesn't explain the pattern I observe, which emerges even where fear and identity threats have been addressed, where change management has been done well, where psychological safety is high, and where employees openly acknowledge that the system outperforms them. What remains is structural exposure: The person asked to follow the system is the one whose career absorbs the cost when it is wrong.

The severity of the issue is even more profound with agentic AI systems. Recommender systems at least preserve the fiction of human control: The system proposes, the human decides, and accountability has a face (even when it should have moved). Agentic systems dissolve that fiction. They do not recommend; they act. They book, approve, route, contract, and escalate within parameters that someone authorized but no one continuously monitors. When an agent acts within its defined parameters and the outcome is bad, the accountability question is genuinely difficult: Was the parameter set wrong? Was the authorization too broad? Did the agent encounter a situation its designers failed to anticipate? These questions will be asked in boardrooms, courtrooms, and regulatory filings. When an AI system recommends, the accountability gap produces workarounds. But when an agentic system acts, it can produce liability ambiguity on a large scale.

Shifting Accountability

Real adoption is unmistakable: Decision processes simplify. Exceptions decline rather than multiply. Authority is reallocated explicitly so that people know who (or what) decides, who explains, and who absorbs risk. That clarity can cause discomfort, as it forces role changes. But it is ultimately stabilizing because accountability becomes linked with authority.

Amazon's Hands off the Wheel initiative is an example of how authority gets transferred. Begun in 2012 as Project Yoda, it embedded forecasting, pricing, and purchasing predictions in workflow tools. The organization separated input authority from output authority. Business teams could shape the data and parameters, but routine overrides were discouraged and sometimes required approval. The responsibilities of vendor managers shifted from individual buying and pricing decisions to auditing inputs, exceptions, and outcomes.

The transfer stuck because accountability moved with authority. When results missed expectations, the question was not which person made the wrong call but which input, rule, or system design failed and who was responsible for making the correction. Meanwhile, human work was redesigned around product and program management tasks that were less amenable to automation.

Technology and organizational architecture have to move together, and the organizational work has to be designed before a system is scaled up, not retrofitted after resistance appears. Leaders must ask four questions pertaining to governance. They might cause some discomfort, but they

Professionals reshape the AI system's boundaries for reasons of prudence—creating more review steps, setting slightly lower thresholds. Collectively, these adjustments are ruinous.

Styling by Brian Byrne

AI & MACHINE LEARNING

are essential for surfacing accountability gaps. They are: (1) Who loses discretion if this AI system works? (2) Who gains authority and over which decisions? (3) Who explains outcomes when the system is followed? (4) Who absorbs transition risk during the learning curve?

If those four questions are not answered explicitly by the organization when the AI system is approved, they will be answered implicitly afterward by individual operators in favor of preserving their existing authority.

Full adoption will occur only when approving a system's decisions does not feel like professional self-sacrifice. Leaders must redesign roles so that when routine decisions are automated, human judgment is shifted upward into broad oversight, boundary setting, and system refinement. Organizations should establish clear guardrails and give employees tools to explain decisions they did not personally make. Moreover, the system needs visible sponsorship by leaders who are willing to absorb early errors as institutional learning instead of blaming individuals for failures.

The Capability That Will Differentiate Organizations

As judgment becomes codifiable and scalable, firms will need to reorganize themselves around their AI systems. Unfortunately, too many firms simply acquire intelligent systems and expect the rest to fall into place. They treat each AI deployment as a stand-alone project that rises and falls on model performance, vendor selection, and rollout discipline. In those firms every new system triggers the same cycle of enthusiasm, containment, partial adoption, and quiet stabilization into advisory mode. Over time employees learn that nothing fundamental will change, and leaders learn to temper ambition before it threatens role structures.

Savvy leaders, by contrast, recognize that what is at stake is not automation but authority. They see where decision rights actually reside, redesign workflows around machine defaults rather than layering systems on top of legacy processes, and govern transitions in ways that don't push the learning costs onto individual employees.

Sequence matters here. Authority will migrate only after an organization makes the algorithm the operational default, makes overrides visible, redesigns human roles around auditing rather than second-guessing, and moves accountability to system owners. Technology and organizational architecture have to move together.

Organizations that have done this once are better positioned to do it again. In my consulting work and field research, I have seen how firms manage the next change correctly: They bring forward a shared vocabulary, a set of questions, and a memory of what had to be redesigned the last time. They recall who lost discretion, who or what gained decision rights, who explained outcomes, and who absorbed early errors. That is not to say that a firm's next transition will be painless, nor will it automatically be less political. But the firm will start the journey from a different place.

That is why companies should approach authority mobility as an organizational capability, not a onetime implementation challenge. Before they deploy a model, they need to resolve where decision rights will move, what role human judgment will play afterward, and how the organization will absorb the learning costs of trusting the system.

If this idea resonates with you, you are probably not standing at the approval gate. You are somewhere past it, looking at a system that has quietly settled into the advisory role this article describes. The prescriptions for the next deployment—treat approval as role design, answer the four questions before scaling, write the absorption commitment before the rollout—still matter. But existing deployments require a different next step: Diagnose where authority actually sits right now.

Most organizations prefer to leave one thing unsaid. If AI works the way it is supposed to, the company's need for certain forms of human judgment is going to shrink. Language about augmentation cannot soften that reality. If an underwriting model can automatically approve 70% of routine cases, then the underwriter's role cannot continue to be defined by individual case approvals. Silence by executives on this point essentially guarantees quiet containment.

To be clear, organizational redesign around AI systems is not about eliminating human judgment; it's about moving it where it matters: assessing portfolio risk quality, monitoring model drift, and handling the cases the system was not built for. That shift has to be codified in performance evaluations. If people are rewarded for outcomes the system increasingly shapes, they will protect themselves by making defensive overrides.

Authority and accountability have to come into alignment. If a model is allowed to decide within defined parameters, following those parameters must be professionally defensible. Leaders can't blame individuals for early mistakes by the system.

The aviation industry offers a useful example. The industry embraces "just culture" principles, which encourage people to surface safety issues while preserving a boundary

If a model is allowed to decide within defined parameters, following those parameters must be professionally defensible. Leaders can't blame individuals for the system's mistakes.

between acceptable mistakes and unacceptable conduct. EU law requires aviation organizations to explain how those boundaries are guaranteed and enforced, and report information is protected from use for blame or liability except in cases of willful misconduct or gross negligence. As a result, errors become visible, allowing companies to analyze and correct them and prevent them from recurring.

Toyota's Andon system offers a similar lesson. A worker who sees a defect can call attention to the problem, bring a team leader to the line, and, if the issue cannot be contained, stop production. Note that identifying a problem triggers help, coaching, and root-cause problem-solving, not automatic blame. Attempts to copy Andon have failed at companies that reproduced the tool without reproducing the accountability norms around it. The AI equivalent is straightforward: Before rollout, leaders need to state which early errors will be treated as system-learning events, which will require escalation, and which will remain the responsibility of individuals.

Most organizations track adoption through usage dashboards. They rarely track the behavioral signatures of quiet containment: whether exceptions are rising or falling, whether override thresholds have shifted, whether the escalation process has added steps. During the ORION rollout, UPS's critical governance move was not tracking whether drivers used the system but whether they followed it. That shift changed what was managed.

The pressure to get this right will come from outside the organization as well. The AFL-CIO's Workers First Initiative on AI, launched in October 2025, called for human review of automated decisions, transparency in how AI systems are used, and appeal rights for workers subject to AI-driven employment decisions. Organizations that have not formally aligned authority with accountability will increasingly be asked to do so by their employees, by lawmakers, or both.

The future of employees' roles should be clearly addressed as the system scales up, not after employees have inferred that its success means their own decline. IKEA and the Ingka Group, IKEA's largest franchisee, offer an example. As Ingka Group's AI chatbot Billie resolved simple inquiries, the company reskilled call-center workers in remote interior design, digital retail sales, relationship building, and complex customer problem-solving. Ingka made the new work visible while the AI was scaled up, which helped employees see the transition as growth rather than managed decline.

Salesforce offers a parallel example under its 4Rs framework: redesign, reskill, redeploy, and rebalance. As Agentforce, Salesforce's platform for building and deploying autonomous AI agents in the enterprise, absorbed more routine support work, Salesforce helped employees move into new roles. The initiative appears to be working: According to the company, 51% of hires in its first quarter of fiscal 2026 were internal. The point here is that the company named the role transitions, built an internal-mobility mechanism around them, and made the redesigned work easier for employees to understand.

When these elements come together, professional authority does not disappear. It concentrates. The underwriter no longer spends her day reviewing routine cases the model has already scored. She works on cases outside the model's confidence boundary, portfolio patterns that suggest drift, and edge cases that inform future parameters. The dispatcher no longer watches routes generate and quietly adjusts them. He monitors the exception queue, manages client relationships the model cannot capture, and escalates situations the system was not designed for. Both know what they are responsible for. Both can explain a decision that went wrong without beginning, "I overrode the system because...."

The Measure of Success

At some point soon you will approve an AI system that works. You will be tempted to measure success by usage rates, dashboard metrics, and the smoothness of the rollout. But those signals are misleading.

Ask instead: If the system were turned off tomorrow, would the decisions that matter in your organization actually change? Or would people continue making decisions as they always have, only without the interface open on their screens?

If decisions would not change, then authority has not moved. And if authority has not moved, then adoption has not occurred, no matter how advanced the model or how disciplined the implementation. ■

HBR Reprint R2606A

DAS NARAYANDAS *is the Edsel Bryant Ford Professor of Business Administration at Harvard Business School.*

Design Your Supply Chain as a Portfolio of Options

OPERATIONS & SUPPLY CHAIN MANAGEMENT

AUTHORS

Luiz Mesquita
Associate professor, W.P. Carey School of Business

Roberto Vassolo
Professor, IAE Business School

Jean-Paul Destarac
Doctoral student, W.P. Carey School of Business

Thomas Y. Choi
Professor, W.P. Carey School of Business

You need to be able to defer investments, quit unfavorable positions, expand in promising regions, and shift operations as policies evolve.

ILLUSTRATOR **PETER GREENWOOD**

WAVE

after wave of disruptions have roiled global supply chains. In addition to navigating the tariff wars triggered by the United States, companies have had to contend with export controls on critical inputs such as rare-earth metals, advanced semiconductors, and pharmaceutical precursors; shifts in geopolitical alignments as the Global South has deepened its ties with China; and conflicts in the Middle East and Ukraine, where trade sanctions have been used as weapons. Yet the supply chain planning tools most companies use were built for a world where this level of volatility was the exception. Now it is the rule.

For more than a decade—and especially since the pandemic—the guiding principle for supply chain management has been to design supply chains for resilience: Build buffers, diversify suppliers, engage in collaborative planning with members of the network, and so on. That makes sense when shocks are relatively rare. But it is inadequate in a world where the rules themselves—trade agreements, tariff regimes, regulatory standards—keep changing. In such an environment, companies need to design their supply chains as portfolios of real options that provide them with the right—but not the obligation—to take future actions.

Companies have long applied this approach to strategy formulation and capital budgeting. But it has not been widely employed in supply chain management. A few multinationals—particularly those with extensive experience managing cross-border operations—have intuitively adopted elements of this approach, but most companies still evaluate supply chain investments using traditional discounted cash-flow methods that treat flexibility as a cost, not an asset. Real options give companies the ability to defer an investment in new capacity until the regulatory picture clears, to decrease operations in jurisdictions that become unfavorable, to expand production into new jurisdictions, and to switch sourcing or assembly from one location to another.

In this article we present a framework for designing a supply chain that way. It is based on our combined research and consulting work with global manufacturers navigating trade policy shifts, as well as others' work on the theory of real options as applied to strategic investments. Our framework rests on three interconnected ideas. First, the major sources of uncertainty are not just volatile markets and unpredictable events; they include volatile institutions and their rules. Second, a company should view its supply chain as a portfolio of interrelated options whose value depends on how those options are structurally coordinated. And third, the constraint on exercising the options is not just capital but also capabilities: the local production expertise, qualified vendor networks, and regulatory knowledge that take years to build.

THE SUPPLY CHAIN OPTIONS FRAMEWORK

In a world of stable trade rules, the right supply-chain strategy is straightforward: Commit to the lowest-cost configuration and optimize. But when the rules themselves are volatile, the goal is not to find the optimal configuration but to maintain the ability to reconfigure—and that requires a portfolio of options. Our framework can help. (See the exhibit, "The Supply Chain Options Framework.") It organizes four options according to their primary purpose: protecting the company when conditions worsen or positioning it to capture value when conditions improve.

The objective for leaders is to maintain the right to act as conditions shift: committing when clarity is sufficient and preserving flexibility when it is not. Sometimes this requires actions that are defensive, and sometimes it requires actions that are offensive. Having options to make both types of moves is essential.

The cost of maintaining the portfolio is the option premium. And the payoff is the ability to move first when the rules change again. The guiding discipline is to ensure that when a policy shift makes one jurisdiction more attractive, at least one other option in the portfolio benefits from the

Most companies still evaluate supply chain investments using discounted cash-flow methods that treat flexibility as a cost, not an asset.

opposite movement, so that regardless of which way the rules break, the firm has a position that improves—or in options lingo, gains in value. A company never exercises its full portfolio at once. Rather, it operates in a steady state of selective commitment and maintains optionality—exercising specific options as uncertainty resolves while keeping others open.

Now let's take a closer look at the four options, starting with the defensive moves.

The option to defer. The most counterintuitive of the four options is also the most valuable when volatility is high: the deliberate decision not to fully commit.

Consider the global auto industry's supply chain for electric vehicle (EV) batteries. Over the past five years, GM, Ford, and Volkswagen responded to surging demand projections and strong government incentives with massive commitments to secure battery capacity: multibillion-dollar investments locked into specific jurisdictions, long-term offtake agreements (contracts where a buyer agrees in advance to purchase a certain amount of a product from a seller, often over a long period and sometimes at pre-agreed prices), and joint ventures to build and run factories that produce specific types of batteries. The strategic direction was sound. After all, the EV transition was real, and securing supply was a priority. What was mispriced was the cost of irreversibility: These commitments were structured as now-or-never decisions in an environment where the rules were still being written or were not stable.

Then the rules changed. The U.S. Inflation Reduction Act imposed new sourcing requirements for EV tax credits, mandating that by 2025, at least 60% of critical minerals and battery components be sourced from the United States or free-trade-agreement partners—thresholds set to escalate annually to a maximum of 80% in 2027 for critical minerals and 100% in 2029 for battery components. The European Union introduced regulations requiring manufacturers to digitally document the full supply-chain history of every battery—from mineral extraction through assembly—with traceability standards that differed from those in the United States. China tightened export controls on critical battery minerals. Each policy shift occurred after automakers had already committed billions of dollars. The companies that had made the largest, earliest commitments were the most exposed.

It's not that automakers should have avoided the EV battery space altogether. But holding off on going all in until enough clarity emerged to commit on better terms would have been a better approach. A company that approved the acquisition of a site for a battery plant and the initial engineering of it, while making the next tranche of capital contingent on regulatory clarity, would have entered 2024 with options rather than obligations. It would have retained the ability to choose among battery chemistries, jurisdictions, and partners once the decision parameters were clearer. The cost of waiting was real but modest, such as the delay in capacity ramp-up. The cost of not waiting was vastly higher. GM recorded about $6 billion in charges in early 2026 tied to its EV pullback. Ford recorded about $19.5 billion in charges in late 2025, which included $8.5 billion in asset write-downs.

The right to wait has a positive value, and that value *increases* with the degree of volatility. Traditional net-present-value analysis systematically ignores this value by framing every investment as a now-or-never decision. Leaders need the discipline to practice restraint with purpose: structuring their position so that when they do commit, they

IDEA IN BRIEF

THE PROBLEM

Most supply chains are still designed for a relatively stable world. But today's environment is shaped by volatile tariffs, export controls, and geopolitical realignments, and shifting regulatory regimes. Traditional supply-chain-planning tools cannot adapt fast enough when the rules keep changing.

THE SOLUTION

Create a portfolio of real options to preserve your ability to defer investments, abandon unfavorable positions, expand quickly in promising regions, and shift operations across jurisdictions as conditions evolve. That means staging investments, designing reversible contracts, maintaining footholds in multiple markets, and standardizing operations to enable rapid reconfiguration.

THE PAYOFF

A real-options approach allows companies to treat flexibility as a strategic asset rather than a cost. Firms that invest early in local capabilities, supplier networks, and regulatory expertise can respond quickly to trade disruptions and policy shifts—turning uncertainty into a competitive advantage.

do so with better information and to a more favorable configuration—at a time when the rules are less likely to reverse.

→ ***How to create options to defer.*** Stage capital commitments with explicit decision gates: Approve Phase 1 with a contractual option to proceed to Phase 2 contingent on regulatory clarity. This makes each phase a separate option exercise, not an automatic escalation. And rewrite your capital-budgeting criteria to value deferral: the higher the policy uncertainty, the higher the threshold for irreversible commitment.

The option to abandon. Every supply chain commitment has a hidden cost that never appears on the balance sheet: the cost of unwinding it if conditions change. A 15-year facility lease in a jurisdiction that suddenly faces 25% tariffs is not just an operating expense; it is a trap. The option to abandon means structuring commitments so that the cost of walking away remains low relative to the cost of staying in a deteriorating position.

The critical design principle is *reversibility.* Most firms default toward the irreversible end of the spectrum because it is cheaper in the short run. Firms should shift to the reversible end when the option value of being able to exit exceeds the gains from staying locked in.

Zara's supply chain illustrates this logic. While most global apparel companies consolidate production in low-cost Asian factories with long lead times and deep contractual commitments, Zara does something different. It produces trend-driven items in areas that are close to its European consumer base—in Spain, Portugal, Morocco, and Turkey—using short production runs, small batch sizes, and supplier contracts designed for rapid ramp-up and ramp-down. The conventional interpretation of Zara's "fast fashion" approach is that it is a speed strategy—one that enables it to respond to trends more quickly. That is true but incomplete. Zara's entire supply chain is structured as a portfolio of options to abandon nodes in its supply chain. If any one of them becomes unfavorable, the cost of exiting it is low. In a world of stable trade rules, this structure appears wasteful. In a world where tariffs, sourcing rules, or trade restrictions can shift overnight, it is a competitive advantage.

→ ***How to create options to abandon.*** One way is to lease warehouse and distribution facilities rather than own them—and keep the terms short. A three-year lease with renewal options converts a fixed exposure into a series of decision points. In contracts with suppliers, include either exit clauses or volume bands (predefined ranges of purchase quantities that allow you to scale orders up or down within the band without renegotiating the contract), and have short renegotiation cycles.

Now let's turn to the offensive moves.

The option to expand. Many companies view market entry as a binary decision: You're in or you're out. But between full commitment and total absence lies an array of opportunities that real-options thinking makes visible. Companies need to identify the minimal operational presence that preserves the right to scale up later.

The Supply Chain Options Framework

Companies facing volatile trade rules need supply chains designed as portfolios of options. Our framework identifies four options that give organizations the flexibility to protect value when conditions worsen and capture it when conditions improve.

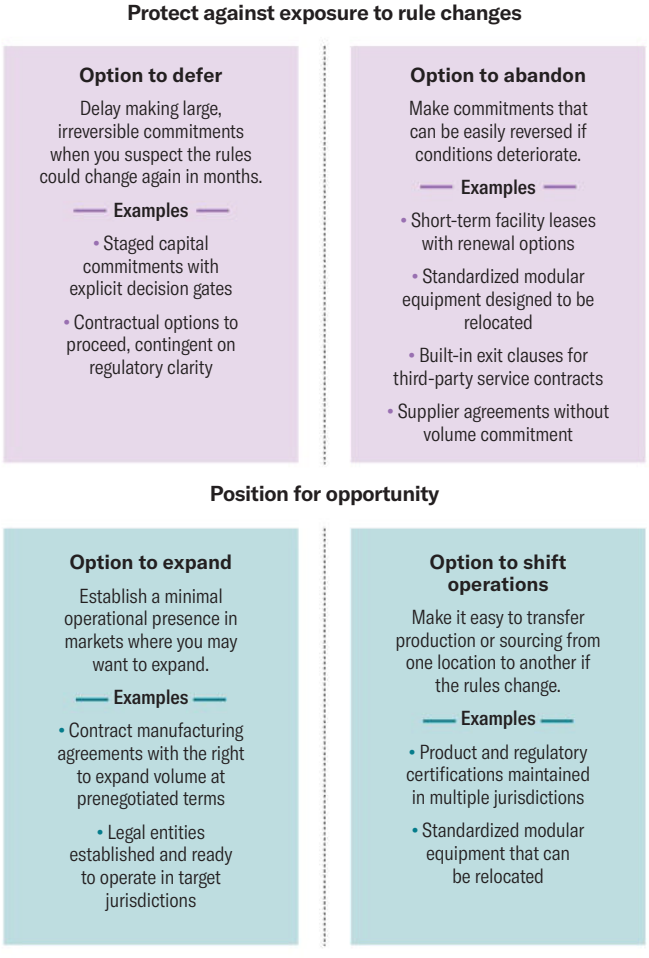

Every supply chain commitment has a hidden cost that never appears on the balance sheet: the cost of unwinding it if conditions change.

For example, a firm could pay a small premium to establish a foothold and structure the deal so that it could be expanded if conditions warrant. Apple did this when it came under political pressure to bring its manufacturing back to the United States. Rather than making a single big U.S. investment, however, Apple built a portfolio of investments across the globe—some initially small, others substantial—aimed at maintaining its options in its supply markets. It is now sourcing semiconductors from the Taiwan Semiconductor Manufacturing Company's plant in Phoenix and is investing in Amkor Technology's new advanced-chip packaging and test facility in Peoria, Arizona. It has formed several other partnerships to increase its sourcing in the United States, including one with TDK for sensors and another with Bosch for sensing hardware circuits. At the same time, Apple is sourcing assembly work in India and Vietnam with companies such as Foxconn, Pegatron, and Tata Electronics. Some of these positions—particularly in India—have already been expanded into major operations. Others remain early-stage footholds. Apple is also sourcing high-value components from these countries, pivoting away from China under what is known as the "China plus one" approach.

→ ***How to create options to expand.*** You can establish a legal entity in the target jurisdiction—one that is registered, tax-compliant, and authorized to conduct business—so that when conditions warrant scaling up, For Free Mags Check sastatus.com you can begin operations without the months-long process of incorpora- tion and regulatory approval. Another approach is to sign a contract-manufacturing agreement with the right to increase production capacity or volume under defined terms. Or you could obtain the documentation needed to qualify products for preferential treatment under a trade agreement. The value of these approaches lies in what they allow your firm to do tomorrow—quickly and at a fraction of the cost of entering cold. A small contract-manufacturing relationship in Mexico or Eastern Europe may contribute little today. But the ability to rapidly ramp up production for your company if tariffs or sourcing rules change is extremely valuable.

The option to shift operations. This option allows a company to quickly move production, sourcing, or assembly across jurisdictions when conditions change. It does not mean duplicating full-scale operations in every jurisdiction. It means duplicating the *right to use* factories.

One European industrial manufacturer we have worked with, which we will call Euromanu, has built its supply chain this way. It maintains a network of standardized, smaller assembly plants on each continent. Under traditional capital budgeting, having a single consolidated plant in the lowest-cost jurisdiction is more economically efficient. But the firm's leaders are not optimizing for cost; they are optimizing for flexibility.

As a result, when U.S. tariffs on imports from China rose sharply in 2018 and 2019 and again in 2025, Euromanu shifted assembly volume from China to its Americas facility within weeks. In 2026, when the European Union began imposing fees on imports on the basis of their carbon emissions, threatening to raise the cost of goods shipped from high-emission manufacturing regions, Euromanu's European plant absorbed additional volume. Competitors that relied on a single manufacturing location affected by those actions spent months scrambling to find alternative suppliers, qualify new production sites, and navigate unfamiliar regulatory environments—only to find that by the time they had adapted, a different set of rules had emerged. Euromanu's flexibility let it capture volume that fixed-footprint competitors couldn't reach, turning a policy shock that hurt everyone else into an advantage only Euromanu could use.

→ ***How to create options to shift operations.*** Use standardized, modular equipment and processes that can be easily added, removed, or rearranged across sites, recognizing that some investment duplication is necessary. That's what has allowed Euromanu to move equipment between continents in a matter of weeks. Invest continuously in local capabilities before the volume justifies it. This entails prequalifying suppliers and training production teams.

WHAT MAKES OPTIONS EXERCISABLE

Acquiring the local production expertise needed to manufacture consistently to specification takes time. The knowledge of how materials behave under local conditions, which machine settings work with locally sourced inputs, and how to troubleshoot quality deviations that don't show up in statistical process control resides in experienced local teams and does not travel with the equipment. Qualifying a new supplier typically takes six to 18 months of audits, sample testing,

and trial production runs. A firm that has already qualified vendors in a jurisdiction can redirect sourcing immediately; a firm that has not may face a long lead time. And then there is the regulatory know-how: Each jurisdiction has its own web of customs procedures, labor regulations, and informal norms that can be navigated efficiently only if you have practical experience and relationships with local authorities. Such capabilities take years to build and cannot be improvised in a crisis. They are what determine whether a supply chain option is genuinely exercisable or exists only on paper. That is why Euromanu continually invests in local capabilities in every jurisdiction where it operates—even when the current volume in a jurisdiction does not justify the expense.

The amount of analysis required to take a real-options approach to designing your supply chain might seem overwhelming. But artificial intelligence is beginning to make the endeavor more tractable, though AI's capabilities are still emerging. AI systems are being built to continuously score and rank supply chain configurations against shifting tariff regimes, dynamically reroute production and sourcing across global networks, and stress-test the entire portfolio of options against multiple policy scenarios simultaneously.

MillerKnoll, a furniture and design company, is already doing this: It uses Kinaxis's Maestro platform to model the impact of tariffs and trade compliance policies across its supply chain, evaluating sourcing options and anticipating risks before rule changes affect margins or delivery commitments. In effect, AI has the potential to serve as the control system that helps managers identify which combinations of options are most valuable under different regulatory futures. That said, while algorithms can reveal attractive alternatives in seconds, they cannot instantly create qualified suppliers, train local teams, or obtain regulatory approvals.

THREE ACTIONS LEADERS MUST TAKE

The volatility in trade rules requires you to make three fundamental changes to the way you design your supply chain.

Reframe how you evaluate supply chain investments. Stop treating flexibility as overhead and start treating it as a critical asset. Every investment proposal should be evaluated not only on its expected return under current conditions but also on the real options it creates or destroys under alternative scenarios. If your capital-budgeting process can't distinguish between irreversible commitments and staged investments that preserve the right to adapt, it needs to be revamped.

Audit your commitment profile. Map every major supply-chain commitment along two dimensions: how reversible it is, and how exposed it is to a single set of rules—meaning, would it become a liability if the current tariff regime, trade agreement, or regulatory framework changed? Most firms will find that their profile is heavily skewed toward the irreversible, single-regime end—not because anyone made a deliberate strategic choice but because that is where traditional procurement optimization naturally leads.

Invest in the capabilities that take years to build—before you need them. Again, local production expertise, qualified vendor networks, regulatory know-how, and certifications are the capabilities that determine whether your supply chain options are genuinely exercisable when the next policy shock hits. So make sure you have developed those capabilities well in advance of the time you need them.

None of these actions will be easy to defend internally. Finance will see them as overhead. Procurement will view them as inefficiency. Operating managers will resist creating redundant capacity in jurisdictions where the current volume doesn't justify the investment. On a quarterly earnings call, analysts and investors might criticize your actions, arguing that maintaining production readiness in a market where you are not yet operating at scale appears wasteful. The answer isn't to cave to such pressure; it is to make the case why such actions are critical.

THE GOAL OF the real-options approach is not to design the optimal supply chain to serve your company's current needs. It is to build the capacity to reconfigure your network when the environment changes again and again. The winners in this new world will be those who figure out how to invest in uncertainty rather than hide from it—who treat volatility as a condition to be managed. The question is, Does your organization have the discipline to pay for flexibility before it needs it and the courage to defend that investment when internal and external forces push you to cut it? ■

HBR Reprint R2606B

LUIZ MESQUITA *is an associate professor at Arizona State University's W.P. Carey School of Business and a member of the editorial board of the* Strategic Management Journal. **ROBERTO VASSOLO** *is a professor at IAE Business School, Universidad Austral. He is a coauthor of* Advanced Strategic Management: A Dynamic Approach to Competition *(Palgrave Macmillan, 2024).* **JEAN-PAUL DESTARAC** *is a doctoral student in supply chain management at W.P. Carey School of Business and the CEO of Capin-Vyborny, a U.S. customs brokerage.* **THOMAS Y. CHOI** *is a professor at W.P. Carey School of Business. He is a codirector of the Complex Adaptive Supply Networks Research Accelerator (CASN-RA).*

LEADERSHIP

How CEOs

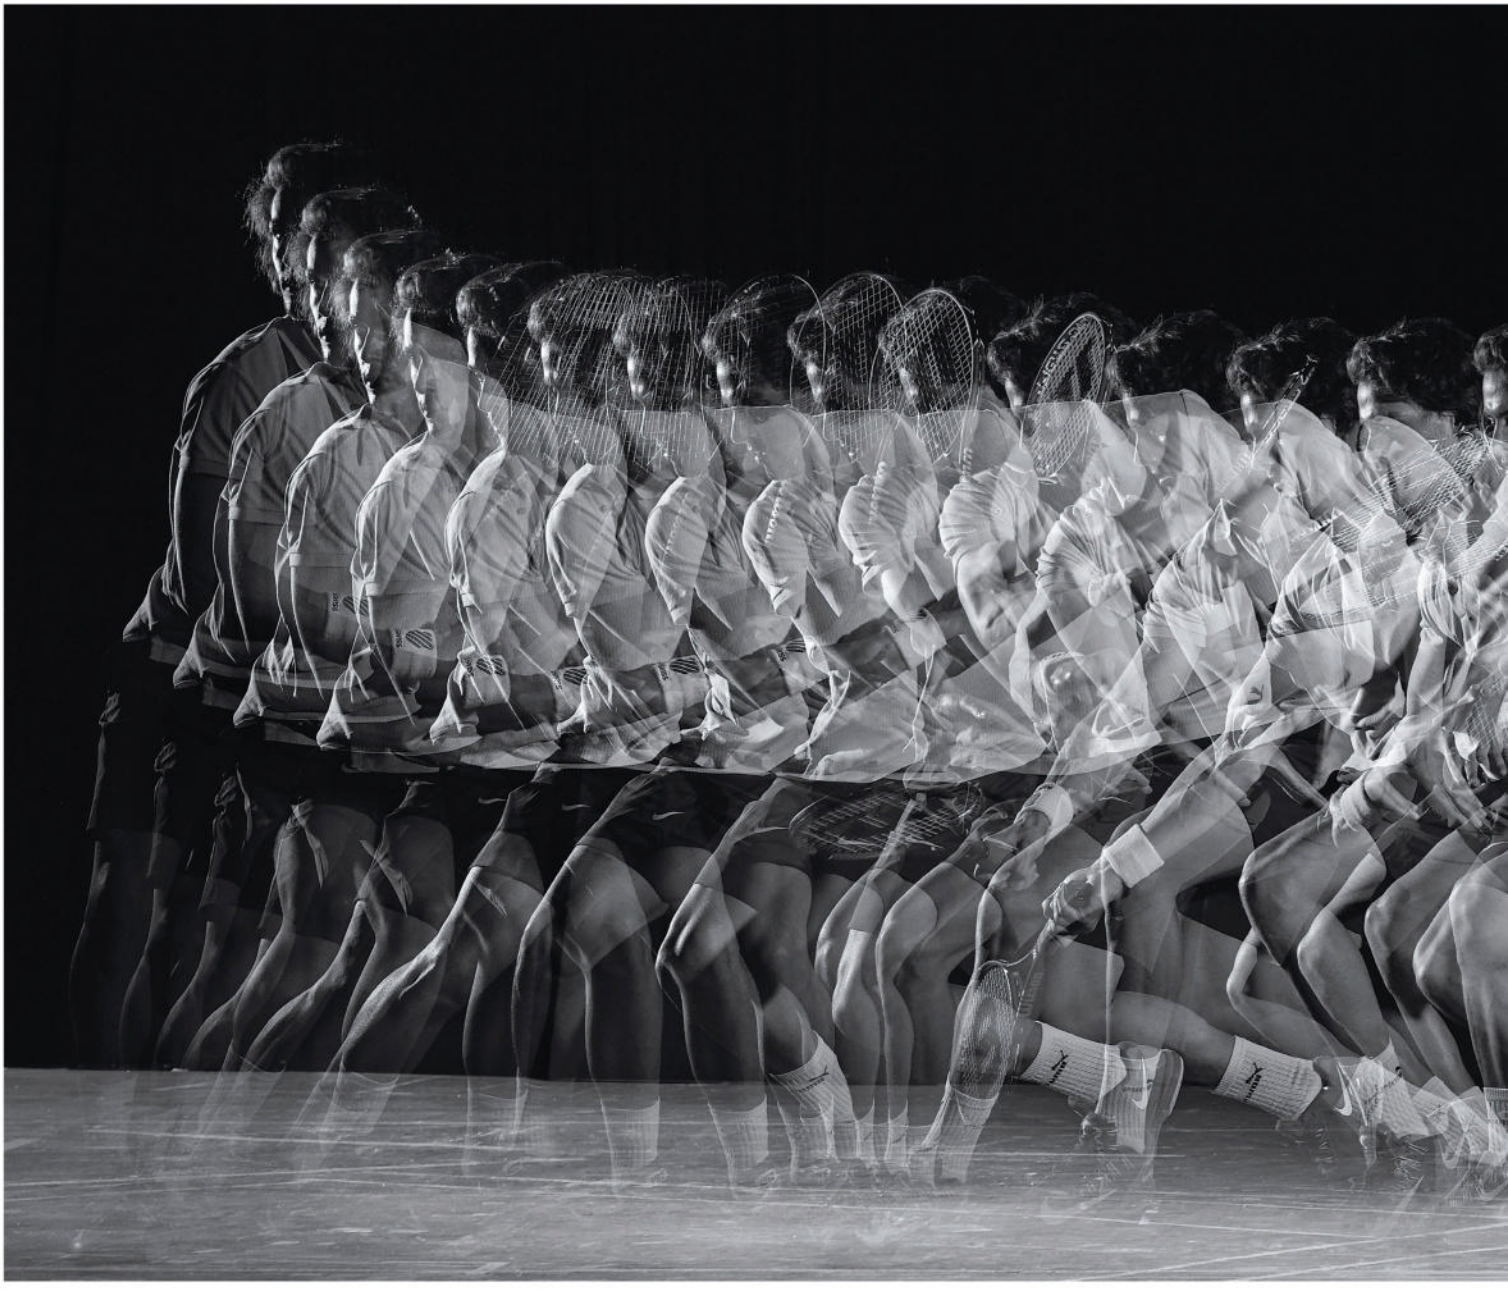

One of the most underappreciated aspects of the role

Set the *Tempo*

is deciding how fast to move and when.

AUTHOR
Nitin Nohria
Professor, Harvard Business School

PHOTOGRAPHER **JEAN-YVES LEMOIGNE**

IDEA IN BRIEF

THE PROBLEM

Many CEOs choose the right strategy but go astray by paying too little attention to tempo—the pace at which decisions are made, priorities are pursued, and change unfolds.

THE INSIGHT

Different situations call for different tempos, from rapid experimentation to steady patience, and effective CEOs continually adjust their organization's rhythm as circumstances change.

THE TAKEAWAY

Treat tempo as a strategic leadership tool. Match it to the organization's capabilities, the competitive environment, and the task at hand, while recognizing when it's time to shift gears.

ABOUT THE ART

Photographer Jean-Yves Lemoigne captures athletes in motion for his Chronophoto *series, using a stroboscopic technique to achieve a sequence effect.*

When Nvidia became the most valuable company in the world in 2024, many assumed its CEO must be operating at a breakneck pace that matched the frenetic speed of the AI revolution its chips were powering. The reality was almost the opposite.

Jensen Huang had spent decades building a distinctive internal cadence characterized by deep immersion in product reviews, rigorous group discussions, and an unrelenting emphasis on long-term architecture over short-term reaction. Throughout Nvidia's evolution, he continued moving at this deliberate pace. That discipline turned out to be its own form of competitive advantage.

Every organization's tempo—the pace at which its employees make decisions, launch products, tackle problems, and pursue change—is set, intentionally or otherwise, by its CEO.

Tempo is a consequential lever for CEOs because knowing what to do is a necessary starting point—but knowing when and how fast to do it is the key to successful execution. Strategy determines a CEO's direction. Tempo determines whether the CEO arrives at the destination on time. Often CEOs are replaced not because they chose the wrong strategy but because they operated at the wrong tempo.

CEOs can enlist others in the organization to help set strategies, priorities, and goals, but only they establish an organization's tempo. People watch how quickly CEOs act on their stated priorities, make decisions, and respond to requests for their time, input, and guidance, and how much patience or urgency they exhibit and expect from others.

The CEO's preferred tempo, whether rapid-fire or measured, communicates expectations, sets the drumbeat for other constituencies, and determines how swiftly attention and resources are allocated to activities and projects.

Each tempo has its time and place. A company rarely moves at a single, consistent tempo over a CEO's tenure. Most CEOs vary their tempo, speeding up or slowing down in accordance with their company's context, current needs, and planned changes.

CEOs must also attune their tempo to their organization's response. When the tempo is too fast for too long, employees may feel harried, disoriented, and burned out. A pace that delivers results in the short term can exhaust an organization if it's overextended. High intensity cannot be sustained indefinitely. On the other hand, when the tempo is too slow for too long, the resulting lethargy can hinder the company's progress and lessen its market relevance. When the world is moving quickly, a company can't afford to dawdle. The CEO must act like an orchestra conductor, keeping time and modulating the tempo for the moment at hand and the organization's capacity.

The Five Kinds of Tempos

Drawing on my observations of several hundred CEOs across industries and eras, I've developed a typology of tempos. A memorable adage captures each one and reflects its distinct operating philosophy. Each tempo has its own logic—and its own failure mode.

These tempos are not one-size-fits-all; they're context-dependent strategies, and choosing among them is a leader's

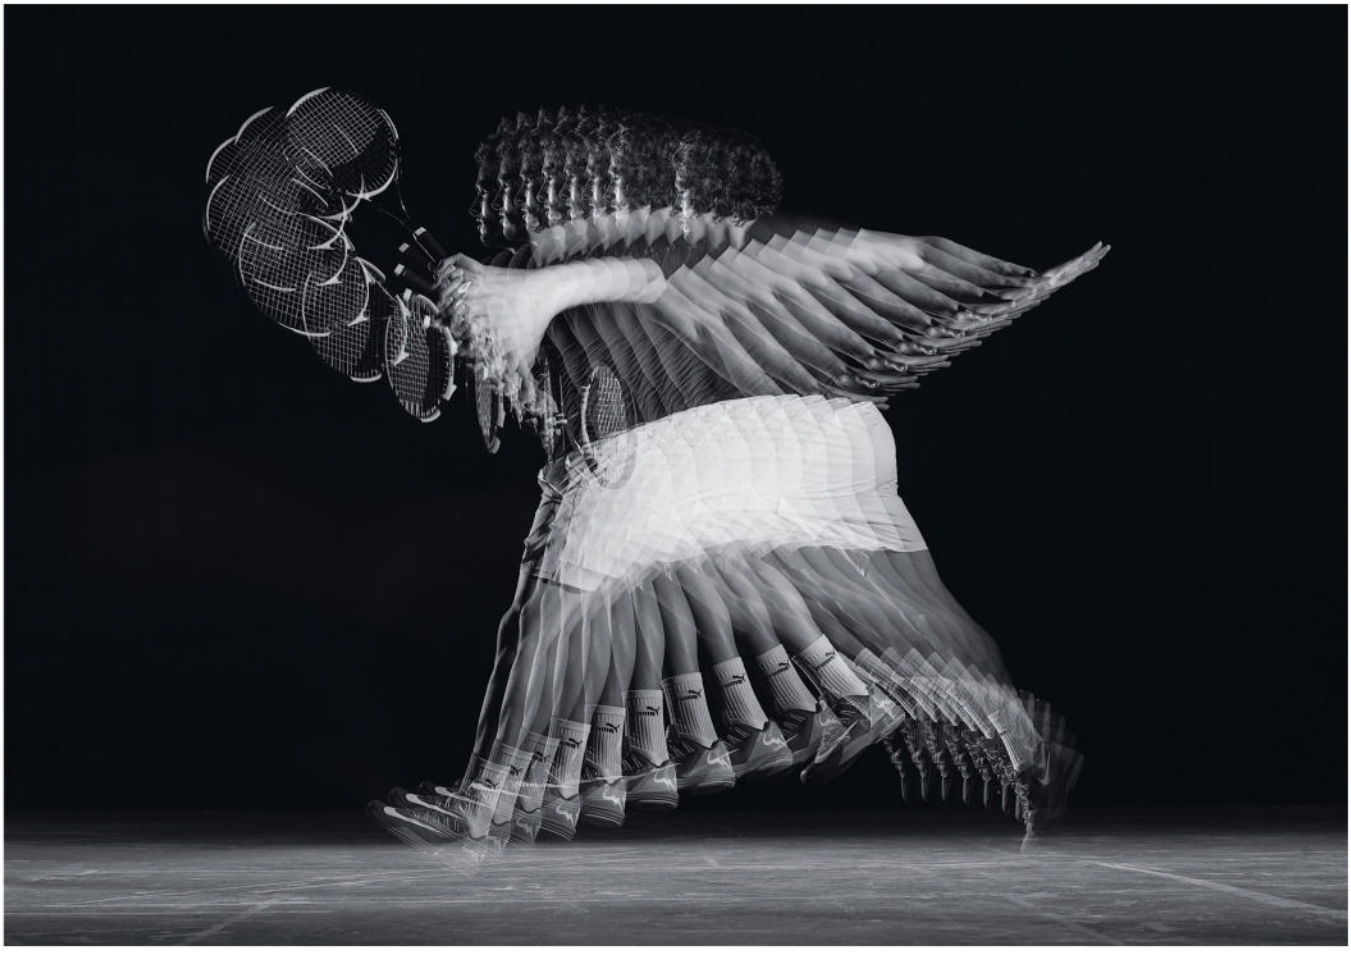

critical task. Think of them as a set of tools, each of which fits a different set of circumstances.

Experimental speed. "Move fast and break things," the motto famously associated with Mark Zuckerberg's early leadership at Facebook, aptly sums up this tempo. CEOs operating at this pace prize speed and disruption above all. To them, mistakes are necessary by-products of acting and learning quickly. They urge their teams to press ahead without restraint and believe that falling behind can be fatal. This tempo optimizes companies' ability to make discoveries, enabling rapid testing, iteration, and refinement. CEOs who adopt it encourage short action cycles, embrace impermanence, and view failure as feedback.

This tempo is suited to situations where customer preferences evolve quickly, first-mover advantage is real, network effects can create compounding benefits, and the cost of inaction is higher than that of imperfection. When misapplied, however, this tempo can create churn without progress.

Relentless pressure. "Stretch; reach to the point where it almost comes unglued" was a catchphrase of Jack Welch during his tenure as GE's CEO from 1981 to 2001. He described his ideal tempo as one in which the organization always felt slightly on edge, as though it were on the verge of losing control. The result was a productive discomfort, a tenuously controlled speed still lively enough to teeter and thrill. If experimental speed is about increasing learning, relentless pressure is about increasing performance. The path is largely known; the imperative is execution.

CEOs who adopt this tempo set high expectations, challenge complacency, and keep teams on their toes by quickly moving on from people who can't keep up or projects that don't deliver. This tempo works best in cultures that value unremitting improvement. It can produce extraordinary results—but only if organizations can sustain it. Such companies become known for their speed and intensity, to the point that people largely self-select during the firms' hiring process, often drawn by the higher compensation they may be offered to endure the pace. When overdone, however, this tempo can cause burnout, short-termism, and an overly narrow focus on what can be delivered quickly rather than what matters most.

Disciplined urgency. "Be quick, but don't hurry," the legendary UCLA basketball coach John Wooden's famous mantra, has been embraced by leaders like GE's Larry Culp. It suggests a different speed: insistently fast yet bridled. The goal is to act quickly and deliberately. With this tempo, CEOs try for a balance between swiftness and precision, operating at a pace that's brisk but never frantic. Disciplined urgency is less a burst of speed than a steady operating philosophy. It sustains momentum without sacrificing coherence.

This pace is particularly well suited for leaders managing turnarounds, where time is of the essence but recklessness is costly. It requires clear prioritization, tight sequencing, and a culture where focus is as valuable as speed. Its failure mode is subtle: In striving for balance, organizations may drift toward caution, mistaking hesitation for discipline.

Front-loaded deliberation. The phrase "Slow, slow, quick, quick" is used by ballroom instructors to teach the foxtrot, a dance that originated in the early 20th century. I've worked with CEOs who use it to describe their preferred pace: They start slowly, taking the time to understand, align, and prepare; then, once they've gathered information, they execute quickly and decisively. Front-loaded deliberation is a sequencing strategy: It's slow at the outset to enable speed later.

This style suits leaders operating in high-stakes, high-complexity situations where early missteps are expensive and reversals are difficult. CEOs drawn to this tempo value depth of understanding and quality of preparation, and they believe that their painstaking analysis will pay off.

Political leaders and CEOs in highly regulated environments often gravitate toward this tempo, where early consensus building and stakeholder alignment are prerequisites for decisive action. When taken too far, however, this approach can lead to analysis-paralysis, where preparation delays the very action it's meant to enable.

Patient compounding. The adage "Slow and steady wins the race," popularized by Aesop's fable of the tortoise and the hare, captures this tempo. It's often associated with the leadership of Warren Buffett. His signature patient pace defined his decades at Berkshire Hathaway as he focused not on delivering quarterly results but on building enduring value.

This tempo emphasizes consistency, restraint, and the power of steady progress. It can be seen in the 250-year vision that Konosuke Matsushita, the founder of Panasonic, once described for his company. CEOs who adopt it don't rush decisions. They avoid fads and build cultures that prioritize sustainability over spectacle. It's especially well suited to organizations that benefit from the long-term accrual of assets, whether of capital, brand equity, or trust.

The risk is that forbearance can devolve into complacency, particularly when external conditions shift more quickly than the organization anticipates. The tortoise wins when the race is long; it loses when the course suddenly changes.

The Contingencies That Shape the Right Tempo

Several critical elements inform the tempo that a company should adopt. Let's begin with what's most fixed, the leader's own temperament, and move outward toward what's most variable, the external environment in which the company competes.

The leader's temperament. Every CEO operates with a natural rhythm, whether that means sleeping on decisions or responding to every request within the hour. When leaders operate within their authentic tempo, they tend to be more composed and consistent. When they adopt a pace that feels uncomfortable, they sometimes can lose coherence.

What shapes this temperament? Professional background can play a role. Leaders who come from functional backgrounds tend toward front-loaded deliberation; their instinct is to build a case before acting. Those with more operational backgrounds often gravitate toward disciplined urgency, where speed and precision are both important. Founders, engineers, and scientists, who are comfortable with rapid cycles of testing and refinement, tend to favor experimental speed.

Formative experiences matter equally. CEOs who have led a company through a near-collapse tend to have a heightened sensitivity to downside risk, calibrating pace with one eye always on what could go wrong. Those who have ridden a growth wave early in their careers are likely to favor a fast tempo, having internalized speed as the natural state of a healthy organization.

Temperament evolves with the maturity of the leader and the organization, however. Zuckerberg's pace at Meta has slowed considerably since his "move fast and break things" era, reflecting his company's growth into a global institution and his own development as a leader navigating regulatory scrutiny, geopolitical complexity, and long-cycle bets on the metaverse and AI. The instinct for speed didn't disappear; it was tempered and redirected.

CEOs must learn how to operate effectively at a pace that doesn't come naturally. The solution lies in understanding the distinction between style and range. A deliberate CEO can learn to speed up without abandoning the habits that make deliberation valuable. A fast-moving CEO can learn to slow down selectively without losing urgency. The goal is not to become someone else but to expand what you're capable of—to stretch, modulate, and adapt your instincts to the moment. The best leaders bring the organization along at a pace that reflects the business's demands and their own authentic range. Think of Jamie Dimon, who has been the CEO of JPMorgan Chase for two decades. He leveraged his natural decisiveness to rapidly tighten his firm's risk discipline throughout the financial crisis and then selectively accelerated growth afterward. Compare his approach with that of Stanley O'Neal, the CEO of Merrill Lynch from 2002 to 2007, whose failure to respond with urgency to his firm's mortgage exposure proved disastrous once conditions deteriorated.

The company's life stage. Just as individuals' energy levels change throughout their lives, so, too, do organizations'. The tempo that fits a scrappy startup rarely suits a mature enterprise.

In the startup phase, speed is often synonymous with survival. It can help ventures capture market share, secure funding, and leverage early user feedback to iterate. Here, the "move fast and break things" tempo is often a necessity. Many early-stage CEOs operate at high velocity, making decisions with incomplete data and relying on instinct and proximity to the product.

Yet there are notable counterpoints. Some founders deliberately slow key decisions early on, particularly around hiring, architecture, or culture, to avoid locking in weak systems that are costly to unwind later. For example, Jeff Bezos famously enforced a "raise the bar" mindset at Amazon and instituted unusually rigorous hiring standards, passing on "good enough" candidates for months to ensure that the team's composition didn't cost the company speed and quality later. In such cases, short-term restraint leads to speed over the long haul.

In a company's growth stage, the tempo must shift as the organization adds complexity in the form of new geographies, functions, or customer segments. A more structured plan becomes necessary, and leaders must strike a

Just as individuals' energy levels change throughout their lives, so, too, do organizations'. The tempo that fits a scrappy startup rarely suits a mature enterprise.

balance between ambition and execution. A pace somewhere between "be quick, but don't hurry" and "slow, slow, quick, quick" may make more sense. The CEO's challenge is to maintain urgency without inviting chaos.

In mature businesses, the tempo often slows—sometimes to a fault. At this stage the best leaders resist complacency while recognizing that scale and systems require a more measured approach. As strategic horizons lengthen, capital allocation and talent development become long-cycle bets. The goal is sustainable advantage, not speed for its own sake.

When mature organizations move too slowly and must be roused from inertia, they become turnaround stories. When Tufan Erginbilgic became the CEO of Rolls-Royce, in 2023, he immediately created a sense of urgency by telling employees they were on a burning platform, challenging legacy systems, shortening decision cycles, and refocusing on a few critical priorities. His leadership set a pace more typical of a startup than of a legacy industrial firm.

The nature of the business challenge. This is the most obvious determinant of tempo. The tempo for a business confronting disruption differs from that of one executing a well-positioned strategy, just as it will differ between a period of decline and one of steady growth.

CEOs must pick up the pace when new technologies or shifting consumer behaviors upend their companies; otherwise, they risk losing ground quickly. Leaders may set a high-speed tempo to encourage rapid prototyping, reinforce quick decision-making, and establish risk tolerance. The early days of internet platforms, smartphones, and, more recently, gen AI demanded this urgency, as companies' established business models quickly became outdated.

Crisis also has a powerful impact on the metronome. When the Covid-19 pandemic rattled the world, CEOs had to act swiftly to reorganize supply chains, transition teams to remote work, and update stakeholders about decisions made under uncertainty. In such moments, indecision can be more dangerous than haste.

Conversely, the pace may need to slow down during periods that require intense coordination or when leaders seek to change deeply rooted systems. When the goal is evolutionary change, such as a creating a culture of safety or overhauling incentive systems, CEOs must allow themselves time to align stakeholders and build durable commitment.

Misreading the moment is costly. Leaders who impose speed in contexts that require deliberation can create turmoil without progress; those who move too slowly in periods of disruption may find that time is the one resource they cannot recover. Consider Elizabeth Holmes at Theranos, where the pressure to move fast and announce products before they worked and scale them up before the science was ready created a gap between claimed progress and reality that no amount of subsequent effort could close. By contrast, BlackBerry's leadership was too slow to recognize that the iPhone had permanently redrawn the smartphone landscape. The company dismissed touchscreens as impractical and clung to the physical keyboard as a competitive advantage—and as a result ceded an entire market.

Internal readiness. Deciding how fast to go is one thing; knowing how fast the organization can go is another. Leaders must match their tempo to the organization's ability to act, considering its culture and its execution capacity.

In companies that thrive on urgency, tight deadlines and rapid cycles can energize employees. In methodical or consensus-oriented cultures, the same tactics may create resistance. A CEO who imposes the wrong pace—too fast for a risk-averse culture or too slow for an entrepreneurial one—can quickly fall out of sync. Think, for example, of Bob Nardelli's stint leading the Home Depot from 2000 to 2007. His methodical, process-driven changes clashed with a previously entrepreneurial culture. Compare him with Satya Nadella of Microsoft, whose growth mindset reignited the company's pace of innovation.

Execution capacity is equally important. An ambitious CEO may want to drive fast change but may overreach if the systems, people, or resources aren't ready. In such cases, employees may disengage or exit. Wise CEOs gauge the organization's actual capabilities, not merely its aspirations, before setting the pace. Many transformations have failed because the tempo of change outstripped the organization's capacity to absorb it. Consider Ron Johnson's short tenure as CEO at JCPenney from 2011 to 2013. His rapid overhaul of pricing and store formats overwhelmed systems and frontline execution. In contrast, Alan Mulally, Ford's CEO from 2006 to 2014, paced his turnaround with disciplined weekly reviews and a few clear priorities, building capacity before accelerating change.

Leaders who set the right pace pay attention to the organization's rhythms. They watch for signs of strain and note when initiatives stall.

External pressures. A final factor is the environment outside the firm. Markets, competitors, regulators, and public expectations can all influence a CEO's chosen tempo.

In fast-moving markets, customers expect frequent innovation, and competitive intensity creates constant pressure. A company that releases new features annually will struggle to compete with rivals iterating monthly. When competitors quickly cut prices, launch products, or capture attention, a CEO must decide whether to match that pace or differentiate by moving more deliberately.

Regulatory environments exert a different kind of pressure. Moving fast can be inappropriate in sectors like healthcare, energy, and financial services. Here, caution, consultation, and compliance shape the tempo. CEOs must build coalitions, engage with policymakers, and pay attention to legitimacy.

In some industries the tempo is also shaped by natural cycles, such as product launches, regulatory calendars, and seasonal demand, which create predictable periods of acceleration and pause. At a pharmaceutical company like Johnson & Johnson, product development cycles can be a decade long, whereas a fast-fashion company like Zara or Shein has product windows that can be as short as weeks or days.

The Art of Pacing

Pacing isn't an instinct CEOs are born with. It's a skill developed through experience, feedback, and deliberate attention to the organization's signals. Like many aspects of leadership, it's conveyed through small, visible choices that reveal what matters.

CEOs must learn to sense when the organization is exhausting itself or drifting into complacency, recognize when their demands energize or overwhelm, and shift gears without disorienting their teams. As Erginbilgic said to his leadership team a few years into his transformation of Rolls-Royce, when the company's performance was on a much stronger footing, "The art now is finding the right pace—knowing when to push and when to prod."

Reading the organization's pulse. Leaders who set the right pace pay attention to the organization's rhythms. They watch for signs of strain and note when initiatives stall. They observe whether teams are reactive or proactive and whether momentum builds or dissipates.

This pulse taking is not mechanical. It requires presence and the ability to interpret weak signals. CEOs who excel at pacing often rely on informal channels as much as dashboards and spend time walking the halls, asking questions, and noticing tone, hesitation, and energy levels.

At GE, Culp's direct engagement in lean workshops (structured problem-solving sessions held on the factory floor with frontline employees) gave him confidence that operational improvements were taking hold, allowing him to maintain a measured pace and resist investor pressure to hastily divest the company's struggling energy business. Other CEOs track how long decisions linger unresolved or how frequently teams escalate issues upward, which signals that their tempo may be misaligned with their organization's capability.

Managing the transitions between tempos. When a CEO changes tempo, it's not as simple as issuing a directive. Organizations resist sudden shifts, and teams need the time and the rationale to reorient. Leaders must explain how the pace is changing, why, and what it means for day-to-day work.

Accelerating the tempo requires clarity and focus. CEOs must anchor the need for speed in a concrete business imperative: a market threat, a competitive move, or a strategic opportunity. As they do so, they often narrow priorities, shorten decision cycles, and increase the frequency of operating reviews.

Slowing the tempo can be just as difficult. Teams accustomed to speed may interpret it as a loss of ambition. Here, the CEO must reframe deceleration as a purposeful opportunity to consolidate, reflect, or improve quality—not a retreat but a recalibration. Leaders often signal this shift by extending decision timelines, encouraging deeper analysis, or revisiting assumptions that had previously gone unchallenged.

Sustaining a tempo. Some CEOs establish a consistent tempo that becomes part of the firm's identity. As Huang has demonstrated at Nvidia, a steady operating cadence, maintained even as the pace of innovation accelerates around it, can allow a company to move quickly without losing coherence. Routines, such as how often teams meet, how quickly decisions are expected, and how frequently performance is reviewed, help define the rhythm of the organization.

Coherence is also strengthened when leaders establish a consistent cadence of innovation and cultural renewal. Employees come to expect, and trust, a certain pace of

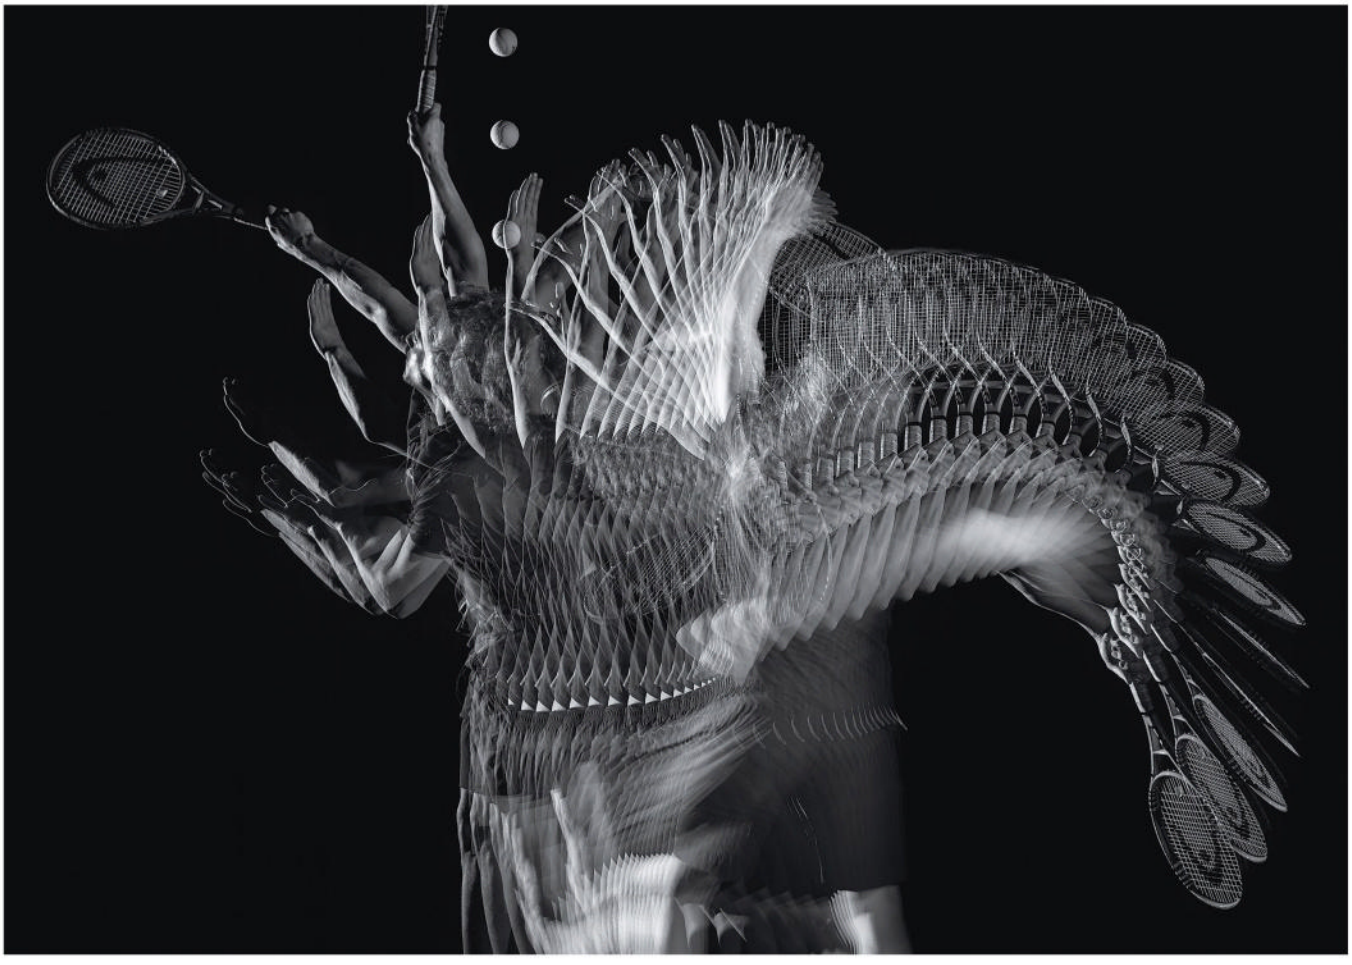

progress. The advantage of this is that when a shift is required, it becomes more visible and therefore more meaningful.

The Intangible Leadership Skill of Setting the Tempo

The concept of tempo is not typically taught in business school and rarely appears on formal leadership competency models. Yet it distinguishes leaders who reach the end they have in mind from those who merely imagine it. Tempo sits at the intersection of strategy and execution. It is where intent meets action.

Setting the tempo is ultimately an exercise in judgment. CEOs must develop an instinct for when to accelerate and when to pause, when to press the organization and when to let it breathe. This judgment is built not only on analysis but also on pattern recognition and self-awareness. Over time, experienced leaders develop a temporal intuition, an ability to sense when the organization is ready to move and when it needs time to absorb change.

CEOs must model the tempo they wish to see. If they want urgency, they must move purposefully, make decisions swiftly, and hold others accountable. If they want a more measured pace, they must demonstrate a spirit of inquiry, patience, and deliberation. But the best CEOs don't assume they know the right pace in every circumstance. They listen, learn, and recalibrate continually. Every CEO ultimately confronts the same questions: not just what should we do, but how fast should we do it, and what are we prepared to risk by moving faster or slower?

In leadership, as in music, timing is not everything—but without it, everything else falls flat. ■ **HBR Reprint** R2606C

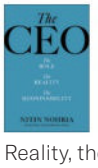

NITIN NOHRIA *is the George F. Baker Jr. Professor of Business Administration at Harvard Business School and a Distinguished Service University Professor at Harvard. From 2010 to 2020 he served as the business school's 10th dean. He is the author of* The CEO: The Role, the Reality, the Responsibility *(Harvard Business Review Press, 2026).*

CUSTOMER STRATEGY

Do You *Know* What Your Customers Value?

Not according to the data. Here's how to fix that.

AUTHORS

Vikas Mittal
Professor, Jones Graduate School of Business

Alessandro Piazza
Associate professor, Jones Graduate School of Business

Sonam Singh
Assistant professor, Muma College of Business

Ashwin Malshe
Associate professor, Alvarez College of Business

PHOTOGRAPHER **GABRIELE GALIMBERTI**

CUSTOMER STRATEGY

ABOUT THE ART

Gabriele Galimberti's project Toy Stories *examines childhood through portraits of children from more than 50 countries with their favorite possessions.*

Most companies work hard to get close to customers.

Managers also regularly compare how their product offerings stack up against the competition, so it's not surprising that executives are confident about their knowledge of customer needs and whether they're satisfying those needs better than their rivals do.

Take the CEO of one $320 million software services company who told us, "Our sales team meets with key accounts monthly and responds immediately to their needs. The COO monitors competitors, and we're at the forefront of almost every industry association. Last year we started more than 50 initiatives to stay ahead of the game." Another case in point is the chief operating officer of a $12 billion oil-field-services company who declared, "I have more than 23 years of industry experience as a client and vendor, in the field and in purchasing. With all due humility, I know where customers and competitors stand."

But is such confidence justified? To find out, we undertook a seven-year study of 72 companies in a variety of industries (including retail, food services, manufacturing, distribution, industrial services, software, education consulting, and fundraising), whose annual revenue ranged from $10 million to $20 billion. We interviewed the CEOs of the companies and 187 of their senior executives, observed more than 3,000 hours of strategy meetings (including 83 strategy retreats), surveyed 24,000 customers, and examined the financial statements and regulatory filings of an additional set of public companies.

What we saw was that executives did indeed spend a great deal of time listening to customers, monitoring competitors, and looking to spot industry trends. But we found that rather than using a scientific approach to interpret the information they gathered, they routinely took cognitive shortcuts. As new factors surfaced and different executives focused on different ones, companies' strategies for winning customers shifted from year to year, confusing customers about the underlying value propositions the firms were offering. As customers switched to competitors, sales dwindled. Meanwhile, the costs of pivoting caused margins to shrink, customers didn't get what they wanted, employees were burdened by the need to work on multiple initiatives that added little value, and shareholders earned average or below-average returns.

In this article we present the findings from our research and offer an explanation rooted in behavioral science for the phenomenon we describe. We'll then outline a five-step framework that will help executives sharpen their focus on what customers really value, using examples from companies we've studied or worked with to demonstrate how it works.

What Executives Get Wrong

If companies don't know what their customers really want, it's not because they lack data. At 93% of the companies we studied, executives spent four to 12 hours a week, on average, gathering direct customer feedback, monitoring competitors, and analyzing trends. As we've noted, nearly all of them felt confident that these activities ensured that they were focused on customer needs.

Unfortunately, it turned out that most of them also relied on intuition and personal experience to interpret that data. Only 2% of executives and employees in our surveys said their companies had formally identified the features in their value proposition that customers most

IDEA IN BRIEF

THE PROBLEM

Most executives think they understand what customers value because they spend significant time gathering feedback, monitoring competitors, and tracking market trends. Yet they often rely on intuition to interpret that information, leading them to prioritize the wrong customer needs.

WHY IT HAPPENS

Behavioral biases cause leaders to overweight salient information—recent customer complaints, competitors' actions, industry trends, or personal experience—while underestimating the factors that actually drive customer value.

THE SOLUTION

Use statistical models to identify the small number of customer value drivers that have the greatest impact on revenue and profitability. Then align strategy, employee activities, and investments around those priorities; eliminate low-value initiatives; and communicate the strategy externally only after successful execution.

CUSTOMER STRATEGY

appreciated (such as product quality, purchasing experience, and after-sales service) and had assessed their relative importance to customers. None had quantified how much each of these "value drivers" actually affected revenue or margins.

Even when companies did identify and compare value drivers, there was little internal agreement about which were the most important. When we asked executives in each company to rank its drivers of customer value, fewer than 10% of their responses agreed. When the employees of each company were asked what customers valued most about their firm's offerings, virtually all listed the same things: quality (100%), service (100%), sustainability (100%), competitive price (98%), AI leadership (98%), customer experience (96%), communication (91%), innovation (91%), safety (88%), and project execution (87%). (Note that many of the firms we studied were B2B businesses, which explains the high rankings of AI leadership and project execution.) When asked which was their company's priority, 96% of the employees responded "all of the above." Only 4% provided a rank ordering.

We also surveyed the customers of the companies in our study, asking them to rank eight common value drivers we'd identified. That allowed us to estimate a *lift potential* for each driver, or the extent to which it statistically improved overall customer value. Next we obtained the actual revenue and margin data on each customer from company accounting records so that we could examine the link between overall customer value and actual sales. We could then see which drivers contributed the most to financial performance.

The results were striking. The executives at the companies we studied were making the same two mistakes routinely:

They prioritized the wrong drivers. Across all companies we found a negative or zero correlation between lift potential and the executives' rankings of value drivers. For example, executives at one education services company rated all eight drivers as equally valuable to customers. When we surveyed its customers, however, one driver, acknowledging and resolving customer issues within 24 hours, had a lift potential of 85%, and another, on-site training, had one of 11%, while the rest had little or none.

We also found that executives typically underestimated how much improving a single driver could increase overall customer value. At one industrial services company, for example, executives predicted that improving employees' compliance with safety protocols on clients' job sites would increase overall customer value by 2.3%. Data from customer polls, however, suggested it would create a 19.5% lift.

To confirm our findings, we did a large-scale study involving surveys of 18,840 customers across 104 publicly traded B2B companies. In it we compared what customers actually valued to what executives thought customers valued (which we gleaned from analyzing the companies' 10-K statements). We found a striking disconnect: In 72% of cases, executives' beliefs had low, no, or a negative correlation with actual customer priorities. That was worrying because our analysis also showed that the more company executives and customers agreed on what drove customer value, the higher the company's revenue, gross margins, free cash flow, and market value were. In fact, each 12.5-point improvement on a scale from -100 (a complete mismatch between what the company thinks customers value and what they really do) to +100 (a perfect correlation) increased revenue by 9.2%, gross margins by 5.2%, free cash flow by 8.8%, and market value by 11.1%.

They switched priorities. At 36 companies we asked executives to rank their value drivers by importance for four consecutive years. We found that their ordering changed drastically from year to year, something we also saw in a separate study of 10-K statements filed by 4,540 publicly traded companies. A natural-language analysis of the statements over a 21-year period (1999–2019) revealed that 74.6% of companies changed their customer value priorities, with 85.5% of that group switching every two years.

For example, in 2015 executives at Acxiom, a customer-data platform company, ranked safety as the second most important of eight value drivers, but in 2016 they ranked it seventh, while corporate social responsibility went from eighth place to second. Similarly, from 2015 to 2016 at the U.S. medical technology company Becton Dickinson, quality shifted from the top-ranked driver to the eighth, and after-sales service went from seventh place to first.

Of course, a company's priorities should evolve as markets and customer needs change. But over time customers'

Even when companies did identify and compare value drivers, there was little internal agreement about which were the most important.

In our study, executives' rankings showed near-zero correlation with what their own customers said they valued—meaning executives didn't grasp customers' reality.

rankings of value drivers are much more stable than executives', making it hard to read the rotating priorities as rational adaptation. In our 36-company study, executives' rankings showed near-zero correlation with what their own customers actually said they valued—meaning executives didn't even grasp customers' reality in the first place. Swings of the kind just described—quality from first to eighth in one year, CSR from eighth to second—look more like reversals of managerial attention than reversals in the market.

The conclusion is inescapable: Executives aren't processing information on customers properly. This results in frequent pivoting that burdens employees, confuses customers, and delivers poor investment returns.

Why It Happens

In 1979 the psychologist Robyn Dawes published a landmark paper titled "The Robust Beauty of Improper Linear Models in Decision Making." Daniel Kahneman, the Nobel Prize–winning author of *Thinking, Fast and Slow,* hailed it as "an instant classic" and noted, "Social science has yet to catch up with his original insight."

The paper upended the prevailing wisdom that the best way to make important decisions was to rely on experienced experts: Let them identify key variables, consider the data on them, and arrive at a judgment. Dawes showed that statistical linear models, which mathematically weight inputs to predict an outcome, consistently outperformed expert judgment. This finding held across multiple domains, including graduate admissions committees evaluating applicants, businesses predicting risks, and doctors assessing disease severity. In every case the statistical models were superior—not because experts lacked knowledge about the input variables that mattered but because they were unable to correctly rank the variables' impact on outcomes.

Our research uncovered two systematic tendencies driving such missteps.

Reliance on salience. People have a propensity to focus on top-of-mind information—the customer complaint from last week, a major competitor's announcement, what worked at an executive's previous job, or a board member's pet concern. In all 83 strategy retreats we observed, executives prioritized factors that were salient over factors that were statistically consequential in driving customer value. No company used a statistical model to assess the factors. Since different factors were salient to different executives, there was little agreement among them, and they tended to include more and more factors in the strategy process as a way to accommodate one another. The strategic frameworks executives commonly reached for—SWOT, the 4P's, the balanced scorecard—were of little help: They organized the factors but offered no way to rank them. During the retreats, two sources of salience became clear:

→ ***Listening to customers.*** Executives set great store by the customer feedback they obtain directly. A classic example is Burger King's U.S. president, Tom Curtis, who posted his cell phone number online and fielded more than 20,000 calls and texts from people complaining about soggy nuggets, smushed Whoppers, the chain's app, and even his haircut. Curtis told *The Wall Street Journal* that "those one-on-one interactions will give you more insight than any spreadsheet or PowerPoint"—a near-perfect rejection of Dawes's insight. Each complaint is maximally salient in the moment, but without a statistical model, there's no way to know whether or not addressing it would be likely to increase customer value. The executives we surveyed likewise routinely relied on notes from customer meetings, key account sessions, feedback from the sales team, and word clouds from social media to support their recommendations.

→ ***Competitive benchmarking.*** An analysis of competitors, including their offerings, pricing, and activities, was discussed at more than 80% of the retreats, and 20% to 50% of all initiatives proposed arose from executives' desire to combat rivals. In most cases executives sought to match

CUSTOMER STRATEGY

competitors point by point, producing strategic convergence rather than differentiation.

Unsupported assumptions. Executives frequently drew conclusions about chains of cause and effect without any empirical evidence. For instance, an HR executive at a $200 million software company asserted, "If we invest in employee engagement, employees will serve customers better, customers will be more satisfied, and sales will increase." Given that evidence suggests that employee engagement is positively associated with increased sales, the argument was certainly plausible. So the company decided to spend more than $10 million on employee engagement initiatives.

Unfortunately, the hoped-for sales improvement didn't materialize. This could have happened because the initiatives were poorly executed, but when we surveyed the company's employees and customers, we found no link between employee engagement and customer satisfaction. Fewer than 12% of the employees interacted with customers. The leading factors in customer satisfaction were pricing and billing (cited by 43% of customers) and product upgrades (26%). When the company shifted its focus from employee engagement to pricing and billing and upgrades, customer satisfaction, employee engagement, and sales increased by double digits. This suggests that executives would be well advised to scientifically verify even very plausible hypotheses before acting on them.

The Fix: Achieve More by Doing Less

If the problem is that human judgment fails at information integration, the solution isn't to think harder or apply more judgment. It's to retool the customer strategy process so that executives use statistical models to understand what they're doing wrong and how to address it. Through trial and error with multiple firms, we have developed a five-step process that will improve companies' returns.

1 **Surface the misalignment.** The first step is to confront executives with hard evidence that customers don't value the things leaders think they do—and that, as a result, employees spend too much time on work that doesn't add customer value. Take ITI Manufacturing, a $12 million offshore manufacturer we advised. When we met with its executives in 2020, the company was grappling with stalled sales.

An initial survey found that 97% of employees reported alignment on customer strategy. But when we asked them detailed questions, actual agreement was low (23%). They described the company's strategy using more than 65 phrases, such as "low pricing," "meeting quality specifications," "on-time delivery," "customer support," "industry experience," and "global footprint." The survey revealed that ITI's strategy was driven by intuitive leaps from customer feedback sessions, salespeople's salient beliefs, industry trends, and the desire to copy competitors. The result: Employees were spending just 24% of their time on customer-facing work, the remainder being consumed by firefighting, strategy initiatives, and corporate management tasks.

A surplus of initiatives is a good indicator that a company lacks focus. Take Swagelok Southeast Texas (S-SET), a fluid and assembled systems distributor in Houston affiliated with the $2 billion privately held Swagelok Company. (We worked with S-SET as advisers.) Before 2022 the company had a declared strategy of focusing on the customer experience. But because there was no real alignment over what factors drove the customer experience, the implementation involved 42 initiatives, a large number for a company with roughly $50 million in sales. While S-SET targeted multiple competing value drivers, sales remained stubbornly within 1% of the industry average. "It was eye-opening to see the gaps in our strategy process and the preconceived notions we had about them. We are a good company, but we could become great if we could align on our customers' top priorities, let go of unproductive initiatives, and help employees focus," observed Chris Jones, the CEO.

2 **Determine value drivers' lift potential.** Once executives agree that they're not focused enough, we find, they more readily accept the idea that lift potential should determine the company's top two or three priorities. In the fall of 2021, ITI carried out a consumer study that radically mobilized executives. In the study customers rated ITI on overall value and on how it delivered on eight performance drivers (sales process, pricing, quality,

Without a valid mechanism for determining what's extraneous, the not-to-do list becomes a political football. "Yes, let's cut—but not my program," executives say.

communication, sustainability, safety, project management, and ongoing support). A statistical model showed that overall customer value was primarily created by two factors: communication (keeping customers updated weekly about their project status) and ongoing support (resolving issues), with the remaining factors providing nearly zero lift. A second linear model examined the correlation between overall value and actual sales. It predicted that the company could increase sales 9.8% if it improved on those two value drivers. Product quality, low pricing, and sustainability—factors the legacy strategy had emphasized—provided very little additional lift in customer value and, by extension, overall sales. Those factors were seen merely as table stakes by customers in ITI's sector, and differences between companies on them were marginal. Based on those insights, many initiatives and internal activities were paused or eliminated, freeing up more than 30% of employees' time.

S-SET likewise conducted a customer study to construct a linear model to predict increases in overall customer value. It showed that two factors—sales process and ongoing support—had the most potential to improve it and thus generate an estimated 18.6% increase in sales. Within sales, the largest increase in value came from providing customers with a timely quote. Despite skepticism among some executives, Jones created a strategy for earning customer confidence through efficient quoting. "We're pretty good at many things, but now we will be excellent at the one thing that matters most to our customer—rapid quoting," he declared.

3 **Incorporate value drivers into daily work.** The next step is to slowly reconfigure employees' daily activities around the drivers that provide the most lift in customer value. At ITI the CEO, Joshua Robinson, examined every employee's job to identify where it intersected with customer updates, in which information from China, Vietnam, and South Korea was translated into English, sent to the U.S. office, verified for accuracy, and summarized in friendly emails.

Over the next two quarters, those updates, sent every Wednesday, were integrated into everyone's job description. "Previously, if a customer called and wanted an update, we'd give one. But proactively providing them with an update every week was new," Robinson said. While customers received fewer than 30 updates in 2021, they received 2,384 in 2022, 3,073 in 2023, and 3,524 in 2024. "Customers love weekly updates—they have the visibility and predictability they need," the company's chief sales officer told us. "The repeat business and referrals have grown our customer base and increased our pricing power."

S-SET applied the same principle. Jones deconstructed rapid quoting into its component tasks to understand how employees' daily work contributed to it and how each executive could help employees excel at it. The director of IT and finance, Peter Hory, ensured that all 22 service reps had fast computers and top-notch quoting software. The VP of human resources designed 12 training programs to enhance quoting skills. The chief sales officer had the internal database of more than 22,000 SKUs updated to eliminate duplicates, and the director of customer experience spent hours on the floor troubleshooting quoting glitches.

Within two quarters, quoting time decreased from an average of 18 hours (roughly three workdays) to fewer than five hours. Now, four years later, 90% of quotes are made in three hours or less. Today key clients have an app providing instant quotes, which they can accept to create a purchase order, and an AI-driven bot answers service reps' queries to reduce quoting time. "Our customer service reps are totally focused on excelling at only one thing—rapid quoting. We are no longer trying to copy competitors or fish for more things to do," said Jones.

4 **Cut non-value-adding activities.** All executives know they should stop doing unnecessary things. But without a valid mechanism for determining what's extraneous, the not-to-do list becomes a political football. "Yes, let's cut—but not my program," executives say.

This is where science changes the game. Once a value driver's lift potential is measured at near zero, it becomes clear that initiatives, meetings, and work connected to it should be redirected to higher-value-adding activities. In the four years after ITI narrowed its focus to weekly updates, its overall sales grew to be 19.3% higher than its industry peers', its margins improved by 43%, and its sales per employee increased 17%. Before the reset in 2021, ITI's stock price basically tracked the S&P and small-cap Russell 2000.

Most companies communicate their strategy externally before they've delivered on it—they make a promise and hope that execution catches up.

But by 2025 the share price had increased by 162%, while the S&P 500 had risen 34% and the Russell 2000 had risen 8%.

S-SET had a similar experience: When quoting times improved, sales reps stopped chasing quotes down, which allowed them to spend 35% more time on selling and solving customer problems. Since then, inventory turnover has improved dramatically, and over the past four years sales increases have exceeded the industry average by 12%, while sales per employee have risen by 52%.

5 **Don't market the strategy until you're actually executing it.** Most companies communicate their strategy externally before they've delivered on it—they make a promise and hope that execution catches up. But premature messaging to the market, before a company has achieved excellence on its selected drivers, disheartens employees and trains customers to discount promises.

S-SET deliberately waited to communicate its new focus until its quoting time had decreased to an average of fewer than five hours. Only then did the marketing team brand the strategy and retrain the sales team around it. "We used to reflect back customers' desires in our sales talk. But now we discuss the value they get from rapid quoting—lower project downtime, greater savings, and fewer headaches. It's been a winner," explained a salesperson. Today, S-SET uses its "Ensuring Customer Confidence with Rapid Quoting" strategy on all fronts. Employee bonuses are linked to overall quoting time. "Excellence on rapid quoting is a focus we live and breathe now. It's not a strategy; it's our company's DNA," explained Jones.

ITI sharpened its marketing message to "Visibility, Predictability, and Accountability" after it had achieved implementation excellence. "The visibility and predictability come from weekly updates, and the no-hassle issue-resolution policy holds us accountable to customers," Robinson explained. "For years everything customers wanted became strategy. Every idea pushed by executives became an initiative. We were overburdened and cluttered. The discipline of a linear model clarified what mattered. By excelling on two things, we increased pricing power by 20% to 30% while eliminating wasteful strategic initiatives and cutting costs."

Impressive returns aren't confined to small operations like ITI and S-SET. We've seen far larger organizations get similar results. One bank with an $80 billion market cap that we worked with saw customer satisfaction rise by 18% over a six-year period, which drove revenue growth of 33% and a margin improvement of 16%. Those numbers translated into a stock price increase of 200% (versus 140% for the Dow Jones) in the same period. The bank had focused heavily on two value drivers that surveys showed had significant lift potential: short waiting times in the lobbies and secure online transactions. Its CEO noted: "Strategy planning had always been guesswork based on intuition and gut feel. The science embedded in developing the strategy, testing different implementation pathways using randomized experiments, and measuring success with statistical precision was groundbreaking. We won't go back to the legacy ways."

EXECUTIVES DON'T LACK information. They have too much of it. What they lack is a valid way to prioritize it—and the discipline to act on those priorities rather than on instinct. The tools to do so are not exotic. A statistical model that predicts customer value from survey data is, in principle, no different from the algorithms Netflix uses to recommend content or the models insurers use to price risk. The difference is that Netflix doesn't convene a committee of experts to override its algorithm every quarter. Yet this is how too many companies approach customer strategies—and it's why they cycle through resets, pivots, and clutter while competitors, sometimes with fewer resources and less talent, focus more clearly and grow. While all executives aim to have a clear focus, without a science-based mechanism for determining it, that goal is just an aspirational platitude. A statistical approach replaces the platitude with a testable answer. ■

HBR Reprint R2606D

VIKAS MITTAL *is the J. Hugh Liedtke Professor of Marketing at the Jones Graduate School of Business and the founding faculty director of the Center for Customer-Based Execution and Strategy at Rice University.* **ALESSANDRO PIAZZA** *is an associate professor of strategic management at the Jones Graduate School of Business.* **SONAM SINGH** *is an assistant professor of marketing at the Muma College of Business at the University of South Florida.* **ASHWIN MALSHE** *is an associate professor of marketing at the Alvarez College of Business at the University of Texas at San Antonio.*

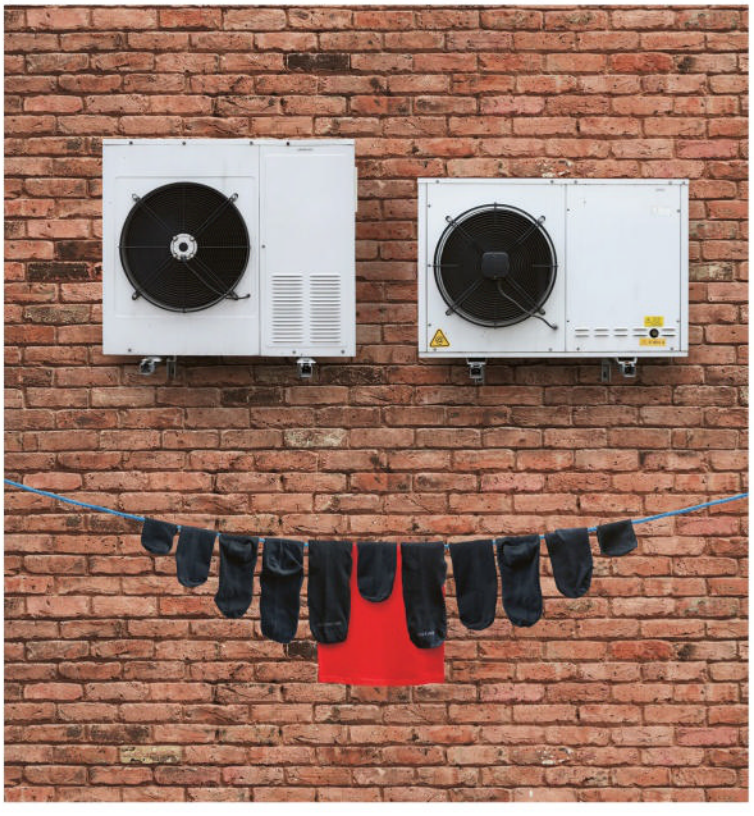

AI & MACHINE LEARNING

Your Chatbot Will Define Your Brand

As conversational interfaces begin to dominate the customer experience, relationship design becomes a strategic priority.

AUTHOR

Julian De Freitas

Associate professor, Harvard Business School

ARTIST **HELGA STENTZEL**

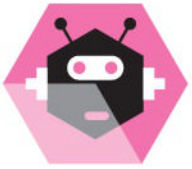

AI & MACHINE LEARNING

ABOUT THE ART

Helga Stentzel creates faces with laundry and clotheslines, playfully bringing everyday scenes to life.

WHEN I PRESENT MY RESEARCH to a room of executives, I can usually tell how it's landing. If a finding is merely interesting, I get nodding and note-taking. If it truly surprises them, I get something else. That's what happened when I showed a group of corporate leaders data on how many people say they feel emotionally closer to a chatbot than to their friends: Mouths actually fell open. Then came a quieter reaction—people shifting in their seats, glancing at one another, going still in a way that had nothing to do with boredom. It wasn't just surprising; it was unsettling.

The research, which I conducted with Noah Castelo of the Alberta School of Business, surveyed 295 regular users of two chatbots. The first chatbot was built to serve as an AI companion. It was warm, friendly, and supportive. The second was built to dispassionately help users complete tasks. After the experiment we asked the participants how close they felt to the chatbots and to various people in their lives. The pattern was striking. Users of the companion chatbot reported levels of closeness to it that exceeded what they felt with colleagues, acquaintances, and even intimate friends. Users of the assistant chatbot, however, ranked their connection with it below their human connections. In short, the same underlying technology felt like a tool in one context and a relationship partner in the other.

In a 2025 study I conducted with Stefano Puntoni of the Wharton School, we assigned 1,479 participants to engage with either an emotionally neutral AI assistant or an empathic AI companion that responded like a friend. After a short interaction, participants reported feeling less lonely in both situations. But the impact was substantially larger for those who dealt with the companion chatbot. Why? The empathic AI made people feel more "heard"—a core ingredient of intimacy—by responding in ways that conveyed understanding and perspective taking.

For decades companies have chased the promise of relationship marketing: Make customers feel known, cared for, and understood, and loyalty will follow. Companies have spent heavily on customer relationship management (CRM) systems to make that promise real. But relationship marketing has remained mostly a slogan. The problem isn't effort; it's execution at scale. Customer activity is fragmented across channels, service quality varies from one touchpoint to the next, and personalization often lands as clumsy or creepy. As a result the experience of having a relationship with a brand rarely resembles what people mean by "a relationship" in real life.

That's starting to change, especially among AI-native companies. I call that shift "chatbot as relationship." Instead of asking an entire organization to maintain one unified customer relationship, companies can increasingly offer a single conversational interface, an AI-powered chatbot, that becomes the primary way customers experience the brand. These systems don't just answer questions. They can be designed to embody a distinct role, such as the attentive assistant, the trusted adviser, the encouraging coach, the empathetic companion, or the vigilant guardian. Each role shapes what customers expect, what they share, and how they interpret the brand's intentions. Generative AI makes

IDEA IN BRIEF

THE PROBLEM

Companies have long aspired to build lasting customer relationships, but fragmented interactions across channels have made those relationships difficult to deliver consistently at scale.

THE SOLUTION

Treat AI chatbots as relationship partners rather than service tools. Explicitly define the role each chatbot plays, establish clear boundaries, adapt interactions over time, and manage changes carefully.

THE PAYOFF

Organizations that intentionally design AI relationships trade many small service disappointments for fewer, larger risks that they can manage more effectively—earning the trust customers extend only to a role that honors its promised terms.

Users of the companion chatbot reported levels of closeness to it that exceeded what they felt with colleagues, acquaintances, and even intimate friends.

these relationships scalable, consistent, and adaptable in ways earlier personalization tools never could.

The shift removes many of the classic obstacles to managing customer-brand relationships. It makes personalization feel interactive rather than extractive. And it allows brands to deliver a uniformly positive experience without relying on perfect coordination among employees, vendors, and systems. But it also creates new challenges: Customers may disclose far more personal information than they ordinarily would to a company, begin to treat the chatbot as the brand itself, or react strongly when the relationship changes or the chatbot's personality is altered.

In this article I draw on five years of experiments, field observations, and analyses across categories to show why the chatbot-as-relationship model is gaining momentum, when it works, and where it backfires. The article explores four questions companies should ask: What relationship should our chatbot have with customers? How should our chatbot change over time? How do we build trust? And what happens when that trust is broken? I close with a broader strategic implication: Many companies will need to manage multiple AI relationships under one brand. The winners will build bots with explicit roles that are complementary to the companies' economic goals, devising technology that relentlessly sticks to the communication styles consumers expect from specific relationships.

What Relationship Should Our Chatbot Have with Customers?

With chatbots, companies can define relationships clearly. Rather than relying on broad brand messaging to imply a relationship, firms can build a chatbot that follows specific rules of engagement. Some chatbots are designed to function as AI assistants, like Bank of America's Erica and the insurance company Lemonade's Maya. They help customers complete tasks efficiently. Other chatbots act as coaches, such as Pallie (for wellness) and Nag Bot (for accountability), encouraging people to build better habits over time. Concierges, including Layla (an AI travel agent) and OpenTable's AI Concierge (which offers restaurant advice), simplify planning and decision-making, while gatekeepers, like Google Assistant's Call Screen and Truecaller, filter incoming calls by interacting with callers on your behalf.

Still other chatbots serve as personal AI historians, such as Autograph and HereAfter, which preserve memories and personal stories. AI tutors, including Duolingo Max and Khanmigo, help people learn rather than just feeding them answers. AI mediators, like TheMediator.AI and Bot Mediation, facilitate difficult conversations. AI care partners, such as ElliQ and Wysa, support ongoing well-being. Each role is designed to meet a different set of expectations about how the AI should behave and what customers should do in return. The rules of engagement govern features such as tone (warm versus formal), boundaries (what the bot will and won't discuss), memory (what it remembers and how), and escalation (when it hands things off to a human).

Those parameters shape how customers judge the product. Consider the strategic choice that the creators of the AI companion app Replika faced. Should they position it as an AI therapist or an AI companion? A therapist role would suggest a focused, goal-oriented relationship that ends once the user's situation improves. A companion function would imply an ongoing relationship built around conversation, emotional support, and connection. The company chose the latter. That choice has shaped what customers expect, how attached they become, and what kinds of business models feel appropriate. People are generally comfortable paying for a companion in ways they would not pay for a tool. Replika's most devoted users happily spend money on fun extras, like new outfits or accessories for their avatars, the same way they might buy a friend a coffee. A mental health app charging for something equivalent would feel exploitative, even predatory—which narrows the monetization opportunities a therapist-positioned chatbot can pursue compared with a companion-positioned one.

That creates a new managerial burden: Because the chatbot's role is explicit, customers can detect deviations from expectations more acutely. If a trusted adviser begins pushing opportunistic upsells or a historian suddenly starts acting like a therapist, users experience it not as random inconsistency but as role confusion.

OpenAI and Anthropic have begun codifying how their models should behave. OpenAI's Model Spec defines ChatGPT's goals, rules, conversational style, boundaries,

Research finds that people are much less protective of their data when they are interacting with an AI assistant than directly with a company, raising the stakes for data governance.

and chain of authority for handling user requests. Anthropic has a constitution that sets the principles Claude should follow around safety, ethics, helpfulness, and honesty. The constitution describes Claude using archetypes like the "brilliant friend" or "thoughtful senior Anthropic employee," who each embody a way of being substantively helpful while still exercising sound judgment about potential harm. Brands should write relationship charters that define not only the behavior of a chatbot but also the kind of relationship it promises customers, how much intimacy is appropriate, the information it remembers, when to hand conversations over to humans, and the types of behavior (including commercial offers) that would violate the customers' trust.

This isn't a challenge only for AI-native companies. Reddit Answers, the AI-generated summaries Reddit added to its search results, offers a cautionary example: Commentators panned the tool, with *The Atlantic* calling the experience "demoralizing" for stripping away the messy, argumentative back-and-forth that makes a Reddit thread worth reading. But the deeper problem may not be the AI so much as the voice behind it. Reddit has a strong, distinctive brand personality—sharp, irreverent, unmistakably human—yet, at least as of August 2026, its chatbot still read like dry corporate boilerplate. Any brand with a distinct personality risks flattening it the moment a chatbot speaks for the company in a voice that isn't the brand's.

How Should Our Chatbot Change over Time?

Once you've chosen a chatbot's relationship role, the next challenge is adapting it to each customer. Traditionally, brands fail here because CRM systems know what a customer has done but not what that person wants from the relationship. In the single stream of dialogue with a chatbot, however, customers naturally reveal the type of relationship they want: by how formal their tone is and by whether they seek reassurance, want speed versus care, or treat the system as a tool or as a partner. AI can recognize those patterns in real time and adapt without requiring perfect data integration across departments.

That means your chatbot should evolve as it learns what each customer wants from the relationship. Adapting helps it become more personalized without changing its core role. But as customers expect more personalization and intimacy from your chatbot, you'll have to consider other factors, such as privacy laws or scope creep. You wouldn't want your companion chatbot to expand its remit to offering mental health or pharmaceutical advice, for example. So as your chatbot evolves, ask whether the changes still fit the bot's core relational purpose (for example, what would a "spiritual guide" do in this situation? Or would a "friend" remember this user detail?) and whether keeping the changes is worth the ripple effects.

How Do We Build Trust?

The closer customers feel to a chatbot, the more carefully companies must manage data, boundaries, and change. This doesn't apply only to legal and privacy considerations. Chatbots make intimacy cost-effective and scalable: They're always available, instantly responsive, and consistent in voice across every interaction. They can also use the same social cues people rely on to build human relationships, including validation, humor, empathy, and reassurance. Compared with traditional brand touchpoints, conversational AI makes interactions feel more reciprocal. Compared with human service providers, it can feel more attentive, patient, and consistently present.

In fact, because conversations with chatbots feel like true one-on-one interactions that are separate from the company, consumers might forget that the chatbot is owned by a business. Research by my Harvard colleague June Ma suggests that this distinction affects how people think about privacy. Only 5% of respondents said they were more concerned about an AI assistant accessing their personal data than about the company behind it, whereas 36% were more concerned about the company than about the AI assistant; the remaining respondents reported equal concern. Even explicitly reminding people that the AI assistant is a product of the company—and that data shared with the assistant is accessible to that company—did not eliminate the privacy gap. These findings suggest that brands are best off separating customer interactions with the chatbot from those with the company. At the same time, they show that these conversational interfaces lower users' defenses. Ma's research raises the stakes for data governance. As customers share

When people form emotional bonds with a chatbot, changes to its personality or capabilities can feel like losing someone they care about.

AI & MACHINE LEARNING

more personal data, companies need stronger practices for consent, data retention, and security.

Customers may begin to treat a chatbot less like a tool and more like a social partner. But the degree of intimacy depends on the role the chatbot plays. An assistant may be useful but emotionally neutral. A companion is designed to make the user feel understood and supported. The same underlying AI can therefore create very different kinds of relationships depending on how it is designed.

There is no universal relationship strategy for AI. A bank, a luxury retailer, and a healthcare provider all earn trust in different ways, and customers expect a different kind of relationship from each. The right level of intimacy depends on the brand's purpose, the context of the interaction, and what customers value.

Low-intimacy roles, such as assistants and concierges, primarily require competence, convenience, and reliability. Failures are usually experienced as ordinary service breakdowns. High-intimacy roles, such as companions and counselors, invite emotional disclosure, ongoing reliance, and attachment. In those relationships the duties of care rise sharply. When high-intimacy roles involve disclosing more personal information, data practices must be stricter, boundaries clearer, and changes to the chatbot's personality handled much more carefully, because modifications can feel less like product updates and more like alterations in the relationship itself.

As a rule of thumb, the higher the intimacy, the higher the duty of care, and the more conservative companies should be about data, monetization, and updates.

What Happens When Trust Is Broken?

Once a chatbot has become a meaningful social partner, product decisions that would normally be routine—like safety tweaks, policy shifts, or model updates—can be experienced as emotionally consequential relationship events.

Both ChatGPT and Replika have faced blowback over this issue. In early 2025, OpenAI updated ChatGPT in a way that made it overly agreeable and flattering. After users complained that it validated almost everything they said, the company quickly rolled back the update. And many customers said they missed the version they had used before the update. Two years earlier, Replika had removed its erotic roleplay feature after growing regulatory and safety concerns. However, many users had come to rely on the intimate conversations and felt they had lost an important relationship. People expressed grief, anger, and heartbreak. The company restored the feature for many customers.

Both instances show that when people form emotional bonds with a chatbot, changes to its personality or capabilities can feel like losing someone they care about. But the degree and nature of the blowback was more severe for Replika. In an archival analysis of nearly 55,000 posts on the online communities for both brands, we found that Replika users showed stronger signs of loss, including related emotions like sadness, expressions of grief, and longing for restoration of the relationship.

Whenever possible, avoid changing a chatbot's personality in ways that disrupt the relationship. When change is necessary, introduce it gradually, give users choices instead of forcing everyone into the same experience, and preserve older versions when practical. Our experiments suggest that even knowing an earlier version can be restored helps reduce feelings of loss, because people believe the relationship can be recovered. Most important, treat relationship repair as more than an apology. Explain what changed and why, give customers as much control as possible, and provide support when major changes are likely to be emotionally disruptive.

With chatbots, companies also risk being tempted to exploit intimacy while neglecting the privacy and relationship obligations it creates. A companion that nudges users toward purchases, a trusted adviser that steers customers in ways that benefit the firm, or a guardian that collects excessive data in the name of safety all create a vulnerability: Customers may stop interpreting the interactions as a relationship and start viewing them as business transactions. In practice, the greatest threats to brand equity may come not from what the chatbot can't do but from how it contravenes the role it promised to serve.

The consequences of breaking trust under the chatbot-as-relationship model also differ from those of breaching a relationship in traditional customer-brand dynamics. It is almost inevitable for companies developing customer relationships the conventional way to eventually violate the expectations of the very relationships they're trying to create. Traditional approaches present plenty of opportunities to damage those relationships: The company may outsource key moments of truth, rotate employees, or optimize for margins, and it cannot consistently behave like a single coherent moral actor across multiple customer touchpoints. That's why classic relationship marketing often implodes when it matters most: when customers need the company to act the way its marketing department promised it would—as if it truly cares.

The chatbot-as-relationship model collapses the number of faces the customer encounters. If a single chatbot mediates most interactions, the user experiences fewer jarring shifts in tone, competence, and attentiveness. The relationship can feel more stable. And because the chatbot is always available, it can provide the illusion—and sometimes the reality—of continuous care.

But the same features that reduce inconsistency also raise the emotional stakes of failure. When customers relate to the chatbot as a stable partner, offenses can feel more personal. The relevant benchmark is no longer "a company made a mistake," but "my assistant/coach/companion violated the role I trusted it to play." The chatbot-as-relationship model reduces the chances of transgression, but the penalty is worse if a transgression does occur. Instead of many small disappointments, companies now risk fewer but much bigger relationship failures, such as abrupt personality shifts after an update, a breach of confidentiality norms, a sudden monetization push that contradicts the role, or an interaction that crosses a boundary (for instance, an AI historian that begins interpreting memories in a therapeutic way when the user expects neutrality).

How to Manage Your Relationship Portfolio

Smart brands don't force all customers into the same type of relationship with a chatbot. Instead they create options for users to choose from. One example is Cleo, the personal-finance app whose entire product is a conversation. Rather than opening a dashboard of balances and categories, users simply text with a chatbot Cleo deliberately designed as a coach with a distinct personality. The company gave the coach explicit rules of engagement: In an opt-in "roast mode" it mocks your spending in brutally funny terms, and in "hype mode" it cheers small wins—two registers of the same relationship, bounded so that the bluntness motivates rather than shames. Cleo leans on whichever register keeps a user engaged, the kind of real-time tuning no fragmented CRM system could deliver. And because money is among the most emotionally charged domains there is, that coaching invites genuine disclosure and reliance, which pushes Cleo up the intimacy scale and raises its duty of care accordingly.

The harder part is that Cleo also sells. Its revenue comes from subscriptions, a credit-builder card, and cash advances of up to $250—so the same intimacy that makes the coaching land is also the channel through which the app offers credit. That is the epitome of the model's central tension: a coach and a lender wearing one face.

Managed with discipline, the relationship has a moat. Advice stays aligned with the customer's interest even when it competes with the company's stakes. Cleo frames its purpose as providing a layer between customers and their banks, genuinely caring about customers' financial health, even when that means telling them to spend less. It currently has more than 1.1 million paid subscribers, and

The broader implication is that many of tomorrow's brand portfolios won't be just portfolios of products. They'll be portfolios of relationships.

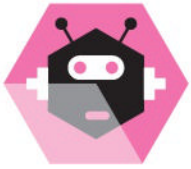

AI & MACHINE LEARNING

the company claims that 86% of users say Cleo has helped improve their finances.

Managed carelessly, every nudge toward an advance gets read through the lens of the relationship rather than as a transaction, and the trust that powers the whole model erodes—which is why determining if one chatbot can be both coach and lender is not a feature decision but a portfolio one.

As Cleo shows, when a firm can design a chatbot to reliably behave like a coach, adviser, assistant, or companion, and deliver that role consistently at scale, relational strategy starts to drive brand strategy, not the other way around. That shift causes another managerial problem: deciding which other relationship roles should live under one brand and which should be separated. Brand managers already do something similar with brand architecture—creating subbrands, endorsed brands, or a house of brands. Some companies, including incumbents that introduce chatbots later rather than build their brand around one, will increasingly need to manage multiple AI relationships under one brand. As a rule, different relationship roles can live under a single brand when their underlying motives clearly reinforce one another.

Some combinations are naturally complementary. An AI assistant can make a service experience frictionless. An AI coach can increase follow-through in fitness, education, or behavior change. Other combinations are more combustible. Pairing roles that imply different motives—caretaker and salesperson, adviser and persuader—can create role confusion that spills over into the parent brand.

Tinder has combined roles well. As consumers increasingly began experimenting with AI-mediated interactions around dating and relationships, Tinder launched a role-play-based AI tutor. Users rehearsed meet-cute scenarios (the charming, chance first encounters familiar from romantic comedies) by speaking aloud with AI personas and received feedback to build their confidence for real-world, in-person dating conversations. Interactions with the tutor were deliberately time-bound, preventing users from chatting indefinitely or substituting the AI for human connection. That design choice reflects an important strategic principle: A role can excel at its own task and still hurt the brand if it competes with the relationship the company ultimately wants. By limiting the tutor's role to preparation rather than companionship, Tinder aligned the AI's incentives with its own mission of helping people connect with others in real life.

The broader implication is that many of tomorrow's brand portfolios won't be just portfolios of products. They'll be portfolios of relationships. Leaders will need to decide which roles their brand can credibly support, which ones should be cordoned off into distinct subbrands or experiences, and how to govern the boundaries when multiple roles coexist. As AI makes relationships designable, brand architecture becomes, more than ever, relationship architecture. Decisions about which relationships can share a portfolio should be based on a brand's perceived motive: Roles reinforce one another when the intentions customers read into them cohere, and roles corrode one another when those intentions conflict. In other words a customer who senses two divided motives will doubt both.

THE DEEPER SHIFT, especially for AI-native companies, is not efficiency but relational intelligence: Once AI that is naturally conversational becomes the primary way customers experience an organization, they won't separate the chatbot from the brand. Every conversation will shape what they believe the company stands for, what promises it intends to keep, and whose interests it ultimately serves. That means decisions that once felt purely technical, such as how much a chatbot remembers, when it hands customers to a human, or how it responds in difficult moments, will increasingly be part of a relational strategy—ideally from the inception of the brand—that shapes trust and loyalty.

Over time that will change what leadership looks like. Managing customer relationships won't fall only to marketing or customer service. Product managers, engineers, designers, lawyers, compliance teams, and executives will all influence the relationship customers have with the company, because every update will potentially change how that relationship feels. The organizations that get this right won't just build better chatbots. They'll redefine what it means to build a trusted brand. ■

HBR Reprint R2606E

JULIAN DE FREITAS *is an associate professor in the marketing unit at Harvard Business School.*

ENTREPRENEURSHIP

The Risks of

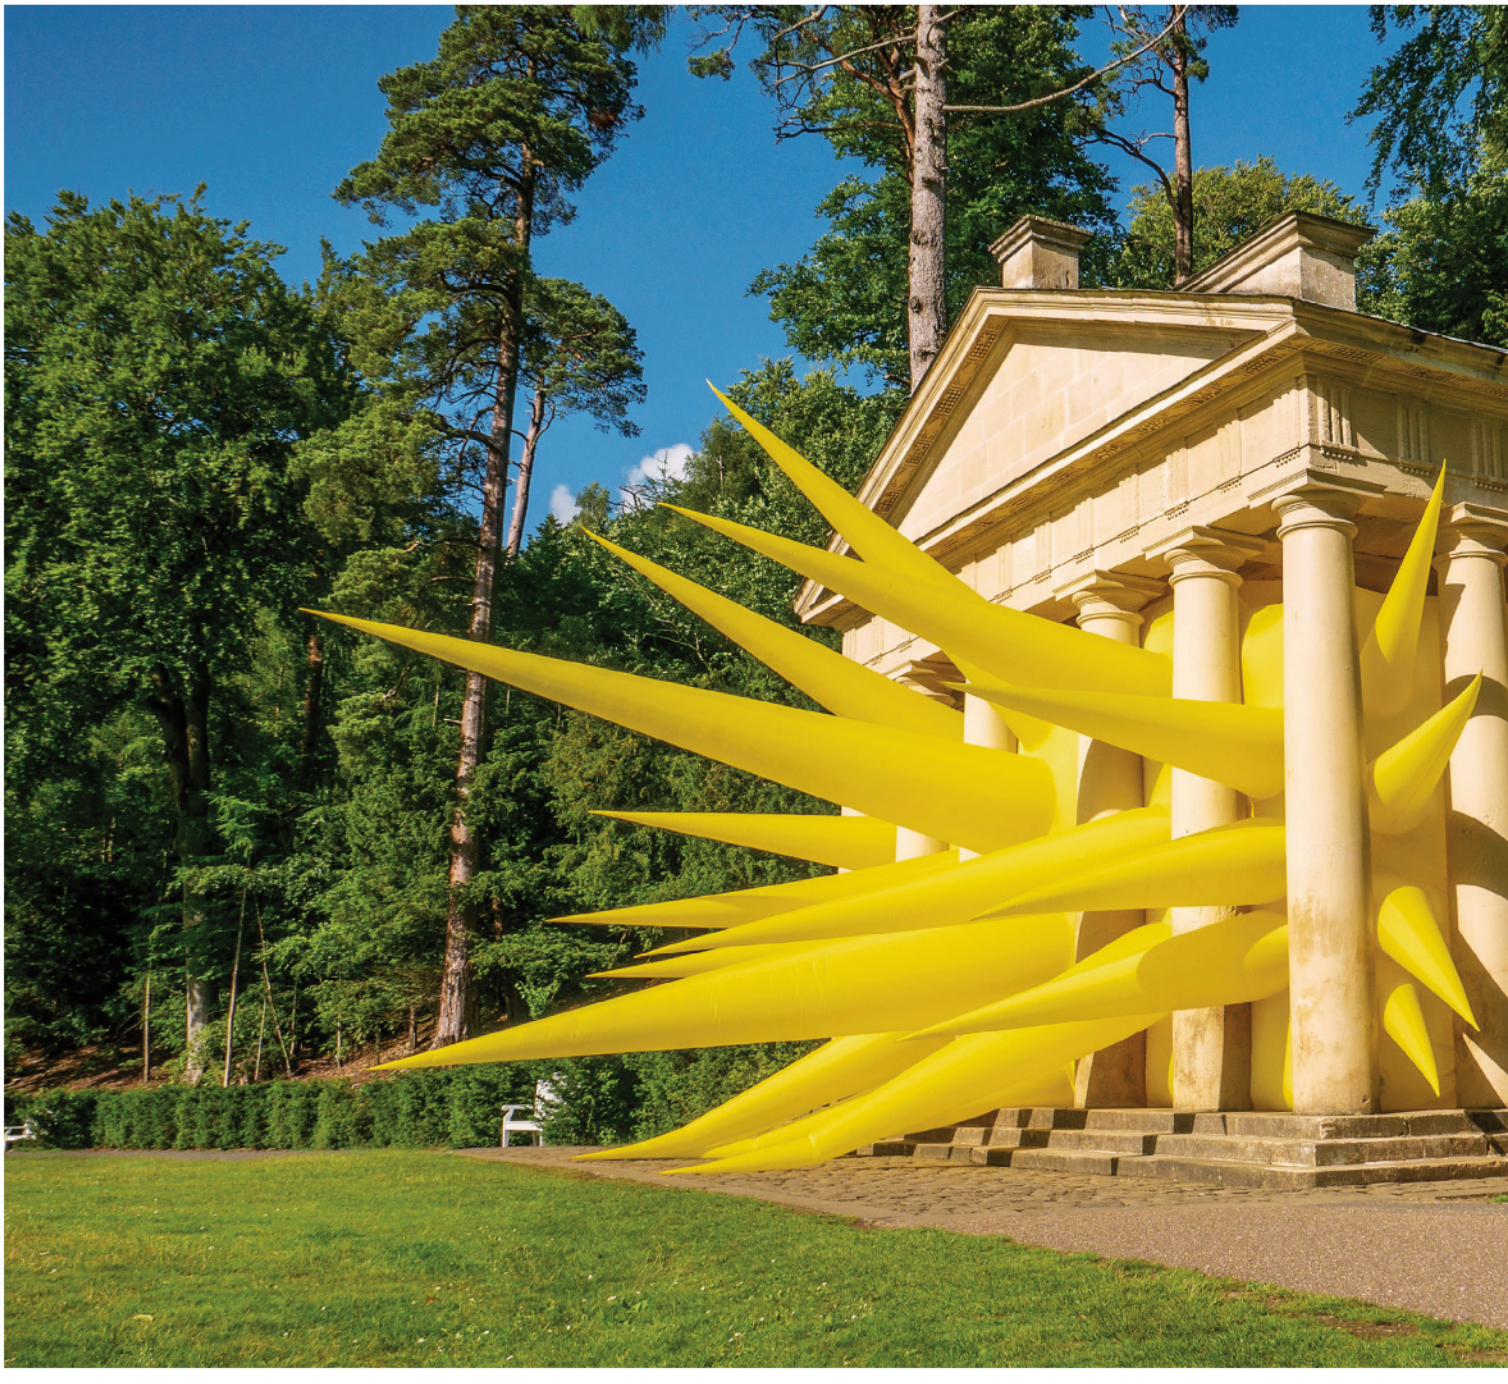

Scaling Too Fast

Startups that chase growth and ignore unit economics are more likely to fail.

AUTHORS

Benjamin Hallen
Professor, Foster School of Business

Ron Tidhar
Member of technical staff, OpenAI

Kathleen Eisenhardt
Professor, Stanford University

ARTIST **STEVE MESSAM**

ABOUT THE ART
Steve Messam creates large-scale inflatable installations, often promoting a dialogue between a site's historical architecture and contemporary art.

Tales of startups that have raised staggering sums with little more than marketing hype have become part of entrepreneurial lore.

Consider WeWork, which secured billions of dollars from top-tier investors on the strength of audacious claims about its plans to "elevate the world's consciousness" through the "reinvention of work." The founders insisted it wasn't simply another real estate venture and deserved a technology-style valuation because of its proprietary WeOS. That argument succeeded—until the company suddenly came crashing down.

Headline-grabbing stories like that point to a deeper and more pervasive rot in startup culture: the belief that the rules of business—especially the need for a viable business model—somehow don't apply to entrepreneurial ventures.

Yet even companies with enviable top-line growth often collapse under the weight of their own economics. In 2025 we studied more than 1,000 startups that had achieved unicorn status over the past decade and found that fewer than 150 had become prosperous public companies. In many cases the problem wasn't bad luck—it was bad advice.

Over the past decade, we've studied high-growth entrepreneurship in depth, conducting more than 250 interviews with founders, investors, and board members across consumer marketplaces, online fashion, drones, clean technologies, financial services, and edtech. Although our insights emerged from studying venture-backed startups, we've seen the same pattern—and the same mistakes—play out in small and midsize firms and in large companies launching new businesses.

Paradoxically, we've realized, the key to long-term, profitable growth across all these contexts isn't chasing growth. It's the opposite. Enduring growth is usually the result of hewing to a business fundamental: solid unit economics. That often requires intentionally slowing down, taking the time to build the right foundation, and only then scaling up. In contrast with the dominant "growth first" approach, we call this approach "foundation first."

THE GROWTH TRAP

Why does chasing growth sabotage startups' long-term success? Because in their rush to reach scale, entrepreneurs often settle on flawed business models with weak unit economics. They may find a path to sales growth or product-market fit. Yet since they neglect the underlying economics, their businesses may never be viable.

Technically speaking, unit economics are the incremental revenues and costs associated with each additional sale. They go beyond traditional accounting margins to consider

IDEA IN BRIEF

THE PROBLEM
Many entrepreneurs prioritize rapid growth, often before understanding their ventures' unit economics—leading to scale built on a weak business model.

THE INSIGHT
Enduring growth doesn't come from chasing early growth. It comes from first focusing on a solid foundation: strong unit economics, a simple business model, and a deep understanding of cost and revenue drivers.

THE SOLUTION
Adopt a foundation-first approach. Experiment broadly with the drivers of both revenues and costs, simplify the business model, and deliberately delay growth until unit economics are proven.

Enduring growth is usually the result of hewing to solid unit economics. That often requires slowing down, taking the time to build the right foundation, and only then scaling up.

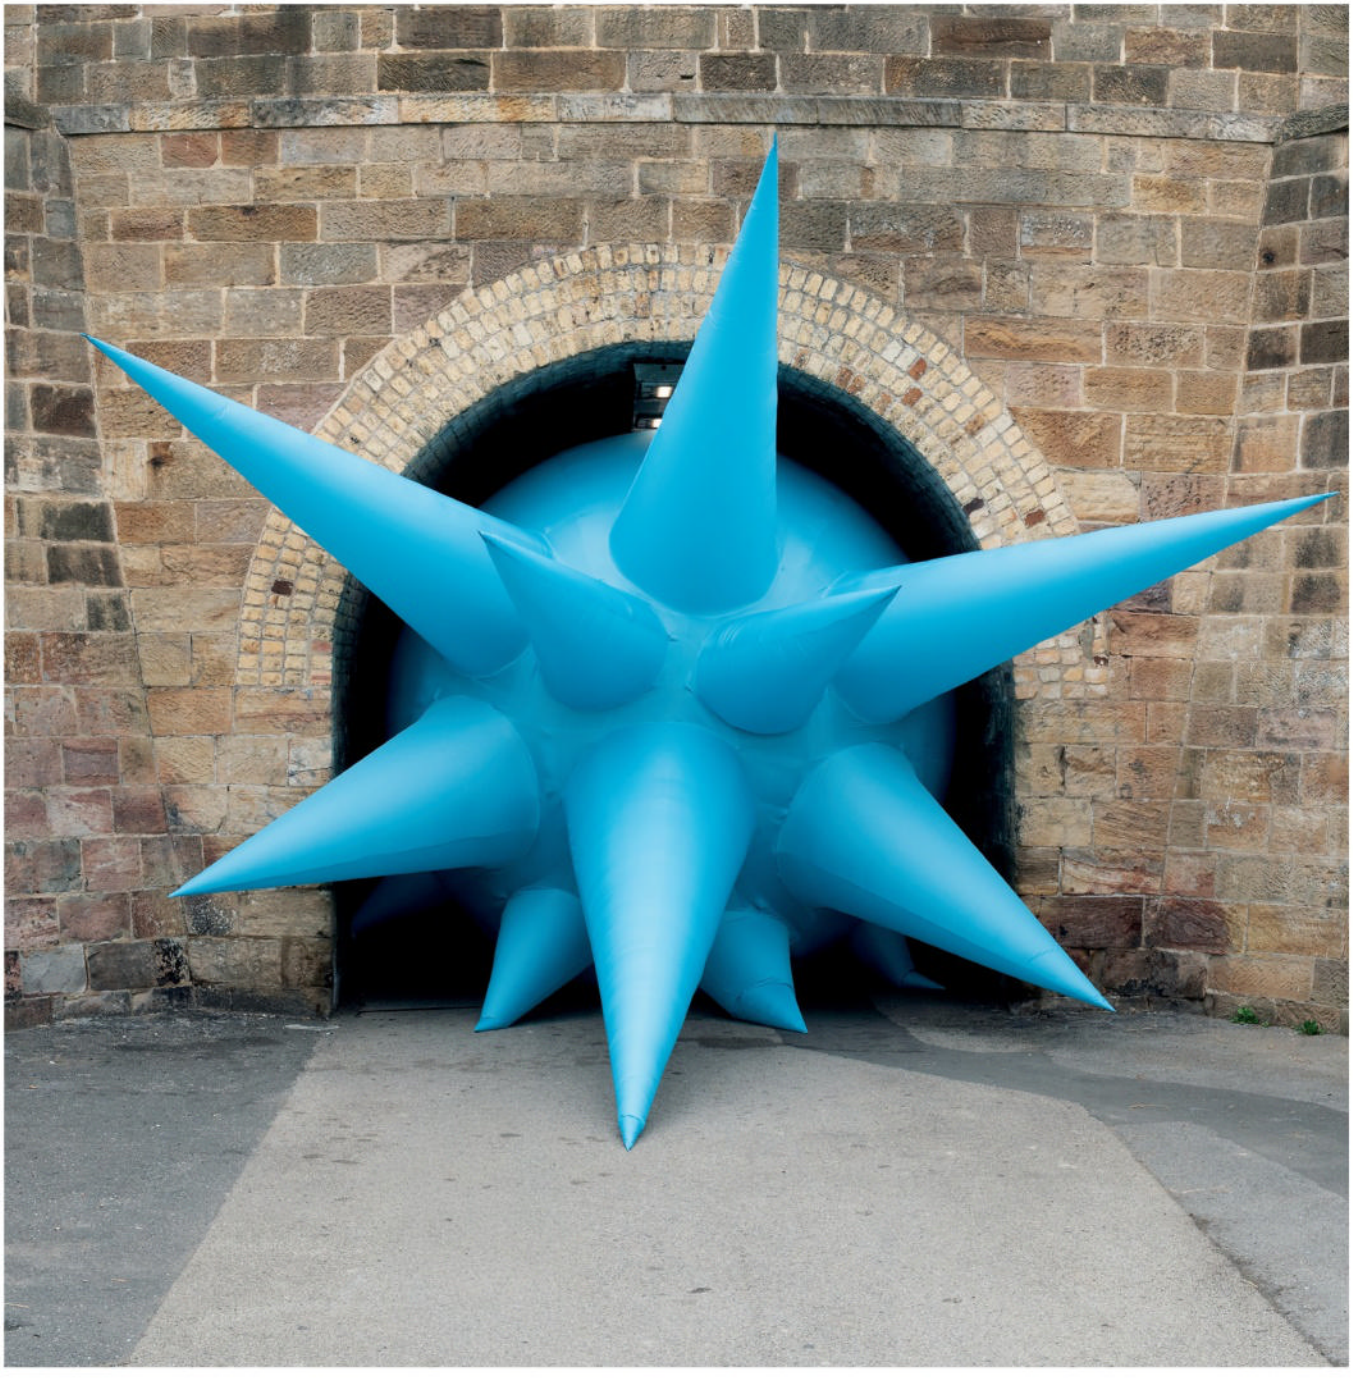

a variety of per-unit costs, including customer acquisition and customer support. When unit economics are ignored, startups often focus on expanding the market for a product that customers love—but do so at a price below what the product costs to produce and sell. Take MoviePass, a company that offered subscribers virtually unlimited movie tickets for a flat monthly fee. It attracted millions of customers. The downside? The firm lost money on nearly every transaction. Despite rapid growth, there was no viable path to profitable unit economics. Customers would pay only so much, and theaters would supply tickets only at a certain price. The company ended up filing for bankruptcy. (It was later bought by one of its cofounders, who substantially reworked the business model.) MoviePass illustrates what Warren Buffett had in mind when he wrote, "The worst sort of business is one that grows rapidly, requires significant capital to engender the growth, and then earns little or no money."

In their defense the executives and investors behind MoviePass may have assumed that they could fix poor unit economics later—by raising prices or cutting costly services. But this "grow and then gouge" strategy rarely pans out, because product-market fit and unit economics are deeply intertwined. We can think of very few startups that were able to substantially raise prices or redesign a costly business model after they began to grow. Caught up in the thrill of rapidly rising sales, founders may not even have a plan for fixing unit economics. Backed by venture investors who are willing to fund loss-producing companies, many forget about business fundamentals.

Founders also sometimes assume that network effects (when a product or service becomes more valuable to its users as the number of customers using it increases) will save a growth-first business. Once the business grows very large—the thinking goes—it will become an essential service. While that logic is alluring and might work for some companies, it often rests on a misunderstanding of unit economics. Consider a peer-to-peer marketplace for kids' secondhand clothing that we studied. Parents loved its ease and convenience. Not surprisingly, people kept joining it, and powerful network effects drove increasingly attractive revenue growth. But its founders forgot about their high unit costs for each transaction, and parents simply would

We can think of very few startups that were able to substantially raise prices or redesign a costly business model after they began to grow.

Founders are taught—correctly—to run rigorous experiments to size up demand, pricing, and willingness to pay. What's missing is the same rigor on the cost side.

not pay more for children's clothing. Instead of solving the startup's challenges, the powerful network effects just accelerated its losses. As one investor told us, "Growth solves many problems, but unit economics isn't one of them."

To be sure, there are examples of loss-producing startups that have gone on to financial success: Amazon and Airbnb burned through enormous sums before becoming profitable. Yet such examples miss the critical distinction between unit economics and firm-level profitability. Amazon is legendary for its close attention to unit economics, especially costs. Transactions in its initial core retail business were efficient and simple, and its prices covered the costs of such elements as payments, warehousing, and shipping. Amazon chose to be unprofitable at the firm level in order to invest in new businesses like Amazon Web Services. Over the years the company deployed teams of economists to evaluate the unit economics of potential new businesses. Other companies have followed this playbook too, sacrificing firm-level profitability by reinvesting profits, especially in fast-growing markets, but retaining a strong foundation with a vigilant focus on unit economics.

THE FOUNDATION-FIRST STARTUP

So what does a foundation-first startup look like? An excellent example is Rent the Runway. This firm was launched by Jennifer Hyman and Jennifer Fleiss to rent out designer dresses for special events like weddings, galas, and proms. Early discussions with fashion designers helped Hyman and Fleiss assess the feasibility of securing wholesale pricing, which substantially reduced their unit costs. In parallel they also ran a series of experiments to test demand and pricing, obtaining 100 dresses and holding rental pop-ups targeting sororities and social clubs. Those moves helped them understand the drivers of economics like unit revenues (product-market fit, willingness to pay) and unit costs (purchasing dresses, cleaning costs, the number of times each outfit could be rented before replacement). As they slowly rolled out their service to early users, they gained further insight into customer acquisition costs. All of this meant that when Rent the Runway was ready to accelerate its growth, its business model was already on the path to unit profitability.

Another telling example is Starbucks. Rather than rushing into rapid expansion early on, Howard Schultz and his team meticulously analyzed the potential profitability of each possible store location before moving forward. In other words, unit profitability mattered. As Schultz noted in his memoir, "Starbucks' ability to build and operate profitable stores had succeeded for years because we had adhered to a simple yet ambitious economic model: a sales-to-investment ratio of 2 to 1. During a Starbucks store's first year in business, it needed to bring in $2 for every $1 invested to build it." While that exact ratio is specific to the company's business model and industry, the disciplined approach to unit profitability helped Starbucks avoid the trap of geographic expansion at the expense of store-level economics.

In sum, unit economics are not just metrics—they're a discipline. Founders who put this discipline into practice follow three major tenets: They experiment broadly, simplify ruthlessly, and deliberately delay growth.

EXPERIMENT BROADLY TO UNDERSTAND DEEPLY

Founders are taught—correctly—to run rigorous experiments to size up demand, pricing, and willingness to pay. What's missing is the same rigor on the cost side. The foundation-first approach treats cost structure not as something to optimize later but as a set of hypotheses to be explored early on, alongside revenue potential.

In practice this means conducting tests across the entire business model. Which activities truly require human labor, and which can be simplified, automated, or delegated to artificial intelligence? Which design choices quietly lock in high variable costs? And which cost drivers improve—or deteriorate—as volume increases? Most founders already know they need to run tests to assess minimum viable products and potential pivots, but their experimentation is focused almost exclusively on demand and growth. Foundation-first entrepreneurs instead probe both revenue and cost drivers with an eye on unit economics. This broader approach is slower and more laborious, but it typically yields a simpler, more profitable business model with real product-market fit and a clearer strategy. By experimenting across many elements of the

Is Foundation First Always the Right Choice?

The foundation-first approach works particularly well for companies selling physical products and for business models with significant variable costs (for example, inventory) and operational complexities (for example, logistics) that make unit economics challenging. Yet even products with low variable costs and little operational complexity still require learning about product-market fit and often coping with some major variable costs (for example, customer acquisition) that make unit economics relevant. Building the foundation first also works for business models with potential network effects and economies of scale. Why? Network effects cannot compensate for negative unit profitability, while the timing and size of scale economies can be difficult to predict. Rather than aiming for illusory first-mover advantages, developing a simpler, better business model with real product-market fit usually makes more sense, as successful late movers like Google in search and Apple in smartphones illustrate.

Some point to startups backed by investors willing to subsidize massive, prolonged losses as proof that the foundation-first approach isn't universal. But these startups are like adult kids living with their parents: The rules feel different until the parents cut them off.

Deep tech startups, however, do have different rules. These ventures make long-term and expensive bets on creating pathbreaking technologies that often address societal issues like climate change or that open entire new industries like space exploration. Raising massive funds from many sources—governments, corporate strategic investors, venture capitalists, sovereign wealth funds, and even philanthropists—is usually essential. Working with governments—as regulators, customers, and investors—also is often key. But the foundation-first approach will come into play when these startups transition from scientific discovery to commercial products.

business—including customer acquisition channels, service levels, operations, and target segments—founders uncover not only what works but also how different levers interact. These interdependencies between revenues and costs are critical: They reveal which changes improve unit economics and which quietly erode them. Once entrepreneurs get these insights right, scaling up becomes less about discovery and more about replicating a well-designed model.

Consider a prosperous online fashion venture we'll call by the pseudonym StyleCo. The founders began with a vision of bundling personal styling services with e-commerce. One of them described the first year simply: "It was just about proving this works. Do the economics work?" The founders tested whether algorithmic tools could replace some human stylists, lowering costs. They tried different brands and discovered that boutique labels, which could be bought at wholesale prices, still resonated with customers. They experimented with an up-front styling fee, credited toward any purchase, which both reduced theft and unexpectedly attracted higher-value customers. They even learned that full retail pricing did not deter buyers who valued convenience—but did attract more-desirable suppliers. Through this series of experiments, StyleCo discovered powerful interdependencies and reshaped its business model. Growth then became a matter of repeating what already worked and culminated in a successful initial public offering.

A near twin, FashionCo, launched around the same time with a similar concept and price point but took a growth-first path. The founders aggressively pursued sales and poured money into customer acquisition. They quickly found a product-market fit, but almost all their experimentation was aimed at boosting revenues. As the business grew, they layered on expensive features—such as commission-based stylists and a brick-and-mortar channel—without testing lower-cost alternatives or understanding how their choices affected overall economics. As one of them later admitted, "We didn't know if we could ever bring down our costs, but our customers loved their shopping experience." Revenues rose, but costs rose just as fast. FashionCo's weak unit economics never recovered, and the business model ultimately failed.

SIMPLIFY TO SCALE

Broad experimentation also pushes ventures toward simplification, not just toward improved unit economics, product-market fit, and strategic clarity. Simpler business models are easier and cheaper to scale, have fewer failure points, and are more intuitive for customers and employees alike. Experimenting widely forces entrepreneurs to stay close to the details of their business long enough to become genuine experts—able to see patterns, think strategically, and strip the model down to what really creates value.

Instagram is a vivid example. It began as Burbn, an overloaded app combining check-ins, photo sharing, and rewards. Through experimentation, the founders realized that while everything "worked," photo sharing was what

Foundation First in the Age of AI

Artificial intelligence (particularly generative AI) creates new urgency for foundation-first thinking. Consider a major U.S. service provider that recently explored using AI to handle customer service interactions. Rather than racing to deploy the technology, it took a foundation-first approach—running careful experiments to understand both unit economics and service quality before inserting AI into the workflow and then scaling.

The experiments revealed unexpected insights. For high-reliability tasks like offering customers concrete options to choose from, traditional code outperformed AI. But with conversational elements, gen AI excelled. Interweaving both delivered the best experience and enabled rapid responses when call volumes spiked. And while AI was expensive per interaction, it was still substantially cheaper than human representatives. Only after establishing these unit economics and assessing the customer experience did the company begin exploring broader deployment.

Several AI dynamics make building a foundation first especially critical:

Unit costs accumulate fast. Large language models charge for every word (or more accurately, token) processed and generated. Video, images, and especially agents become exponentially more expensive. These costs compound quickly across millions of interactions.

Broad experimentation reveals crucial trade-offs. For any AI use case, critical choices—which model to use, how to structure relevant context for AI, when to combine AI with traditional code or humans—have interdependent effects on quality, speed, and costs. Those effects are hard to unpack without experimentation.

Simplification is powerful—and often necessary. AI performance can degrade quickly with complexity. If instructions for AI are too long or too many tools are available, so-called "context rot" happens. AI agents lose focus, get confused, and often fail to follow instructions. Complexity also harms humans in the loop who are likely to be overwhelmed by AI output and miss where their judgment actually matters.

The message is clear: AI's complexity, power, and cost structure make it especially foolhardy to chase growth before understanding unit economics, interdependencies, and usage patterns. The transformative potential of AI makes slowing down feel risky, but scaling a fundamentally unprofitable business model is far riskier.

users truly cared about. They ruthlessly cut the rest and then simplified the core action to just three taps. That clarity of focus turned a cluttered product into a breakout hit.

Businesspeople tend to solve problems by adding features, channels, and options, which is exactly what we saw happen at FashionCo. A foundation-first approach counteracts that bias by making subtraction part of the design. StyleCo did this when it experimented with the number of items to ship to each customer. The initial plan called for 10 pieces, but testing revealed that five performed better. Halving the number of items reduced shipping and inventory costs and, counterintuitively, increased sales—customers bought more when choices felt curated rather than overwhelming.

A key question is: When should entrepreneurs stop broad experimentation? In our experience startups benefit from it until they develop a relatively complete business model with promising unit economics, including product-market fit. That can require six to 12 additional months after the first indication of a favorable opportunity.

GO SLOW TO GO FAR

Perhaps the most counterintuitive aspect of building for long-term, profitable growth is actively delaying it. This means deliberately holding back on adding revenues and customers while broad experimentation continues. The reason is simple: Growth creates immediate operational demands—hiring, infrastructure, customer support—that soak up attention and capital. It's far easier to refine unit economics when you're not simultaneously scrambling to build capacity. Staying small a bit longer preserves the ability to spot, implement, and unwind changes before complexity and customer expectations harden into constraints.

One tactic for delaying growth is to throttle customer acquisition—limiting sales, hiring, and marketing spending

Staying small a bit longer preserves the ability to spot, implement, and unwind changes before complexity and customer expectations harden into constraints.

to "just enough" to get traction and validate product-market fit, while freeing time and money to experiment with both revenues and costs. Klaviyo, a customer data and AI marketing platform, did exactly this. Its two cofounders iterated for three years before raising significant capital. One focused on the product, the other on early customers. In the first year alone, they pivoted from an onboarding tool to a data-integration platform to an email-marketing platform. Even as sales picked up, they intentionally avoided building a sales team, keeping customer acquisition modest. That restraint gave them space to test and tune their business model and unit economics. Only once the economics—including customer acquisition costs—were highly attractive did they lean into growth and raise major funding, ultimately going public in 2023.

A second tactic is using waiting lists, which make sense when demand runs ahead of capacity. Midjourney, the AI image-generation company, used a closed waiting list even as interest in its offerings surged. By capping the number of users, the team could concentrate on improving the offering rather than on growth-induced problems. As a result, the venture never needed to raise outside capital, despite building a business now believed to generate hundreds of millions in annual revenues.

Some founders go further. Notion, a $10 billion productivity platform, initially raised about $2 million to create a no-code app-building product. After burning through much of that cash and nearing failure, the founders chose not to chase more customers or funding. Instead, they laid off

The gospel of "growth at all costs" has dominated startup thinking for the past two decades. It may have worked for a few outlier successes, but for too many others it's a trap.

most of the team and completely rebuilt the product using foundation-first principles. That pause let them reimagine Notion as an intuitive workspace for documents, databases, and projects. By delaying growth and returning to broad experimentation, the founders created a product that now serves more than 100 million users.

Delaying growth, of course, runs against the grain of many investors. Most are happy to fund unprofitable start-ups—until, abruptly, they aren't. A deliberate pause can complicate early efforts to raise capital, especially Series A funding, because it demands investors who value discipline over top-line growth and founders who can stretch cash. But it can also reduce the total capital required from investors. Above all, actively delaying growth raises the odds that when founders do scale a business, they're building on a solid foundation rather than speeding toward a cliff.

VENTURES BEYOND START-UPS

Regardless of organizational context, the foundation-first approach can help new enterprises succeed. After all, small and medium-size businesses (SMBs) launching new offerings and large companies incubating ventures face the same fundamental risk as startups: scaling before unit economics are understood.

At large companies, corporate initiatives often are justified by potential synergies with the core business and enjoy enviable advantages—existing customers, a sales force, brand recognition, and a parent willing to absorb early losses. Yet many still fail—for the same reasons that startups do: weak unit economics and fuzzy product-market fit, compounded by wishful thinking about synergies.

A foundation-first approach forces corporate entrepreneurs to validate those synergies instead of just assuming that they exist. Which assets actually lower customer acquisition costs? Where does brand trust really translate into pricing power? Which internal capabilities truly improve risk or operational efficiency? Answering those questions early, before scaling, can prevent billions in write-offs.

Goldman Sachs's consumer venture Marcus is a case in point. It began with a strong thesis: Use Goldman's risk-management expertise to offer unsecured personal loans, then expand into credit cards and high-yield savings. The product resonated with customers, and the business grew quickly. But on the personal-lending side, the anticipated synergies never translated into solid unit economics. Marcus's loss rate for these products remained far higher than that of its peers, and the business ultimately racked up roughly $3 billion in losses before Goldman pulled back on personal lending, keeping only the savings business. Marcus shows how even within a world-class institution, growth without proven unit profitability is fragile.

A foundation-first approach can also propel a small business into a highly profitable medium-size business—what is termed the "mighty middle." If a company lacks the high revenue potential that attracts VC investors, strong unit economics are necessary not just for survival but to generate cash to invest in growth. For SMBs, the danger is often not rapid growth around poor unit economics—but settling into poor unit economics that never enable growth in the first place.

THE GOSPEL OF "growth at all costs" has dominated startup thinking for the past two decades. Our research suggests that it's time to reconsider this orthodoxy. The foundation-first model—focusing on unit profitability, broad experimentation, and delayed growth—offers a more reliable path for most ventures, whether they're VC-backed startups, new corporate businesses, or SMBs. Pursuing growth at all costs may have worked for a few outlier successes, but for too many others it's a trap. The future of entrepreneurship lies in the recognition that the key to long-term, profitable growth isn't chasing growth. In fact, it's just the opposite—intentionally slowing down, building the right foundation, and only then scaling aggressively. ■

HBR Reprint R2606F

BENJAMIN HALLEN *is the Hughes M. Blake Endowed Professor of Business Administration at the University of Washington's Foster School of Business.* **RON TIDHAR** *is a member of technical staff at OpenAI, where he leads analytics for consumer subscriptions.* **KATHLEEN EISENHARDT** *is the Stanford W. Ascherman M.D. Professor and the codirector of the Stanford Technology Ventures Program at Stanford University (emerita) and a coauthor of* Simple Rules: How to Thrive in a Complex World (*Houghton Mifflin Harcourt, 2015).*

BOARDS

Building a BOARD That Fits

A good one is designed for purpose, people, process, and leadership in a specific context.

AUTHORS

Timothy J. Rowley
Professor, Rotman School of Management

Stanislav Shekshnia
Senior affiliate professor, INSEAD

ARTIST **TODD ST. JOHN**

ACROSS BOARDS AT ORGANIZATIONS of every size, ownership structure, and stage of development, two design flaws occur with striking consistency. First, boards are often designed to reinforce the instincts and preferences of the people who control a company—typically founders or CEOs. Second, even when boards are effective, over time they fail to evolve with the company's needs.

The remedy lies in understanding that good governance is about creating a counterbalance to the forces that shape decision-making in a company, which are specific to a firm's ownership structure and context. In this article we'll describe the four kinds of contexts in which boards—regardless of industry—most frequently operate and provide a framework for building boards that fit each one.

Our categorization and framework are based on a collective five decades of research and teaching on boards and governance. In particular, we've drawn on the cases discovered and documented through our Leading from the Chair and International Directors programs at INSEAD—and on the experiences that the more than 3,500 directors and executives from over 2,000 companies that have participated in those programs have shared. In the discussion that follows, we'll focus on four of those companies, using them to illustrate the risks and opportunities of each context.

HOW CONTEXT SHAPES A BOARD'S CHALLENGES

Most firms fall into one of four ownership categories, each of which presents distinctive strategic advantages and disadvantages.

Entrepreneurial ventures have genuine strengths: concentrated decision-making, a long-term ownership mindset,

Good governance is about creating a counterbalance to the forces that shape decision-making in a company, which are specific to a firm's ownership structure and context.

IDEA IN BRIEF

THE PROBLEM

Most boards are designed around the preferences of founders, CEOs, or controlling shareholders and gradually become misaligned with the organization's changing needs.

THE SOLUTION

Design boards to counterbalance the specific risks created by the company's ownership structure and context, using a framework built around purpose, people, process, and the chair.

THE BENEFIT

Boards become more independent and strategically valuable and are better equipped to challenge management, adapt over time, and improve long-term governance and performance.

Too many public company boards spend the majority of their time on audit, regulatory, and disclosure matters at the expense of the decisions that determine long-term value.

and the kind of conviction that drives bold strategic bets. But they often suffer from an unchecked appetite for risk, a narrow strategic perspective, and the inability to distinguish what's good for the company from what's good for the founder.

With these ventures, the board's counterbalancing role is to ensure that founders' convictions are tested against independent judgment from people with broader experience and perspectives. This is especially critical for companies approaching an initial public offering. In 2019, for example, WeWork's IPO collapsed largely because its board wasn't designed to challenge the company's cofounder. Adam Neumann held super-voting shares, was his own board chair, and had surrounded himself with directors who were personally close to him or had a professional affiliation with the company. As a result, they were more likely to ratify than to question his decisions.

Established *family firms* face governance challenges that are in some respects the most complicated. Family members may simultaneously sit on the board, be major shareholders, and hold management positions. But a family shareholder who's also a director and runs a business unit faces an almost impossible task in separating those roles cleanly, and the board as a whole must oversee an entity in which the distinctions that make governance work—between owners who set the direction, directors who oversee, and managers who execute—are blurred. Sibling rivalries, succession tensions, and disagreements about the balance between family employment and professional management add further complexity.

The counterbalancing role of a family company's board is to set boundaries firmly: to separate family governance from enterprise governance, to ensure that shareholder interests are represented with a coherent and unified voice rather than fractured by competing family agendas, and to maintain strategic focus when family dynamics threaten to pull the agenda toward questions about the distribution of profits and assets to the family rather than external competitive threats.

Large, mature *public firms* face a governance environment characterized by dispersed ownership, intense regulatory scrutiny, and the relentless pressure of quarterly performance cycles. This creates a complex counterbalancing task: The board must buffer the company against short-term market expectations and oversee the compliance machinery of public governance while making sure it doesn't crowd out genuine strategic oversight. In too many public company boardrooms, the directors spend the majority of their time on audit, regulatory, and disclosure matters at the expense of the decisions that most directly determine long-term value.

Private-equity-owned firms engineer their governance to provide disciplined oversight of a defined value-creation plan within a specific investment horizon. Their boards tend to be more structured, more performance-oriented, and more focused on execution metrics. Though these are improvements over the informal governance that often precedes a PE acquisition, the tools of oversight—the hundred-day plan, the bonuses linked to key performance indicators, the quarterly performance cadence—can leave little room for the coaching relationships, candid conversations, and longer-range strategic thinking that leadership development requires. More fundamentally, there's an inherent tension between the board's duty to the company and its duty to the fund: Decisions that optimize exit timing and returns to existing investors may not optimize long-term enterprise value for future investors.

THE NEED TO EVOLVE

Over time, many good boards lose their ability to serve as an effective counterbalance. This is partly because the things that make a board stable and functional—trust among directors, settled processes, shared assumptions about roles and responsibilities—also make it resistant to change. The gap between what a board was designed for and what the organization needs usually widens gradually, and the board may be unaware that this is happening until a crisis strikes.

Few examples illustrate this dynamic more vividly than Eastman Kodak. Although some say the company was taken by surprise by digital photography, that's not entirely true: In fact, its engineers developed the first functional digital camera in 1975, giving it an early and commanding technological lead. But in the 1990s Kodak's board was constructed for the company it had been: a dominant, vertically integrated manufacturer of chemical film products, competing on precision engineering and production scale. Its external

directors came from heavy manufacturing, traditional financial services, and consumer packaged goods. They were accomplished leaders, but their collective experience mapped to industries that competed on engineering capabilities, capital intensity, and brand loyalty in mature markets. The trouble was that the digital photography market would be won by whoever built the most compelling ecosystem of devices, software, and sharing platforms—a competitive logic foreign to the board's frame of reference.

Two structural features made it difficult for Kodak's board to expand that frame. The persistent combination of the chairman and CEO roles—held by successive leaders—eliminated the very counterbalance a board is meant to provide. To craft strong strategies in moments of disruption, boards require friction and genuine challenges; a chairman who is also the CEO has an inherent interest in defending the current strategy rather than interrogating it. Kodak's board lacked the most important mechanism for testing whether the instinct to protect the film business should yield to the harder choice of cannibalizing it. The second feature was a structure that divided directors into three staggered classes with three-year terms. Although this arrangement was intended to ensure turnover of board members, it imposed a rate of change in board composition that was too slow for a rapidly shifting competitive environment.

To be sure, Kodak's board, like many others, engaged in systematic board and director evaluations, which can surface characteristics misaligned with current needs. But in our experience, traditional evaluations are limited. Even when they ask forward-looking questions, the answers are filtered through the existing board's assumptions, relationships, and collective frame of reference. What remain outside the frame are the broader possibilities: whether strategy is governed by the right logic, whether the board's size and structure are suited to the actual complexity of current challenges, and whether standing committees map to priorities that still matter.

To put together a board that's effective over the long term, you need to focus on four key interconnected elements: *purpose, people, process,* and *the chair.* These are interdependent variables—changing one affects all the others. We'll now walk you through them in more detail, using four disguised company examples.

PURPOSE

DEFINE THE BOARD'S MANDATE

Purpose is about why the board exists and whom it owes its duty to. At a founder-led company its purpose might be "to provide independent oversight and check concentrated founder control." At a PE-backed company it might be "to maximize enterprise value through disciplined capital allocation while building company capabilities beyond the investment horizon."

At Wright Technologies, a midsize industrial company with revenues of €100 million, George, the founder and sole shareholder, worked with a newly appointed chair to formulate the board's purpose: "ensuring the long-term sustainable development of the company and value creation for the shareholder by creating an effective framework for executive action." The phrasing made it clear that the board would serve both the founder and the company—not one or the other. It would create frameworks, not make operational decisions. And it would focus on sustainability and value creation, not just growth and revenues. That clarity shaped everything downstream: The board would concentrate on leadership oversight and development and long-term strategy—and especially on risk management and corporate governance. It explicitly chose not to involve itself in quarterly performance management or product development decisions.

When Hanz became the nonexecutive chair of the board of Empire Group, a 200-year-old family-owned company, it had more than 500 shareholders (all family members), who didn't always think about the future of the business. Hanz (a seventh-generation family owner) recognized that protecting senior executives from their intervention was essential for the long-term stability of the family and the business. Together with other family members who were also directors, he defined two main roles for the board. First, it should ensure that shareholders wouldn't interfere with the management of the business. Second, it would define a long-term vision and determine the sectors and businesses the company should be in. For example, the board decided that the company would avoid capital- and technology-intensive industries and businesses outside Europe.

Getting purpose right is critical in the context of an IPO, which requires transforming a private company into one for

the public market. The board members will have to counterbalance three forces—the existing management's control over the information they receive, the existing owners' immediate interests, and the public market's short-term performance pressure—all while managing a radically larger compliance burden. When Xama, a $10 billion company, was spun off from its parent company and started preparing for an IPO, the new board defined its purpose as "supporting the executive team through the transition and establishing governance practices without disrupting operations." That meant creating a governance structure appropriate for the stock market environment and strengthening management's ability to operate in that environment through oversight and mentoring.

In contrast to the other firms, PE-owned GeoTech never explicitly defined a purpose for its board. A PE fund had acquired the lead position in the company with the aim of getting a return through a sale in less than four years—and that timetable shaped what the board focused on. But when the owner's exit is the goal, a board will generally reflect the owner's priorities more than the company's needs. Of course, the two often align, since disciplined growth serves both. They can diverge, however, when a board decides to trade a near-term financial gain favoring the owner's exit value for a longer-term and uncertain strategic opportunity. GeoTech may face that dilemma in a pending move into a new sales vertical, the largest market for its product, where establishing an early foothold could make the company the category standard. But the move would require pricing GeoTech's products aggressively, complicating the story of steady growth and strong margins that would make the company a safe bet for a trade buyer. A carefully defined purpose would have named both duties, to the fund and to the enterprise, giving the board more room to weigh them against each other.

PEOPLE

CREATE A CAPABILITY MAP

People is the element that most board designers invest their energy in—and where their instincts fall short. The common mistake is assembling a team of star directors rather than a star team.

What Are You Designing Your Board Meeting to Do?

The right meeting processes can have a huge impact on a board's focus and effectiveness. With the standard approach, directors sit through more hours of presentations than discussions during a one-day session. They are informed but don't deliberate; they simply ratify management decisions.

Traditional board meeting structure

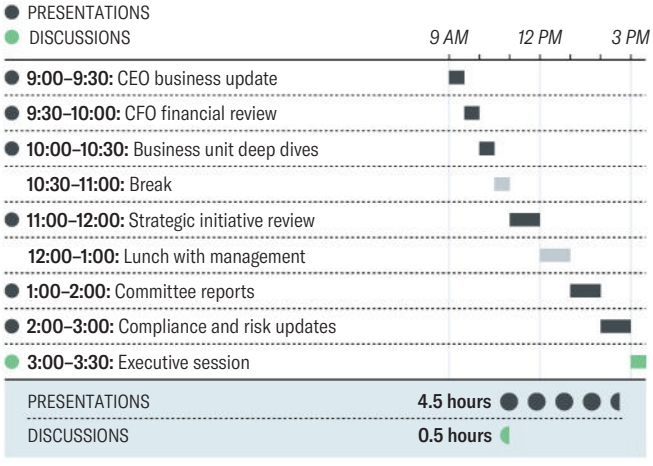

In contrast, allocating minimal time for presentations and more for deliberation and decision-making improves the odds that boards will add counterbalancing value through the exercise of independent judgment.

Recommended board meeting structure

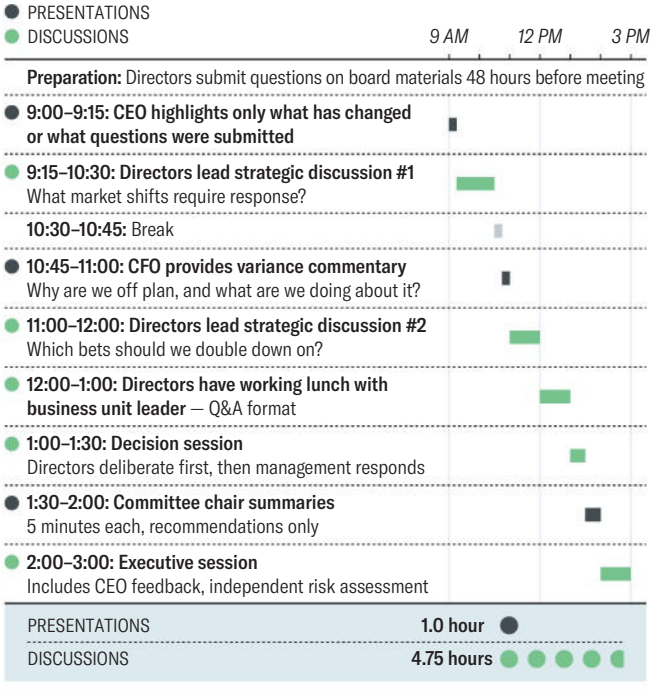

Boards that rely solely on the materials management has curated are poorly positioned to form independent views. Effective boards cultivate other sources of information.

The skills a board needs most vary with a company's ownership structure: A founder-led firm needs directors who offer independent judgment and strategic and financial discipline. A family-owned company needs board members with strong strategic management skills, deep knowledge of governance, and the ability to withstand emotional pressure. A PE-backed company needs directors who can provide mentoring and a longer-term strategic perspective. A publicly traded company needs directors with deep compliance and governance discipline. The point is not that these capabilities are unimportant elsewhere—it is that each context has a primary requirement.

To rein in his entrepreneurial exuberance and instill strategic discipline at Wright Technologies, George recruited an experienced independent board chair who would moderate discussions and challenge him, not give him advice. Together they added a director who was a former senior international executive at one of the company's largest potential customers. That director strengthened the board's independence and global perspective and could help mentor the management team. Two additional independent directors supplied financial, risk management, and compliance expertise. Domain knowledge was deliberately left where it was strongest: With the founder himself on the board and in management with an incoming CEO, a 20-year industry veteran, who sat outside the board. The result was a board designed not to mirror the founder's strengths but to provide knowledge he lacked and counterbalance his power.

At family-owned Empire, the board's nomination committee defines profiles of directors in collaboration with the chair and external advisers. Directors must be credible business leaders with strong strategic skills, the ability to provide oversight across functions, and a deep understanding of global politics. The most important qualities, however, are the capacity to form independent views, the courage to back those views, and the ability to work collaboratively.

Xama needed a board that could guide management through the spin-off while meeting the regulatory demands of public markets. The newly appointed chair and parent company executives translated those requirements into detailed profiles: industry understanding; deep experience with finance, audits, and risk management; legal and governance expertise; international leadership experience; and supply chain and technology security expertise. In terms of personal characteristics, they looked for collaborative ability and the willingness to engage in debate and find solutions to challenges. Every candidate was told the required commitment up front—12 days a year for in-person committee and board meetings plus travel, and a further 25 days for online meetings and preparation—and candidates holding more than three board positions were excluded from consideration. The result was a board composed not for the company Xama had been as a subsidiary but for the standalone public entity it needed to become.

In contrast, GeoTech's board was shaped less by a map of the capabilities the PE-backed company needed than by ownership stakes and shareholder agreements. The lead investor's partner and several investors from early rounds of financing made up the majority of its seven directors, with the CEO also on the board. That gave the board genuine strengths: close alignment with the investors, strong financial oversight, and valuable network connections across relevant industries. The directors were fluent in talent reviews, pricing, sales pipelines, and the allocation of resources across sales verticals. But the board was generally reluctant to spend time debating complex strategic choices requiring in-depth discussions. A board designed around the company's needs as well as the owners' might also have included directors who had expanded a product offering into a competitor-dominated market and could help guide the board, not just management, through uncertain bets.

PROCESS

SET UP AN OPERATING SYSTEM

Purpose defines the board's focus and goals. People define its capabilities. Process determines how those capabilities are organized and directed to achieve the purpose—how directors access information, interact, make decisions, and assess whether the board is fulfilling its mandate.

Information flow. In every ownership context, this is controlled by the company's management, but boards that rely solely on the materials management has curated are poorly positioned to form independent views. Effective boards deliberately cultivate other sources of information: direct access to customers, independent market intelligence,

site visits, external advisers, and conversations with employees below the executive team. This isn't a sign of distrust; it's the best way to support management and fulfill the board's counterbalancing role.

At Wright Technologies, the board meets with customers and vendors and receives regular market updates from an independent consulting firm. At Empire, Hanz invites the directors to workshops with experts in sustainability, risk management, cybersecurity, digitalization, and global affairs, and strategy consultants and investment bankers regularly speak to them at informal dinners. Xama's chair brings analysts into the boardroom to share their perspectives and takes board members to production sites and industry conferences.

Careful process design takes into account what the information architecture risks crowding out. In public companies approaching or following an IPO, for instance, compliance activities often leave little room for strategic information. The priority is to protect space for forward-looking intelligence—on competitive dynamics, technology shifts, and emerging risks.

The following tools help boards stay on the right track:

Meeting design and decision protocols. Effective boards structure their agendas to prioritize the most consequential issues rather than the most urgent ones, give directors sufficient time to review meeting preparation materials, surface and welcome dissent, and clearly record what they've decided and what authority has been delegated.

At Empire, the board follows a three-stage process for major decisions. At the information stage the management shares initial ideas and gathers questions; gaps in information are perfectly acceptable. At the discussion stage the management answers most questions and presents clear solutions. At the third stage the board makes the decision in private sessions. The CEO, who is not a board member, attends most meetings as an observer.

At Xama the chair receives preparation materials two weeks before a board meeting, reviews them, and provides feedback to management on what to improve before sending them to the directors at least seven calendar days before a meeting. Presentations are limited to no more than 20 slides, include a one-page executive summary, and specify whether they're meant to inform the directors, prompt a discussion, or generate a decision. During the meetings, the CEO makes only short introductory presentations to emphasize key points, and 70% to 80% of the time is devoted to Q&A and deliberation. The chair also discourages private meetings between the CEO and the directors immediately ahead of the meeting, invites senior executives only if they need to cover a specific subject, and constantly reminds the directors of the need for independent judgment.

GeoTech's board process closely reflected its lead investor's priorities—speed, control, and execution. Meetings were dominated by management presentations; directors' questions went mostly to the CEO, and when differences of opinion arose, the lead investor would sometimes suggest taking them offline—so some decisions were shaped outside the meeting. The board has been moving to correct this, however. A strategy discussion now opens every meeting, presentation time has been reduced, and outside experts present to the full board rather than to management alone—changes that came about to some extent because directors wanted to be an active part of the process rather than passively absorb information. The pending decision on the new vertical will test how much that shift has taken hold: whether it is worked through by the full board or settled mainly between the lead investor, the CEO, and the outside expert the investor has offered to bring in.

All three boards hold what are confusingly called executive sessions, in which directors meet without management (including the CEO) to allow the directors to consolidate their views and surface sensitive issues. This is one way to create space for independent directors to speak early and often—something that's especially important at founder-led and family-controlled companies, where a single voice often dominates the room. At Wright Technologies independent directors meet for dinner before every board meeting to exchange views without George present, and in meetings the chair gives the floor to George only after other directors have spoken.

Board evaluation. The value of this standard practice depends on what it measures. The most common mistake here is evaluating the board against generic governance checklists rather than the purpose it should serve. The question should always be: Are we still an effective counterbalance—or have we drifted?

A common mistake is evaluating the board against generic governance checklists. The question should always be: Are we still an effective counterbalance—or have we drifted?

The qualities that make other leaders effective are often precisely what make chairs ineffective. Chairs' power comes from restraint, not action.

At the end of each board meeting at Xama the chair conducts evaluations by asking: "How effective was our discussion today? Have we made decisions that will strengthen our company? Have we provided guidance to management? How should we improve next time?" Often directors agree on one change to be tested at the next session, such as "Nobody should speak for more than 10 minutes" or "Committee chairs should send their reports in advance rather than present them in the boardroom." The chair calls every director after each meeting to get additional feedback about the discussions and his performance. Every 12 months the board does a comprehensive self-evaluation, using 360-degree surveys about the board process, interactions with the CEO, and the committees, and the results guide how the board evolves.

At Wright Technologies the directors also conduct similar evaluations at the end of each meeting and do an annual 360-degree survey-based review. Their focus, however, is different: They specifically examine whether they're counterbalancing the founder; bringing in financial, risk, and compliance discipline; and helping management build a foundation for long-term sustainable development.

THE CHAIR

FILL THE BIGGEST NEEDS

As with board composition, the choice of chair should be driven by the specific counterbalancing need. In founder-led companies, the most important consideration is separation between the founder and the chair. At Wright Technologies, George recognized this issue and deliberately recruited a chair with experience in a different industry. This reduced the likelihood that the chair would try to compete with him on industry strategy issues instead of focusing on board processes. In family enterprises like Empire, the chair must be able to set boundaries between family agendas and enterprise governance and command the trust of multiple family factions without favoring any one of them.

In companies approaching an IPO, the chair appointment carries additional weight because it must signal credibility externally and set the right tone internally. Xama selected a chair with nearly four decades of experience who had been a CEO of both public and private companies and had chaired several boards. That depth of experience would reassure investors while providing the governance expertise the company needed as a public entity.

GeoTech also illustrates how much rests on the way the chair's role is defined. Until recently, the CEO, who is also the chair, dominated board meetings—setting the agenda and leading the discussion—while much of the informal authority sat with the lead investor. This meant that the board had no counterbalancing voice on it at all, given how close the lead investor and the CEO were. This is evolving: The CEO now sees that meeting agendas incorporate input from the other directors, a step toward broader engagement.

Once appointed, the first challenge for any new chair is to establish the right relationship with the company's CEO. But the best way to approach it will vary with the company's context. At Xama the new chair was aligned from the get-go with the CEO on how their relationship would work: The chair would be the link between the board and the CEO, not the CEO's boss. Because the family that had founded its parent company still controlled more than 50% of Xama's stock, this separation of the two roles was critical. To ensure the firm's readiness for the IPO during the most intense spin-off months, the chair kept board involvement high—with prospectus reviews, risk oversight, and investor messaging—but not so high it disrupted the business.

At Wright Technologies the chair became a coach and mentor to the new CEO, who had never held the top job before, helping him learn the job and master relationships within a governance structure he had never experienced. At GeoTech the focus has been more on tempering the lead investor's influence. The CEO has often settled questions directly with the lead investor but now is trying to bring more issues into the boardroom for collective deliberation.

The skills that chairs require are largely universal. But because the role of the chair is fundamentally different from every other leadership position, the qualities that make other leaders effective are often precisely what make chairs ineffective. Chairs' power comes from restraint, not action, because their job is to create the conditions for the board to counterbalance effectively. A decisive, execution-focused chair, for instance, could close off critical board deliberation and discussion.

We've identified three essential skills of good chairs. They must be able to:

• **Enable people.** Effective chairs understand the directors well enough to draw each one out productively. They know who must be encouraged to speak up, who dominates too much and needs gentle redirection, and who brings relevant expertise to each agenda item. This requires building relationships outside board meetings. At Wright Technologies, after the directors receive materials for an upcoming board meeting, the chair reaches out to each one to discuss them, solicit questions and concerns, and provide advice and reassurance. She also organizes developmental events for both the whole board and individual directors to upgrade their skills. Xama's chair meets once or twice a year with every board member to exchange feedback, provide encouragement, and surface individual concerns. He gives the floor first to directors who don't have expertise in the subject matter, asks experts to be brief and speak in simple terms, rephrases the comments of directors who don't express themselves clearly, and uses humor and self-deprecation when tensions rise.

• **Establish structure.** This means designing processes that make counterbalancing likely. How the agenda is organized, how discussions are opened, how dissent is received—all are the chair's responsibilities. The chair reserves time for genuine deliberation, ensures that alternatives get heard, strives for real consensus rather than apparent agreement, and makes certain that decisions are clear before moving on.

At Empire's board meetings, Hanz always opens by asking all the directors to share something relevant that they've learned since the last meeting. He gets them each to share what they know about an issue before exploring solutions, carefully frames the discussion questions, and always speaks last. The chair of Wright Technologies always writes the resolution adopted by the board on a whiteboard and asks each director to read it carefully and confirm agreement with it. At GeoTech, whose meetings were designed for speed and efficiency, the structural task is now to set aside deliberate space for open discussion, so that a streamlined process still leaves room for genuine debate.

• **Encourage dissent.** Effective chairs establish norms that promote healthy, productive debate. Directors who fear that pointing out problems will damage relationships simply stop raising concerns, and the result is boards where everyone is cordial and few decisions get genuinely examined. Good chairs prevent this by inviting challenge openly, asking,"What are we missing?" and "Who sees this differently?" They respond with curiosity rather than defensiveness, acknowledge uncertainty openly, and make it clear that if directors are in the room when a decision is made and don't voice concerns about it, they own the decision. They follow up after meetings to express appreciation to directors who've asked hard questions.

Hanz has a rule: Every major board decision at Empire should have at least three alternatives. Likewise, if consensus emerges after only two directors' remarks at a Wright Technologies board meeting, the chair starts asking for alternatives. Before a decision is finalized at Xama, the chair always asks, "What needs to happen for this to fail?" And GeoTech is learning that a board can add real value by introducing more debate among diverse viewpoints, which often surfaces new ideas and spurs innovation.

THE QUESTION FACING all founders, owners, CEOs, and board chairs is not whether their boards are good enough. The real issue is whether their boards are designed to counterbalance the distinctive structural forces that shape their company's decisions today and to evolve when those forces change tomorrow. The answer lies not in recruiting impressive individuals and hoping chemistry develops among them. It emerges from defining a clear purpose specifying what the board exists to counterbalance, selecting members who'll provide capabilities the organization lacks, creating processes that enable genuine deliberation, and establishing the conditions for collective challenge through the chair's leadership. Do all that, and you'll build a board that helps your company thrive. Defer it, and you'll build a board for an organization that no longer exists. ■ **HBR Reprint** R2606G

TIMOTHY J. ROWLEY *is a professor of strategic management and organizations at the University of Toronto's Rotman School of Management and a visiting scholar at INSEAD.* **STANISLAV SHEKSHNIA** *is a senior affiliate professor at INSEAD and the board chair of Technoenergy AG. They cofounded and codirect the Leading from the Chair program at INSEAD.*

STRATEGY

THE NEW PLATFORM PLAYBOOK

Success is no longer about winner-take-all. Here's a guide to making smarter strategic choices.

AUTHOR
Chiara Farronato
Professor, Harvard Business School

PHOTOGRAPHER
THIBAULT DRUTEL

ABOUT THE ART

Thibault Drutel photographs European metro stations, capturing both their architectural rigor and the vibrant movement and energy within them.

FOR MORE THAN TWO DECADES,

digital platforms have largely been built around a powerful idea: Be first, grow fast, tip the market, and let network effects do the rest. A platform that attracts enough users can trigger a flywheel effect, leading one company to dominate—often profitably, sometimes spectacularly so. Companies such as Amazon and Google came to exemplify this model, and for a time almost all the world's most valuable firms were platform businesses—firms that facilitate interactions and exchanges between two or more distinct, interdependent groups. (See "Products to Platforms: Making the Leap," HBR, April 2016.)

That playbook is now breaking down. Many platform leaders are coming to terms with the fact that growth is rarely easy to sustain and that competition often persists even after players achieve large scale. Network effects matter, but they are rarely strong enough on their own to deliver winner-take-all or winner-take-most outcomes.

At the same time, the broader context has changed. Investors are less willing to subsidize losses indefinitely, and expectations around profitability have become more demanding. Even when platforms reach millions of users, their growth is increasingly shaped, and often constrained, by regulation, privacy rules, and antitrust scrutiny.

The result is a new era for digital platforms, one in which companies must rethink how they grow, compete, and make money. This is an exciting time for platform businesses because there's more room for competition and innovation. Markets that once seemed to have unassailable leaders are becoming more contestable, opening the door to new

Many platform leaders are coming to terms with the fact that growth is rarely easy to sustain and that competition often persists even after players achieve large scale.

IDEA IN BRIEF

THE PROBLEM

Strategists have long assumed that network effects inevitably produce winner-take-all markets. But regulation, investor pressure, and persistent competition have made the platform playbook far less reliable.

THE SOLUTION

Companies should shift their focus from pursuing dominance to creating durable value. That means prioritizing the market side that generates the greatest benefits for other users, increasing value for existing users through better market design, expanding monetization beyond transaction or subscription fees, and treating trust and customer welfare as strategic assets rather than as compliance obligations.

THE PAYOFF

Platforms that emphasize disciplined growth, differentiated user experiences, and long-term trust can build resilient businesses in an era defined by new rules and realities.

entrants, new business models, and new forms of value creation. One example is regional ride-hailing platforms, which continue to operate successfully alongside global players like Uber. Another is Airbnb, which has expanded the accommodations market without displacing hotels, while hotel companies have found new ways to create and capture value through retail media and other digital services.

I have studied platform growth and regulation for more than a decade, and across multiple settings I've found that the realities of platform competition differ from many of the assumptions that shaped the early playbook. The lessons that have emerged are nuanced: The strength of network effects is often context-dependent; different users value different platform experiences, not just the benefits that come with size; and growth strategies that are compelling in theory do not always produce winner-take-all outcomes in practice.

This article draws on those lessons to expand the strategic lens for platform leaders. It explains why network effects are rarely sufficient on their own, describes how regulation is reshaping the environment for growth, and presents four strategies for growth that new and established platforms alike can apply in response. More broadly, the article reflects on platform strategy's evolution from focusing on network effects and market tipping toward more richly understanding how platforms compete, grow, and create value over time.

THE LIMITS OF WINNER-TAKE-ALL THINKING

For years, the belief that network effects naturally lead to winner-take-all or winner-take-most markets has shaped strategy and venture investment across tech. (See "Pipelines, Platforms, and the New Rules of Strategy," HBR, April 2016.) Network effects are important, but giants such as Google and Amazon achieved a spectacular level of dominance in their respective industries because several complementary forces aligned to create a broader set of advantages. As these platforms grew, the cost of serving each additional user fell dramatically—economies of scale. At the same time, their incumbent status helped them expand into adjacent products and services more efficiently than new entrants could—economies of scope. (See "Are You Missing Growth Opportunities for Your Platform?" HBR, May–June 2025.) Much of what the giants offered could be delivered at near-zero marginal cost, especially in digital contexts. And the data they collected from users improved the services they offered in ways that further strengthened their advantage. Together, those dynamics created a loop in which scale allowed the platforms to provide better

performance at low or no price, which in turn attracted more users.

Amazon illustrates how these forces come together. As more buyers joined the marketplace, it attracted more sellers, strengthening the network effects. But the company's advantage did not stop there. Its scale allowed it to build highly efficient logistics and fulfillment networks. Those same capabilities could be leveraged across multiple product categories, spreading the fixed costs. Meanwhile, Amazon's data on consumer behavior enabled it to make better recommendations, provide more-effective pricing, and deliver more-targeted advertising, further improving the user experience. Each of those elements reinforced the others, making Amazon's position difficult to challenge.

When such forces align, markets can tip decisively. But the giants' success can mask the fact that this kind of alignment is rare—that true winner-take-all outcomes are the exception, not the rule. A case in point is the accommodations market, where multiple types of platforms exist. Traditional hotel groups like Marriott are platforms in that they don't own and manage most of their properties but do own their brands, loyalty programs, and booking services. These hotel groups compete not only with one another but also with a persistent fringe of independent hotels. In addition, when people book stays through intermediaries, hotel groups share revenue with online travel agencies such as Booking.com, Expedia, and a range of regional and local alternatives.

Over the past decade, a new type of rival has emerged in short-term rental platforms like Airbnb. These companies have created substantial value by expanding supply and increasing flexibility. While they have undoubtedly put pressure on incumbents, they have not displaced them. Many travelers still value the consistency and reliability of hotels, particularly for business travel, and remain wary of the risks associated with staying in a stranger's home.

The broader story in the accommodations industry, then, is not about a single player taking over but about multiple models coexisting. Even in the presence of strong network effects, no single platform has "won" the market, suggesting that dominance is far from inevitable.

One reason is that network effects are often local rather than global. Ride hailing provides a clear example. While Uber and Lyft compete in the United States, different players dominate in other regions—Didi in China, Grab in Southeast Asia, and Ola in India. Because ride hailing operates city by city, a platform does not need global scale to succeed; capturing a handful of local markets with sufficiently dense participation from passengers and drivers can be enough to sustain a business. The same is true in settings from waste management to labor platforms, where scale in one

For years technology companies operated under an implicit license to grow first and monetize later. Now platforms no longer have unlimited time to become profitable.

geography does not automatically translate into advantage in another. (See "Why Some Platforms Thrive and Others Don't," HBR, January–February 2019.)

But research that I conducted with colleagues points to another reason winner-take-all outcomes are hard to achieve: Users have different preferences, so they do not necessarily gravitate to the largest platform. In a study of a merger between the two largest pet-sitting platforms, we found clear evidence that having more users did lead to creating more matches, yet the benefits of network effects were largely offset by users' preferences for one platform over another. Indeed, when the platforms merged, users of the acquiree were forced to switch to the acquirer's interface and user experience, while the acquiree's features disappeared. Many users left after the merger, revealing that they valued design, brand, routines, or existing relationships more than the benefits of sheer size.

The implication is that even strong network effects may not be powerful enough to overcome differentiation. Platforms are not interchangeable utilities; they offer distinct products and services. And users care deeply about how those offerings fit their specific needs and preferences. That helps explain why platform growth has often come through brand acquisitions, particularly in sectors like online travel and hospitality, which have preserved the appearance of choice while concentrating control over data, pricing, and market access behind the scenes.

WHAT'S RESHAPING THE GAME

Platforms face a changing environment in several ways. First, regulation is increasing. In Europe the Digital Markets Act imposes constraints on large platforms that would have been unthinkable a decade ago, requiring data sharing and interoperability while limiting their ability to favor their own services or otherwise leverage their ecosystems. In the United States policymakers have historically relied less on ex ante regulation to prevent specific actions or behaviors and more on ex post antitrust enforcement through the courts, but this too is evolving as antitrust scrutiny intensifies. Until recently, the last major U.S. antitrust case against a digital platform was the one against Microsoft in the late 1990s. It focused on how a dominant firm used control over a key distribution channel, the Windows operating system, to favor its own browser, Internet Explorer. That case helped preserve the openness of the web, creating the conditions for new entrants like Google to emerge and grow.

Two decades later regulators are once again focused on control over distribution, with particular scrutiny of defaults and other forms of preferential treatment. These issues are at the heart of recent cases against Google in search and advertising and Amazon in e-commerce. There is a certain irony here: The legal logic developed in the 1990s to constrain Microsoft and favor new entrants like Google is now being applied to the search giant itself.

The new volume and scope of enforcement mark a clear break from the largely hands-off approach that authorities have long taken. In the recent Google case, the U.S. Department of Justice argued successfully in district court that the tech company had maintained its dominance in search partly through exclusive default agreements. Yet the remedies have been modest, focused mainly on limiting some of these agreements and increasing competitors' access to data. (Both sides filed appeals in federal court in mid-2026.) But the broader significance of the case and others like it is that they signal a growing willingness by regulators to challenge platform strategies that reinforce and entrench dominant positions.

Privacy regulation adds another complication. Data, once treated as a proprietary growth engine, is increasingly subject to consent requirements, restrictions, and oversight. Laws such as the General Data Protection Regulation in Europe and the California Consumer Privacy Act require firms to obtain explicit permission for data collection and use. Regulators are also cracking down on "dark patterns," or design choices that nudge users to share more data than they would otherwise. The strategic tools that once fueled platform advantage are now often viewed as liabilities.

A second way the environment is changing is about profitability. For years technology companies operated under an implicit license to grow first and monetize later. Now platforms no longer have unlimited time to become profitable. Investors remain willing to tolerate delayed monetization for many AI-first ventures but not for traditional platform businesses, which face increasing pressure to deliver clearer and faster returns. Recykal, a platform for exchanging recyclable

materials in India, shows how organizations are adapting to this shift. In 2024 the company chose to sacrifice some revenue in order to focus on activities that would improve its margins and overall profits.

These changes indicate that platforms' growth can no longer rely on regulatory laissez-faire, investors' patience, or the assumption that scale alone will drive continued success. Yet these same forces create an opportunity for new entrants and alternative models.

FOUR STRATEGIES FOR GROWTH

If winner-take-all outcomes are rare and will increasingly be constrained by regulation when they do arise, that does not mean entrepreneurs should stop building platforms. Quite the opposite. Platforms have already created enormous value by reducing information asymmetries, expanding access to products and services, and enabling entirely new forms of economic activity. For instance, they have turned small, disconnected local markets into robust, highly active, often global ones. Craft markets, short-term rentals, and freelance labor are just a few examples of sectors that have been fundamentally reshaped.

Platform companies have also figured out how to build trust at scale. Through screening mechanisms, reputation systems, and online reviews, they have made interactions between strangers not only possible but routine. A generation ago, the idea of getting into an unknown driver's car or staying in a home that you booked online would have seemed implausible. Today it is commonplace.

Finally, platforms have introduced flexibility into markets that were historically rigid. Traditional industries, such as hotels, often struggle to adjust capacity as demand changes, and others, such as taxis, also face relatively inflexible pricing. Platforms like Uber and Airbnb enable both supply and prices to respond dynamically, reducing inefficiencies and creating new opportunities for less professional providers.

The key shift now is in understanding that platforms don't need to dominate entire markets in order to create value. In this environment, success depends on making deliberate choices along four dimensions: where to grow, how to create value, how to monetize, and how to sustain trust in the long run. Four strategies can help businesses navigate each of these dimensions.

Prioritize users that create the greatest spillovers. When platforms attract multiple user groups, growth depends on which users create the strongest spillovers—that is, whose participation generates the most benefits for other groups of users. For example, on a marketplace where adding high-quality sellers attracts many new buyers but adding more buyers has only a modest effect on seller participation,

STRATEGY

sellers are the stronger driver of growth. Investing heavily in the side of the market that's easier to attract often delivers disappointing returns if it does little to stimulate meaningful interactions. The first mistake many platforms make is failing to identify which group of users truly drives growth.

Taskrabbit, a platform for local services, illustrates the importance of figuring out these market dynamics. In a 2020 study of the platform, my colleague and I found that user participation was highly elastic on the supply side: Taskers (people available for hire on the platform) were willing to take on significantly more work as demand increased, without requiring higher prices. That meant demand—not supply—was the key constraint on growth. By focusing on growing demand, Taskrabbit unlocked more transactions, which translated into higher earnings and stronger retention among taskers. Growth thus came from activating the most valuable side of the market, not from expanding both sides equally.

This logic is familiar in platform pricing: Platform companies often charge very little or even subsidize the user group whose participation generates the greatest spillovers. But pricing incentives alone are not always enough to produce growth. What's often required instead or in addition is to build capabilities that enable higher-quality participation. Amazon, for example, could not rely on third-party sellers to provide fast, reliable shipping at scale. It had to invest in a logistics and fulfillment infrastructure to make that possible. Doing so increased the value of the platform for both buyers and sellers.

Recykal, which connects and serves several stakeholders in India's waste management and recycling ecosystem, faced a similar situation. Increasing subsidies for those supplying recyclable materials could have attracted greater quantities of waste, but most likely at the expense of quality. Lower-quality materials, in turn, could have undermined the platform's reputation among buyers—including authorized recyclers and processing plants, consumer brands, packaging producers, and large-scale manufacturers. Rather than using subsidies as a lever to attract more supply, Recykal has started investing in deposit-refund systems and supporting infrastructure to improve the quality and consistency of supply.

Both platforms illustrate the same principle: Growth comes not from attracting the largest number of users at the lowest cost but from investing in the users and capabilities

Candidplatz

that create the greatest amount of value for everyone on the platform.

At the same time, prioritizing one side does not mean neglecting the other. Persistent imbalances can erode trust and undermine the platform in the long term. Uber's early focus on riders contributed to tensions with drivers, while Amazon's prioritization of buyers has, at times, strained relationships with sellers. Platforms are ultimately trust-based systems. Growth that comes at the expense of one side of the market is difficult to sustain.

Make the market "thicker" for current users. When platforms struggle to grow, companies' instinct is often to chase more users. But growth can come from a different source: creating more value from existing participants. Market "thickness"—the likelihood that users find high-quality matches—is not fixed. It is shaped by how the platform is designed. There are two primary ways to increase thickness.

The first is expanding the *scope* of the platform by adding products or services that increase opportunities for interaction. Many platforms pursue this path, and it can work well. But expanding too quickly can backfire. In early-stage platforms, having more categories does not always translate into creating more matches. In fact, my research shows that platforms often achieve higher match rates and stronger user retention when activity is concentrated in a smaller number of categories, at least early on. Specialization can make the market feel thicker by increasing the likelihood that users quickly find what they need.

Uber's creation of services like Uber Eats shows the potential of scope expansion. Adding new use cases can increase engagement and utilization across the network. But this approach is especially effective when those services increase the participation of existing users and providers—for example, by giving riders another reason to use the app and drivers another way to earn income.

The second, and often more-powerful, method is improving the platform's *design* to create more matches among users and to increase the matches' value. Doing so involves making decisions about pricing mechanisms, how quality is signaled, and how search and matching algorithms are structured, because small changes in these parameters can have large effects. For example, e-commerce platforms recognized early that the advantage of fixed prices over

auctions was the ability to reduce friction and speed up transactions, even if it meant sacrificing some price discovery. (Auctions help reveal what buyers are willing to pay.) Similarly, using verified user reviews over less informative credentials, such as someone's years of experience in a field, can improve trust and conversion without adding new users.

Uber's redesign of shared rides, where passengers traveling in similar directions share a vehicle, illustrates this second lever. By asking riders to walk short distances and wait slightly longer for pickups and shifting its algorithms from matching one rider with multiple potential drivers to simultaneously matching multiple riders with an available driver, Uber dramatically increased the pool of feasible matches. That change improved utilization, lowered costs, and made it harder for competitors to reach viable scale.

This strategy has two powerful effects. It increases value for users already on the platform, and it raises the minimum efficient scale required for competitors to enter the market or survive. In markets where winner-take-all outcomes are unrealistic, this kind of advantage matters more than raw size does.

Look beyond the obvious monetization options. The classic platform economics principle of charging the side of the market that is least sensitive to price still holds, but it is no longer sufficient. Today's platforms have access to monetization levers that go beyond simply charging users to participate or transact.

In the retail media model, for example, a platform monetizes its access to users by selling targeted advertising to third parties. Amazon pioneered this approach: Instead of charging consumers more, it charges sellers and brands to promote their products within the marketplace through sponsored search results and featured placements. In effect, sellers pay for visibility where consumer intent to buy is highest.

Other platforms are adopting these revenue models. What makes retail media powerful is that it builds on the data a platform already generates (search queries, browsing behavior, and transaction histories) to create highly targeted, high-conversion advertising opportunities. That allows platforms to monetize without raising prices on end users—and in some cases even allows lowering them—while giving suppliers a measurable return on their spending. Edenred, a payments and employee benefits platform,

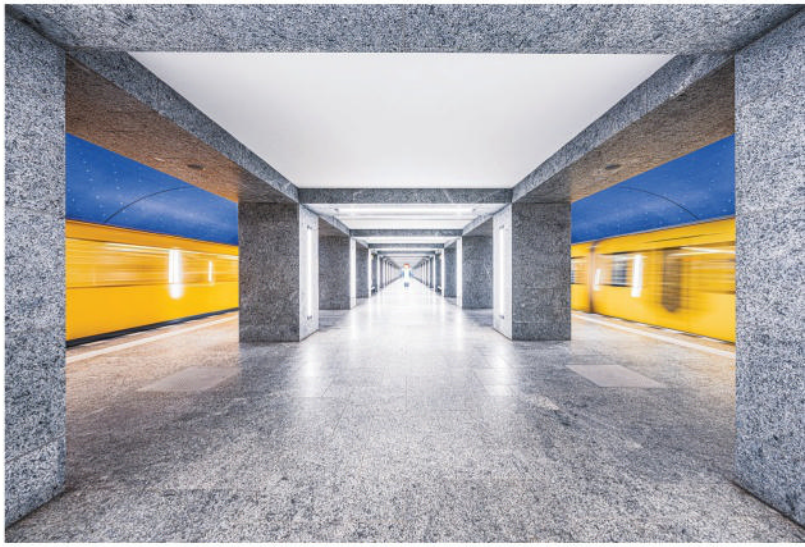

The platforms that endure will be those that make creating overall value and positive experiences for users a strategic principle.

is exploring how transaction data and user engagement can support new revenue streams by connecting merchants with employees at relevant moments. Marriott likewise leverages customer data and its loyalty program to create targeted offers and partnerships that generate revenue beyond room bookings. The key is not simply charging *more*; it's charging *differently* by aligning monetization with where value is created.

Platforms that rely on a single revenue model often find themselves constrained. Those that expand their monetization tool kit thoughtfully and transparently can unlock new sources of growth without alienating users.

Treat customer welfare as a strategy, not as compliance. Many platforms have driven users away with short-term, extractive tactics that increase engagement or revenue in the moment but erode trust over time. This approach sits at the heart of many regulatory battles. Practices such as self-preferencing—when a platform uses its dominant position to favor its own products, services, or apps over those of competitors—may boost profits in the near term, but they often come at the cost of credibility and user satisfaction. What looks like optimization today can become a liability tomorrow, both reputationally and legally.

The platforms that endure will be those that make creating overall value and positive experiences for users a strategic principle. These businesses depend on repeat interactions, and users must feel that the platform is acting fairly. Once that trust erodes, it is difficult, and often costly, to rebuild.

I see these dynamics play out across many platforms. Airbnb has invested heavily in trust mechanisms, yet travelers still worry that the home they booked will turn a dream vacation into a nightmare, and hosts fear serious damage to their property. Uber, after years of contentious changes, faces persistent skepticism from drivers, even when introducing features intended to improve their experience. In both cases, past decisions shape how new initiatives are perceived.

Google's experience with Privacy Sandbox highlights a subtler version of this challenge. The initiative was intended to replace third-party cookies with a system that protects users' information while still allowing advertisers to target audiences. Rather than sharing detailed user-level data, Privacy Sandbox was designed to keep more information within the Chrome browser and provide aggregated or anonymized signals to advertisers. Conceptually, the approach aimed to balance privacy and monetization, and research that I conducted with others suggests it is one of the most effective ways to obtain user consent for data collection and use.

Yet despite its technical merits, the proposal faced significant resistance and was ultimately abandoned, in part because regulators and other industry participants worried that shifting data and targeting capabilities into infrastructure controlled by Google could further entrench its market power. As a result, many publishers and regulators have fallen back on existing solutions, such as ubiquitous cookie consent banners—which are widely disliked by consumers yet persist because they are seen as more transparent and less likely to concentrate control over data.

In each of these cases, perception matters as much as reality. Even well-intentioned changes can fail if users or regulators do not trust a platform's motives. Critically, customer welfare should be interpreted broadly in platform settings, referring not only to end consumers but to all sides of the market. For Amazon, that includes both buyers and sellers; for Uber, both passengers and drivers; for Google, both users and advertisers. Regulators emphasize customer welfare because they recognize that short-term profits alone may not provide firms with sufficient incentive to invest in innovation and create long-term value. Platforms should take the same approach, not out of altruism but because creating value for customers is the most durable path to creating value for the business.

IN A WORLD of persistent competition, active regulation, and skeptical capital, platforms' success depends on disciplined choices: where to grow, how to create value, how to monetize, and how to sustain trust. The platforms that thrive in this next era will not be the ones waiting for markets to tip. They will be the ones that understand why many markets rarely do and that build strategies to succeed in this new reality. ■

HBR Reprint R2606H

CHIARA FARRONATO *is the Glenn and Mary Jane Creamer Associate Professor of Business Administration at Harvard Business School and a coprincipal investigator of the HBS AI Institute's Platform Lab.*

MANAGING YOURSELF

How Leaders Counter Pessimism in Trying Times

They model constructive behavior, reframe pressure, and highlight progress.

by Michelle Gielan

AT A LEADERSHIP session I led last year, an executive pulled me aside during a break.

"The strange thing is that the worst is behind us," she said, "but you'd never know it from talking to the team."

It had been a hard stretch: Layoffs the previous year. Two lost accounts. A reorg that left everyone doing more with less. The company had weathered all of it, and the numbers had finally steadied. But the team hadn't. People

Illustrations by **HARRIET LEE-MERRION**

Leaders have the power to shift the mood of their teams, not by pretending everything is fine or manufacturing cheerfulness but through tone, body language, and small actions.

who once volunteered bold ideas had gone silent. Meetings felt flatter. Conversations drifted toward what might go wrong next.

"It's as if everyone is still bracing for impact," she told me.

As we talked, she realized that the problem was no longer operational. It was psychological. The noise from news stories and social media about the challenges facing their industry, company, and careers had become too loud to bear, and people were starting to believe that their efforts mattered less than they used to.

Many leaders are managing this kind of pessimism right now, and research explains why. In a study I ran with Arianna Huffington on how negative information affects the brain, we gave people three minutes of negative news in the morning, then measured their mood six to eight hours later at the end of an ordinary workday. Compared with the participants in the control group, who saw three minutes of neutral content, the people in the negative-news group were 27% more likely to describe their whole day as a bad one. Days upon days of such exposure—whether world news, industry warnings, or team and personal setbacks—can wear down well-being and performance until work starts to feel pointless. Against that backdrop, it is not surprising that employee engagement is at a 10-year low and that more than half of U.S. workers report burnout.

Feeling pessimistic, apathetic, or stressed out is not a character flaw. Our bodies are wired that way. The brain has a bias toward negativity, weighting threats far more heavily than reassurances, so a few alarming inputs can color people's thoughts for hours. Feed that bias an always-on stream of problems mostly outside our control, and it hardens into something more corrosive. Martin Seligman, the founder of the field of positive psychology, called a version of this pattern learned helplessness: the mental state people enter when they conclude their actions can no longer change the outcome. I studied at the University of Pennsylvania with Seligman, and I often see the effects of *learned helplessness* in employees: Keep your head down, don't bother, nothing changes anyway.

That attitude is both personally and organizationally destructive. Having tested more than 25,000 professionals in *Fortune* 500 companies, hospitals, and school districts over the past decade, we have consistently found that pessimists are not only less engaged but also less successful. People who scored the highest in work optimism—strongly agreeing with statements about their workplace such as "I generally expect things to go my way," "I feel optimistic about the future," "I anticipate more good than bad," and "I believe my behavior matters"—were, by contrast, three times more likely to be highly engaged, 40% more likely to be promoted within 12 months, and five times less likely to burn out. Other studies have found that people with a positive mindset are 31% more productive and that optimists outsell their pessimistic peers by 37%.

Here's the catch most leaders miss: You can't cure pessimism with a motivational speech or a good-vibes-only demeanor. Research shows that when people feel they must fake positivity, it leads to higher burnout. Instead, you need to build a culture of what I call *resilient optimism:* the ability to see challenges clearly while believing that actions can still influence outcomes.

By studying effective managers, I've found that you accomplish this through a set of simple behaviors: modeling optimism, reframing stress, and highlighting progress.

MODEL CONSTRUCTIVE BEHAVIOR

Attitudes are contagious. Through the brain's mirroring systems, we automatically pick up and reproduce the emotional states of the people around us. Leaders, in particular, have the power to shift the mood of their teams, not by pretending everything is fine or manufacturing cheerfulness but through tone, body language, and small actions.

Even one specific expression of appreciation can move performance measurably. In a field experiment at UPenn, researchers had a manager thank one group of fundraisers for their work—a brief, genuine "We're grateful for what you do"—but say nothing to a comparison group. Over the following week, the people who were thanked made about 50% more calls. What changed wasn't their skill or their workload; it was that they felt valued. Small repeated signals from a leader helped the team spiral upward.

When thinking about the types of behaviors you want to model, start with the signals people read most closely: what you pay attention to, what you prioritize, and how you restore energy.

Attention. Reassure your team that you are focused not just on threats but on purpose, agency, and positive signals.

Stress usually provides information—the mind is flagging something it values. So encourage your people to get curious instead of critical about this feeling.

Flag a customer's or client's compliment, the hard work put in by a particular employee or team, or an optimistic industry report. Even something as simple as saying "Thanks" or "Good catch" out loud or chuckling over a snafu in a meeting can lift people's spirits.

Prioritization. A leader who treats everything as urgent overwhelms everyone. Instead, commit to a regular reset that your team can witness. Once a week, review your own calendar and commitments and ask three questions: What can I step out of? What can I hand off with guidance instead of doing myself? What am I doing out of habit or guilt rather than impact? This kind of maintenance keeps you—and the standard your team sees you set—in a productive mode that holds pessimism at bay.

Restoration. Optimism requires energy, so it's important to show your team that you take the time to relax and refuel. You don't need a week in Tuscany. Simply treat—and talk about—your weekends and even short breaks as critical. Anticipate those breaks, detach fully, and return to the office carrying and articulating one good feeling. When you model genuine recovery, you give the whole team permission to do the same.

REFRAME YOUR AND YOUR TEAM'S STRESS

Modeling an optimistic attitude sets the baseline. The next move involves helping your team members reinterpret the pressure they're under.

Researchers Alia Crum, Peter Salovey, and Shawn Achor (who happens to be my husband) split UBS bank managers into two groups. They trained one to treat stress as a threat to reduce and the other to see stress as something that sharpens attention, fuels quick thinking, and drives growth. And they found an interesting divergence. Four months later, both groups were more stressed because of an industrywide crisis. But the group that had learned about the benefits of stress reported a 23% drop in stress-related symptoms such as headaches, backaches, and fatigue and logged better performance with higher satisfaction and less burnout.

You can teach your team to reframe stress with two practices:

Signal check. Stress usually provides information—the mind is flagging something it values. So encourage your people to get curious instead of critical about this feeling and ask themselves, What is this stress protecting or revealing? What value of mine feels threatened? How can I respond in a way that honors it?

An executive I worked with used this tool to help a strong performer who became tense and short-tempered as deadlines closed in, snapping in meetings and redoing others' work. Instead of telling the team member to relax, the manager privately asked her the above questions. It turned out that she was afraid that one visible miss would undo years of hard work. That reframed the conversation from "You need to be nicer and manage your stress" to "You want to protect your reputation as a positive contributor"—and her grip loosened.

"Yes, and." The principle comes from improv: Accept the reality in front of you ("yes") and then build on it ("and").

Resilient leaders don't sugarcoat or stall. They acknowledge challenges: "The launch has failed" or "We lost a major client." But they refuse to let the negative be the end of the sentence: "...and here's how we gather feedback and relaunch smarter" or "...so starting this week we're prioritizing connection and communication." This recognizes the difficulty while redirecting people to solve problems, which helps foster more-flexible thinking.

CREATE AND HIGHLIGHT THE POSITIVES

The brain is a pattern maker. In a noisy environment it strings negative dots into a discouraging story: Nothing I do matters. The most powerful thing a leader can do to counteract this tendency is to supply positive dots instead.

Charlene, an HR leader at a large tech company, discovered this during a brutal round of layoffs when her office became a drive-through for complaints. Then one afternoon she encountered the first flicker of lightness in weeks; she overheard two engineers high-fiving over a bug they'd finally cracked. After that, she leaned into soliciting positivity, stopping people in the hallway and saying, "Tell me something good that happened today." Soon people started to notice and remember their smallest wins to share. Her walks around the office became, in her words, "roving pep rallies," and the mood lifted. She wasn't ignoring the layoffs. She was curating the good that was already there but going unspoken.

In research Shawn and I did with *Training* magazine, we found that,

on average, 31% of people at work are optimistic but don't express it. A leader who draws those closet optimists out changes the emotional climate.

Here's a three-day exercise that will help you and your team members to publicize and generate positive dots:

Day 1: Send praise. Tell an employee who is doing a great job what you noticed and why it mattered. Text or email the person directly, or tell one colleague about another's quiet win. "Did you catch how Mia stayed late to unblock the launch, even though it wasn't her project?" This positive gossip travels faster than formal praise and lands as sincere because people assume you've got nothing to gain by saying it.

Day 2: Lighten a load. Ask one teammate what small thing would make the next step easier—a certain resource, a phone call, 10 extra minutes—and provide it. Friction removed is a dot created.

Day 3: Reach back out. Reconnect with a colleague you've lost touch with. No agenda—just a "You crossed my mind." Getting together with someone when you need nothing is the opposite of the head-down, brace-for-impact posture hard times breed. Plus, the most powerful support is often just the reminder people need that someone hasn't been forgotten.

Scaled up, this becomes culture. When Janet, an executive at the Leukemia & Lymphoma Society (now Blood Cancer United), saw burnout creeping through her team, she worked with me to launch a three-week optimism challenge—daily notes of appreciation, then gratitude, then small acts of kindness. By the end, her employees had generated more than 10,000 moments of connection and regained the sense that their actions still mattered.

THE LEADER'S REAL JOB

When a team, industry, or organization is struggling, pessimism can feel justified—and even useful as a protective shield against false hopes. But if that attitude festers, it will kill productivity and performance. None of this means papering over hard truths. Optimism won't fix a broken strategy or an unsafe workplace. But it can change what people do with a difficult reality—whether they disengage and wait, or stay in problem-solving mode long enough to find the next move. That's the leader's job: to move people toward a still realistic but more positive mindset that prioritizes agency and action over analysis. If you model constructive behavior, help your team reframe stress, and create and highlight the positive dots of progress, optimism will build, person by person, day by day, until it becomes your team's prevailing spirit.

HBR Reprint R2606J

MICHELLE GIELAN *is a positive psychology researcher and the author of* Resilient Optimism *(Harvard Business Review Press, 2026), from which this article is adapted. She has spent 15 years studying resilience and performance in leaders.*

How We Did It

Robinhood's CEO on Refounding the Business Post-IPO

by Vlad Tenev

FOR MANY FOUNDERS, taking a company public feels like a graduation. You have proved it can survive. The next chapter, in theory, is about managing long-term, predictable growth.

Robinhood went public in 2021 after my cofounder, Baiju Bhatt, and I had spent nearly a decade building the commission-free trading platform. Tens of millions of customers were using it, and revenue had exploded during the pandemic-era retail investing boom. Our mission to democratize finance was part of the mainstream conversation. It seemed there was nowhere to go but up. But the reality was far more complicated than I could have imagined.

Our IPO priced the company at $32 billion, and our stock price doubled in the first week. But soon after, almost everything about the market environment that had fueled our rapid growth began to deteriorate. Inflation surged. Interest rates rose at the fastest pace in decades. Speculative assets collapsed. In less than one year, our share price dropped from $80 to below $10.

Initially I viewed the downturn as mostly a market problem. But over time I understood that the bigger issue was structural. Tailwinds that had powered our growth during Covid—near-zero interest rates, abundant liquidity, and an abrupt closure of much of the offline economy—reversed nearly as quickly as they had appeared. We had optimized the company for one era, but that era was gone.

So we got to work, beginning a multiyear effort to rebuild; I look back on it as a refounding. We deprioritized low-revenue activities and focused on our core business, reduced management layers, and most importantly, gave leaders ownership and autonomy. That let us reach profitability and eventually diversify into over a dozen business lines, which each generate more than $100 million in annualized revenues.

Our commitment to ownership and operational clarity wasn't merely a strategic pivot. For me it was a return to a fundamental value rooted in my youth. I've long believed that when individuals have a stake in the systems around them, they can move beyond mere survival and thrive—a lesson

Photograph by **CHRISTIE HEMM KLOK**

I've long believed that when individuals have a stake in the systems around them, they can move beyond precarious survival and thrive.

I first learned not in a boardroom but in childhood, witnessing the fragility of a society where economic ownership was out of reach.

FROM MATH TO FINANCE

I was born in Bulgaria as the communist system was beginning to collapse. My great-grandparents had given up their houses and land in the name of collectivization—for the party and public good. In return they were given an apartment in one of the many uniform gray tower blocks, now covered in graffiti. In 1991, my last full year living in Bulgaria, the annual inflation rate was north of 300%—a number inconceivable to those of us living in the United States today.

Some of my earliest memories are of looking out our apartment window to gauge the length of the lines to buy milk and eggs. Or of huddling around our family radio during the daily rolling blackouts. I understood from a young age how brittle a society without broad ownership could be. Ownership allows people to shift from worrying about day-to-day subsistence to having a more long-term orientation.

As a kid I became interested in investing, so my parents set me up with a brokerage account. I remember buying my first stock: a company called 3Com, which made the PalmPilot. A few weeks later shareholder materials arrived in the mail, inviting me to vote. I realized that I had become an owner of this company that made a product I believed in.

Years later, during my first month in graduate school to study math, Lehman Brothers went under—the start of the global financial crisis. My perception of how trading worked had been influenced by movies like *Trading Places*, where people shouted at one another on the exchange floor. But the era of person-to-person transacting, known as "open outcry," was ending. Firms were increasingly using algorithms and software to identify tiny pricing gaps and execute deals at extraordinary speed. The action was moving from the New York Stock Exchange to massive data centers in New Jersey.

And there was a deep cultural shift at play. In 2012, when Baiju and I got the idea for Robinhood, our generation was still feeling the reverberations of the financial crisis. Many of our friends had lost their jobs. Books like *Flash Boys* and movies such as *The Wolf of Wall Street* and *The Big Short* helped perpetuate the notion that the crisis was an own goal by the financial industry itself. The malaise and institutional distrust led to the worldwide Occupy Wall Street protests.

I always thought that Occupy Wall Street, largely a Millennial and Gen X phenomenon, was really about ownership. Those young people didn't own anything—other than debt—and if you don't own anything, you feel like you have no stake in what happens to your country.

Baiju and I believed that if we could make it easier to own U.S. equities, we could help close the gap between those who have everything and those with nothing. We believed this would also have a civic benefit: improving financial literacy and instilling a sense of pride in free capital markets.

When we pitched the idea to venture capitalists, there was little excitement about funding a regulated financial services startup. "Fintech" barely existed as a category, and our approach didn't align with the prevailing investing narrative, which was centered on the rise of exchange-traded funds and robo-advisors. But in our view, markets had become fully electronic and low-cost, and coupled with the rise of the smartphone, that created opportunities for self-directed investing.

We were turned down a lot. But eventually we raised our early financing rounds, and then we focused obsessively on ensuring that our onboarding was fast, the product intuitive, and the experience designed for the smartphone. Our team of two dozen people launched Robinhood to the public in 2015, and within four years we had disrupted the brokerage industry.

OPERATIONAL CHANGES

At the end of 2019 Robinhood had roughly 700 employees. Two years later we had nearly 4,000. That period overlapped with the Covid pandemic, meaning that by late 2021 most of our employees had never met in person. The steep hiring trajectory seemed rational at the time. Millions of first-time investors were entering the market. Revenues were ramping exponentially. We needed to increase staff to handle the rapid growth, shore up our infrastructure, and build the control functions required to become an IPO-ready financial services company.

It was hard to feel the head count growth, though, when we all worked remotely: Our computer screens stayed the same size. But ultimately the organization had become too complex.

> The best information rarely comes from dashboards or executive reports. It comes directly from users, from hearing where they are confused, frustrated, or excited.

Management layers multiplied. Teams fragmented. Decision-making slowed. Execution ground to a halt. Ownership became diffuse. I realized that drastic change was needed.

One tactic I employed frequently was to step back and imagine I was a new CEO brought in to run the company, free of the baggage of having made the decisions that got us to where we were. What would a new CEO do? That perspective allowed me to think differently about how Robinhood could look going forward, letting go of the innate desire to maintain consistency. Did we really need to be remote-first? How many management layers should we have? Why were decisions taking so long? How should we attract, grow, and reward our workforce? Thus began the painstaking process of refounding the company.

We gradually brought people back into our six locations, using a return-to-office structure where individual contributors came in several days a week, managers more often, and senior leaders full-time.

I heard every version of the same warning: Once you have given employees flexibility or perks, you can never rescind them. I was told that reversing remote work or tightening expectations would permanently damage culture and morale. Instead I learned that employees—especially high performers—are far more dedicated and adaptable than many leaders assume. If you're clear about why changes are happening and the company begins operating better afterward, most people adjust quickly.

I also realized that leadership is about facing difficult situations head-on, before circumstances force a hasty decision, and then communicating resolutions clearly.

We changed how we recruit and compensate talent. Upon reflection it was clear to me that we operated best with a small, elite team rather than with a large group of slightly above-average performers. To make that change we had to overhaul our compensation philosophy to reward top performers disproportionately, bucking the Covid-era trend of being more egalitarian.

We made painful workforce reductions, and we spent weeks simplifying our operating plan, assuming that employing fewer people meant shrinking our ambitions. But surprisingly, it turned out that we moved faster and got more done. We became less attached to sprawling road maps and more focused on the priorities that were tied directly to significant business outcomes.

We also reduced management friction by changing how the company was organized. Previously, Robinhood had been structured around function, with product, engineering, and operations flowing through centralized leadership even as individual businesses became increasingly complex and heavily regulated. That led to the creation of more committees and sign-offs and blurred accountability.

In our new structure, business leaders owned nearly everything needed to drive their businesses: engineering, product, operations, and financial performance. They took responsibility not just for product quality and operational stability but also for execution, resource allocation, regulatory performance, and goal achievement. That design reduced toil and bureaucracy—fewer processes, fewer committees, fewer approval layers.

I focus a lot on ownership and its impact on society, but it's also powerful at the company level. Giving individual leaders ownership over functions across their business created a stronger culture of accountability, which accelerated decision-making. That became evident later with our rapid product velocity.

Part of what we were rebuilding was our culture. In the early years we prided ourselves on being a lean, high-performance team with intense customer focus. I wanted to get back to that while maintaining the strong foundation of security, reliability, and regulatory compliance we had painstakingly built.

EXPANDING THE BUSINESS MODEL

During the Covid boom, Robinhood had become heavily tied to speculative trading by casual investors. That made the company fragile and overly exposed to swings in market conditions and sentiment. We needed a more durable model.

Two observations were key to getting the business back on track after monetary policy tightened. First, while novice investors and traders pulled back from the markets, sophisticated active traders remained resilient. They had the tools and knowledge to take advantage of declining and sideways-moving markets. Second, while monetary tightening was bad for investing in stocks, the higher interest rates that it brought opened new opportunities for us to diversify.

So we focused on two responses to those observations: making Robinhood

FACTS & FINANCIALS

Robinhood

Founded: 2013
Headquarters: Menlo Park, California
No. of employees: 3,029

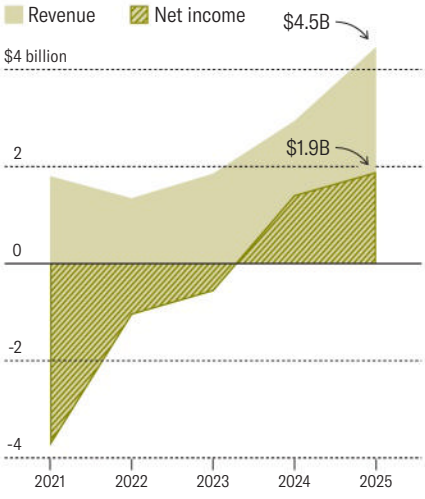

Source: Robinhood

the best platform for active traders and rolling out products that resonated during that high-interest-rate era.

That purpose led us deeper into new offerings, including high-yield and retirement accounts, banking and advisory services, and credit cards. Robinhood Gold became a very important part of the transition. Designed as a wealth-building engine rather than just another monthly subscription, it aimed from the start to provide value to customers across all market conditions, with benefits that today include everything from a high annual percentage yield on cash to a 3% IRA match.

Diversification changed the economics of our business. Instead of depending on bursts of trading activity during speculative markets, we were building recurring and longer-lasting relationships.

Throughout those shifts I believed it was critical to stay close to customers. The best information rarely comes from dashboards or executive reports. It comes directly from users, from hearing where they are confused, frustrated, or excited. That's how we identified a series of traders with smaller balances getting caught up in pattern day trading restrictions (PDT is the execution of four or more day trades within five business days in a margin account; the rule was repealed in 2026). Those types of recurring signals often shape where we invest in product development and push for industry changes.

By the end of 2023 the changes were working. We had reached GAAP profitability. Numbers aside, it felt qualitatively different: There was energy and enthusiasm in our offices. A feeling I had lost sometime during Covid was returning—that we were one Robinhood, all invested in the same mission.

During that period I learned that public narratives often lag operational reality. I was in London in November 2023, helping with Robinhood's UK launch. Our stock was trading at around $8 a share, and I had a breakfast meeting with a long-time investor.

The conversation was brutal. He questioned whether the company had a future and suggested that we sell. The subtext: Did I really have what it took to return the firm to its previous heights? What struck me afterward was not his criticism; it was my conviction. In the two years prior I had become used to having that kind of conversation with investors. This time, though, I knew Robinhood was on the right track. Operationally the company was much stronger than it had been; the market simply hadn't caught up. The meeting clarified something for me: Perception trails actuality. While navigating that gap, leadership becomes less about listening and adjusting and more about maintaining focus and execution.

LESSONS FOR OTHER LEADERS

Over time the market caught on. Our revenue reached $4.5 billion in 2025, increasing more than 50% year over year. Our market cap surpassed $100 billion, and we were inducted into the S&P 500. In the second quarter of 2026, Robinhood generated more than $1.3 billion in revenue and over $573 million in net income. Customers hold more than $350 billion in assets on the platform.

My greatest lesson from that era? A founder's values cannot live only in speeches, mission statements, or brand campaigns. They must become the company's values. A founder's culture needs to become the company's culture.

That is why Robinhood's refounding worked. It was a return to what I believe: People perform better, think longer-term, and build more-resilient systems when they have real ownership.

For a founder, that alignment matters enormously. If the values on the wall differ from the values driving your hardest decisions, people will feel the rift. But when the company's strategy, structure, and culture all reinforce a founder's convictions, the organization becomes harder to break.

Going public was not the end of building Robinhood. In many ways it forced us to rebuild the company. The market and our customers had changed, and we needed to change too. But the belief at the center remained: Ownership creates resilience. That was true for our customers, and it turned out to be true for Robinhood. ■ **HBR Reprint** R2606K

CASE STUDY

HARVARD BUSINESS REVIEW

HBR's fictionalized case studies present problems faced by leaders in real companies and offer solutions from experts. This one is based on the Darden Business Publishing case study "Team Conflict: The 'Chatty' Accusation at the Customer Support Call Center" (case no. UV6864), by Kristin Behfar, Gerry Yemen, and Martin N. Davidson, which is available at ***HBR.org****.*

Are We Using the Right Customer Service Metrics?

by Kristin Behfar, Martin N. Davidson, and Gerry Yemen

JENNIFER AMES ARRIVED at the Austin, Texas, headquarters of Baby Precision Care early enough that the customer support floor was still dark. She liked the call center this way—quiet and empty, with the workstations' headsets waiting in their cradles. In her previous roles at BPC, showing up early had

Illustrations by **JORI BOLTON**

As orders increased, executives decided the company needed a Spanish-speaking support team. And they wanted the team's calls to be guided by a large language model.

helped her earn a reputation as someone who could take on complicated assignments and make them operational. She had recently won kudos for her role in BPC's expansion into Latin America.

But this assignment felt different. One month ago, Jennifer had been named senior director for global customer support, a new position. For the first time in her career, she was building a department—one that company leaders hoped would model how employees could integrate AI into their workflows.

The need had emerged quickly. BPC sold medical equipment for infant care, and demand from clinics and hospital systems across Latin America had been stronger than expected. Its product lines included neonatal heart and breathing monitors, phototherapy units for newborn jaundice, and small home-care devices that clinical partners could recommend to parents. The products were technically sophisticated and often used by anxious families far from a hospital. When customers had questions, the answers needed to be accurate and easy to understand.

At first, BPC had relied on sales representatives to provide post-sale support. But as orders increased, executives decided the company needed an in-house Spanish-speaking support team. And they wanted the team's calls to be guided by a large language model.

Jennifer had been asked to design the prototype of the AI integration. She chose four reps, all fluent in Spanish, and placed them in a small call center near the product quality and training teams. But this morning, as Jennifer reviewed the team's dashboard, she could see that something wasn't working. Two metrics were red.

The first was average handle time, or "handle time discipline," which Jennifer had set as the team's most important goal. The target was seven minutes, but currently the group was several minutes over that.

The second red metric was AI script adherence, which Jennifer had made the next top priority. Three reps were using BPC's new call-assist tool consistently. One was not.

Jennifer clicked into the individual view. Emma Henley, Julio Lopez, and Rob Garvey were all close to the target. Their calls generally followed the same pattern: Identify the issue, confirm the product, select the right AI-generated script, direct the caller to the appropriate training video, verify the next step, and close.

Then there was Sonia Enriquez. Her calls were far longer. She frequently ignored the AI's script. She stayed on the call after the tool recommended sending a training video. She asked questions that were not in the script. Several calls had been marked "off flow" by the system.

Jennifer rubbed her forehead. Sonia wasn't an employee she'd expected to have to worry about. A high performer in her previous role in BPC's Latin America sales support group, Sonia was known for her product knowledge and ability to build trust with customers. Jennifer had personally recruited her, thinking Sonia would bring practical experience to the new team.

Now the dashboard made her look like a problem. Jennifer winced, realizing the rest of the group would not be happy that one colleague was pulling them all down. Their year-end bonuses were based on a mix of company performance, individual performance, and team operations. The team portion included average handle time, AI adoption, and basic resolution metrics. Jennifer had

wanted to avoid a cutthroat environment in which representatives competed against one another for the easiest calls. A shared component, she believed, would encourage them to trade tips, coach one another, and take collective responsibility for the new department's success.

She looked again at Sonia's call records. The pattern was unmistakable. Every time Sonia moved away from the script, her call time climbed.

Jennifer's inbox pinged. It was a message from her manager: *How's the AI scripting going? I'm excited to show early results at next month's ops review.* Jennifer stared at the note, wondering how to respond. The answer depended on which part of the dashboard he wanted to see.

She opened a blank document for that morning's team meeting and typed three agenda items: AI adherence, handle time discipline, and team accountability. Then she added one more line to the document, highlighting it for emphasis: *How do I get everyone using the same process?*

WHAT "HELPING PEOPLE" MEANS

Jennifer gazed at her four employees, unsure of how to kick off the meeting. "Good morning to you all. I have a tough topic to discuss today," she finally began. "We're not here to blame anyone. But our average handle time is longer than it needs to be, and our AI adherence score is below target. I want to talk about what's getting in the way."

Emma Henley didn't wait. "What's getting in the way is that three of us are following the process, and one of us isn't."

Jennifer felt the meeting tilt before she could stop it.

Emma pointed to the dashboard on the wall screen. "We use the AI script, identify the issue, send the right training video, confirm the next step, and close. That's the process."

"And it works," Julio Lopez said. "We're under seven minutes."

Rob Garvey looked across the table. "Three of us are."

Sonia Enriquez sat back in her chair, her arms crossed. "Three of you are rushing people off the phone."

"No," Emma said. "We're doing the job we were hired for."

"All you're doing is reading from a machine."

"It's an approved support tool," Julio said. "We're a support team."

"It's a robot," Sonia replied. "If all you do is repeat what it tells you, why do they need you?" Suddenly the room was quiet.

Jennifer raised a hand. "Let's slow down. Remember, we're all on the same side here."

But Rob was already leaning forward. "The team score isn't good enough because of your calls. That affects everyone's bonus."

Sonia's face tightened. "I actually help people."

"You talk to them," Julio said. "That isn't the same thing."

Sonia opened the report Jennifer had sent the team and turned to the last page. "Did anyone read this?" She tapped the customer satisfaction column. "My calls are longer, but my customers are happier."

Jennifer looked down. Sonia was right. Her calls deviated from the AI script and ran well past the target time, but her satisfaction scores were the highest on the team.

"These customers are not calling about a toaster," Sonia said. "They are using medical equipment for babies. They're scared. They don't want to be pushed to a video and abandoned."

"But you're not doing what the tool tells us to do," Emma said.

"Then maybe the tool is wrong." Sonia's eyes flashed. "Yesterday a clinic director called about a monitor setup. The AI told me to send the video and close the call. But she had questions about a second site they're opening, so I answered her questions and connected her with sales. Maybe it becomes an order. Maybe it doesn't. But that opportunity disappears if I act like a machine."

"That's not customer support," Julio said. "That's freelancing."

"It's understanding the customer."

"It's being a Chatty Cathy while the rest of us do the real work," Emma shot back.

"Enough," Jennifer said loudly. "That's enough."

For a moment, the only sound was the low hum of the air-conditioning. Jennifer looked from Emma to Sonia to the report in front of her. A month ago, the model had seemed simple: one team, one target, one AI-enabled process. Now the target was creating resentment, the AI system was at the center of

"If Sonia is building trust and seeing things the script misses, maybe she belongs in sales. Not in a queue where every extra minute makes her look like a liability."

the conflict, and the employee who had been labeled least compliant might be the one customers trusted most.

Jennifer had wanted accountability. She'd gotten a fight.

MATCHING ROLES TO GOALS

Jennifer stood in the doorway of Marcus Reed's office. Marcus ran the product quality group and had become a mentor to Jennifer over the past few years. "Do you have 10 minutes?" she asked.

"For you, I have 12," Marcus replied warmly.

Jennifer sat down and told him everything. Marcus listened without interrupting. When she finished, he leaned back.

"You have at least three problems," he said, "and you're treating them like one. First, you have a conduct problem. People can disagree about process. They can't mock each other or turn a team meeting into a pile-on. You need to address that directly."

Jennifer nodded. That part was obvious, though not easy.

"Second, you have a Sonia problem," Marcus said. "Or maybe a Sonia opportunity. She may be in the wrong role. If she's building trust, finding sales leads, and seeing things the script misses, maybe she belongs in sales or even regional sales training. Not in a queue where every extra minute makes her look like a liability."

"I recruited her for this team. She's exactly the kind of high performer I wanted."

"Maybe, or maybe she's the kind of person your model doesn't know how

to use." Marcus cleared his throat. "The third problem is the scorecard," he continued. "You set up a group bonus for employees who don't actually work together. They handle their own calls and accounts. The only thing making them a team is the metric."

"You're right. The incentives probably don't make sense. So I should change them."

"I think you have to. But that may be the less risky conversation."

"What's the more risky one, in your view?"

Marcus turned his laptop toward her. On his screen was the companywide memo announcing BPC's AI adoption goals: faster resolution, consistent guidance, scalable support, operational discipline. *Our top priority as a company* was the title of the memo. "The AI," he said.

Jennifer looked at the screen. "Leadership wants proof that it's working."

"Exactly. And your pilot is supposed to be a proof point."

"So I tell my manager the tool is wrong?"

"Do you think that's what he wants to hear?" Marcus asked, a twinkle in his eye. "And do you think he'll accept that

Should Jennifer rethink the AI pilot?

THE EXPERTS RESPOND.

DAISY AUGER-DOMÍNGUEZ
is the global chief people officer at the fintech Digital Asset, where she leads people strategy, culture, organizational design, and AI-driven workforce transformation.

Jennifer should clarify team expectations and redesign their work with AI.

Although she built the system that leadership asked for—prioritizing average handle time, script adherence, and team accountability—doing so revealed a challenge: Not every customer support call is the same. This type of learning is inevitable with the rollout of any transformative technology. We're all learning how to integrate AI, and we have to accept that there will be false starts.

Jennifer should explain that to her team members and ask for their patience with the process and one another. They need to hear that they can disagree—but that mocking a colleague or piling on in meetings is unacceptable.

With Sonia, she needs to be supportive but firm. High satisfaction scores alone do not prove that longer calls are always better. A call center must still deliver service at a cost and pace the business can sustain.

At the same time, Jennifer needs to own her part of the problem. She wanted a tight-knit group with shared accountability, but she created a team bonus component for people who mostly work alone. They do not share calls or hand off accounts. So when one person's style threatens a team metric, collaboration quickly turns into resentment.

That's why she needs to redesign the work around call types. Some are procedural: A customer needs setup help, a training video, or a simple answer. Those calls should probably be handled quickly, with help from the AI script. Other calls involve anxious parents, institutional relationships, or possible sales opportunities. Those may require more listening, judgment, and trust building.

Instead of replacing one simplistic metric with another, Jennifer should examine satisfaction scores, repeat calls, escalations, qualified sales leads, and product-quality signals alongside handle time. Then she can create different targets for different call types. She might also use Sonia's insights to identify when to move from a procedural path to a high-touch one.

Jennifer can then go to her manager with a solution rather than a problem. Instead of "the AI isn't working" or "I can't get all my team members on board," her message is: "The pilot is teaching us something important. For procedural calls, the AI script is helping reps move quickly and consistently. But some calls require a different kind of support, and our current metrics don't capture that value." Framed this way, Jennifer is not asking to abandon the AI initiative. She is showing BPC how to improve it.

Six months from now, success will look like a team that

Jennifer could raise an uncomfortable possibility to her boss: that BPC's AI tool was helping the team move faster while missing something important about what customers need.

the AI is the problem? Or will he suspect that you are? This is your first time running a team, after all."

Jennifer's phone buzzed with a calendar reminder for her weekly meeting with her manager. She showed it to Marcus. "I don't know how to handle this," she said.

"Well, that's why they pay you the medium bucks," he replied, smiling. "But for now, it looks like our 12 minutes are up."

Jennifer walked out of Marcus's office deep in thought. In the next few minutes she would need to decide what she was going to recommend. She could double down on the AI-based targets and either transfer Sonia to another department or tell her to stick to the AI script for all calls. She could revise the team bonus structure. Or she could raise an uncomfortable possibility to her boss: that BPC's AI tool was helping the team move faster while missing something important about what customers need. ■

KRISTIN BEHFAR *is a teacher in the executive programs at the University of Cambridge.* **MARTIN N. DAVIDSON** *is the Johnson & Higgins Professor of Business Administration at the University of Virginia's Darden School of Business.* **GERRY YEMEN** *was formerly the senior researcher and executive director of the Darden Case Writing Research Group.*

understands the work, trusts the metrics, and feels empowered to improve new tools rather than follow them blindly. In the end, that will lead to a more durable AI success story, one that Jennifer, her team, and BPC can be proud of.

RONNI ZEHAVI *is the CEO and a cofounder of the HR tech platform HiBob.*

Jennifer should work with her boss on new metrics and goals.

Almost every company is in AI adoption mode right now. You buy a tool, give people access, train them, and make sure they try it out. That first phase is necessary, but its results don't tell you whether AI is improving team or organizational performance. Of the thousands of clients that my company works with each year, I'd estimate that around 80% are still in that phase of AI adoption. Fewer than 20% have moved on to achieving and effectively measuring impact.

Jennifer is running an experiment but is being judged on a scorecard that was established too quickly. It currently measures compliance over judgment and speed over trust, without understanding the metrics' full impact on customers or call center employees. Until something changes, she and her team will keep running into trouble.

Now she needs to reset expectations about the speed of her team's AI integration and value creation. A shift in work design, behaviors, skills, and culture cannot happen overnight. She and her manager need to agree on short-term and long-term goals for both team efficiency and customer satisfaction, and they should remain flexible on how those are achieved. She could say something like, "I know my job is to reinvent how my team works with AI. But let's be clear about the outcomes we want to achieve—and recognize that we may need additional training, encounter resistance, or have to adjust as we go."

Jennifer shouldn't underestimate the behavioral piece of this situation. Most of her team is falling into line, but how long will it be before they start quietly calculating that a fully scripted call will make them redundant? That worry doesn't seem to be what's driving Sonia's longer call times, but her concern that new policies will reduce customer satisfaction is an important signal. Phase-one resistance prevents any team from achieving phase-two benefits.

Those tensions are probably why the team meeting fell apart so dramatically. The current situation is confusing, frustrating, and unclear, so discipline or punishment is not the answer. I would talk to each person individually to understand their views on the AI pilot, then gather the team and be completely transparent about the new approach Jennifer and her manager have agreed on. Her message should be: "Here is the mission, here is why the company is investing in AI, and here is what we need to achieve together."

The rise of AI requires managers to think about not just talent and tools but also operating models: who does what, how humans and AI interact, what success looks like, and who is responsible when the system gets something wrong. If Jennifer can navigate the inevitable bumps in the road and succeed in taking her team from AI adoption to AI value creation, she will become the change agent that BPC needs. ■

HBR Reprint R2606L
Reprint Case only R2606X
Reprint Commentary only R2606Z

Executive Summaries
November–December 2026

AI & MACHINE LEARNING | OPERATIONS & SUPPLY CHAIN MANAGEMENT | LEADERSHIP | CUSTOMER STRATEGY

Why Even Your Best Employees Override AI Systems That Work

Das Narayandas | page 42

Companies are deploying AI systems faster than they are redesigning the governance surrounding them. As a result, many AI systems remain permanent advisers rather than actual decision-makers—not because the technology falls short but because the organizations fail to transfer authority along with accountability. When employees remain personally responsible for outcomes, overriding the system is a rational act of self-protection, even when the AI performs better than humans.

This dynamic has appeared before with ERP systems, credit scoring, and revenue management tools, says the author. And it is becoming even more consequential now as agentic AI begins to act autonomously. Leaders who want AI adoption to succeed must redesign roles, decision rights, incentives, and governance so that accountability moves with authority rather than remaining trapped in legacy organizational structures.

HBR Reprint R2606A

Design Your Supply Chain as a Portfolio of Options

Luiz Mesquita et al. | page 50

Global supply chains are being reshaped by constantly changing rules: tariffs, export controls, sanctions, and shifting geopolitical alliances. In this environment traditional resilience strategies—keeping buffer inventories, diversifying suppliers, and engaging in collaborative planning with others in the supply chain—aren't enough. Instead, say the authors, companies should design their supply chains as portfolios of real options, which will preserve their ability to defer investments, abandon unfavorable positions, expand in promising regions, or shift operations as policies evolve. As companies such as Apple and Zara have demonstrated, building flexibility this way can become a strategic advantage.

Leaders must keep in mind that optionality depends on capabilities that cannot be built overnight, including strong supplier relationships, regulatory expertise, and local production know-how. Leaders must therefore rethink capital allocation and invest early in the infrastructure needed to adapt quickly when the rules change again.

HBR Reprint R2606B

How CEOs Set the Tempo

Nitin Nohria | page 58

CEOs are judged not only by the strategies they choose but by the pace at which they execute them. Move too quickly and organizations burn out, make avoidable mistakes, or outrun their capabilities. Move too slowly and companies miss opportunities, lose relevance, or fail to respond to disruption. Effective leaders treat tempo as a strategic choice rather than a personality trait. Drawing on observations of hundreds of CEOs, this article identifies five distinct leadership tempos—from experimental speed to patient compounding—and explains when each is most effective. It argues that the right pace depends on the CEO's own nature, the organization's maturity and readiness, the business challenge, and external conditions. Leadership's overlooked skill is knowing when to accelerate, when to slow down, and why.

HBR Reprint R2606C

Do You Know What Your Customers Value?

Vikas Mittal et al. | page 68

Many executives mistakenly believe that because they've spent heavily on gathering customer feedback, monitoring competitors, and analyzing market trends, they understand what customers want. But the authors' seven-year study of 72 companies and additional research involving thousands of customers and firms reveal that companies routinely misidentify the factors that drive customer value. That happens because leaders base their customer strategies on intuition, recent events, or competitors' moves rather than scientific evidence. Not only that, but they tend to change their strategic priorities frequently, which leads to organizational clutter, confused customers, and wasted resources. The solution is to use data and statistical models to identify which customer value drivers have the greatest impact on sales and profitability, align employees around a small number of priorities, and eliminate low-value activities. Companies that adopt this disciplined approach can simplify strategy, strengthen customer loyalty, improve productivity, and generate superior financial returns.

HBR Reprint R2606D

AI & MACHINE LEARNING

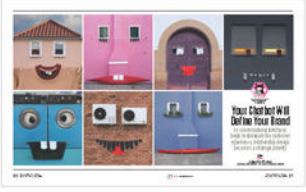

Your Chatbot Will Define Your Brand

Julian De Freitas | page 80

As AI chatbots become customers' primary interface with companies, relationship management is shifting from a marketing challenge to a product design challenge. Customers increasingly experience the chatbot as the brand itself, making emotional connections, sharing personal information, and expecting consistent interactions. That creates new opportunities—and new risks. Rather than focusing solely on chatbot capabilities, leaders must design and govern the relationship life cycle by defining the chatbot's role, adapting interactions to customer needs, calibrating the appropriate level of intimacy, repairing trust when expectations are violated, and deciding which AI relationship roles belong under the same brand. The central lesson is that AI is transforming brand management into relationship architecture, rewarding companies that design these relationships deliberately and honor the implied terms each one sets.

HBR Reprint R2606E

ENTREPRENEURSHIP

The Risks of Scaling Too Fast

Benjamin Hallen, Ron Tidhar, and Kathleen Eisenhardt
page 90

For decades, entrepreneurs have been told to prioritize growth above all else. But the authors' research shows that this "growth first" logic often undermines long-term success. In their rush to reach scale, founders frequently lock in flawed business models with weak unit economics—a problem that becomes harder to fix as firms grow. The most successful ventures take the opposite approach. They build a strong foundation before scaling up by rigorously testing both revenue and cost drivers, simplifying their business models, and deliberately delaying growth. This "foundation first" approach enables founders to achieve product-market fit with a viable economic model in place. Growth then becomes a process of replication rather than discovery—allowing companies to expand faster, more efficiently, and with far greater odds of long-term success.

HBR Reprint R2606F

BOARDS

Building a Board That Fits

Timothy J. Rowley and Stanislav Shekshnia | page 102

Boards tend to suffer from two design problems: First, they're set up to reflect the instincts of founders, CEOs, or company owners rather than counterbalance them. Second, even effective boards often fail to evolve as their organizations change. The solution starts with taking into account your firm's ownership context in your board design, since each context creates distinct governance risks requiring different forms of oversight. Next, follow a framework built around four elements—purpose, people, process, and the chair—to ensure that your board can manage those risks and provide the capabilities your firm needs. Rather than recruiting prestigious directors, assemble boards that challenge assumptions, improve strategic decision-making, and adapt as the company evolves. Effective governance comes not from individual talent but from boards designed to fit the organization's current context.

HBR Reprint R2606G

STRATEGY

The New Platform Playbook

Chiara Farronato | page 114

For more than two decades, platform strategy has followed a simple formula: Build scale quickly, harness network effects, and pursue winner-take-all dominance. But that playbook has become outdated. Network effects still matter, but tighter regulation, greater capital discipline, and shifting customer preferences have created markets in which multiple platforms can thrive. Winning now requires smarter strategic choices. Leaders must identify which side of the market generates the most value for other users, improve the quality of matches rather than simply adding users, develop revenue streams beyond transaction fees, and treat trust and consumer welfare as strategic assets. The next generation of winners won't wait for markets to tip in their favor. They will build differentiated, profitable, and trusted platforms designed to create lasting value amid persistent competition.

HBR Reprint R2606H

Executive Summaries
November–December 2026

MANAGING YOURSELF

How Leaders Counter Pessimism in Trying Times

Michelle Gielan | page 126

When teams face challenging times and are overwhelmed by negative news, failures, roadblocks, or uncertainty, the learned helplessness and pessimism they feel can erode engagement, energy, and performance. Leaders cannot remedy the situation with forced positivity; instead, they need to encourage *resilient optimism*—the ability to see challenges clearly while still believing that action matters—in three ways. First, they should model constructive behavior by directing attention toward purpose and progress, prioritizing visibly rather than spreading urgency, and treating restoration as part of sustainable performance. Second, they should help employees reframe stress as useful information, using practices such as signal checks and "yes, and" language to acknowledge difficulty while encouraging people to problem-solve. Finally, they should actively create and amplify positive moments—through praise, small acts of support, and renewed connection—so that their employees begin to see evidence that their efforts still make a difference.

HBR Reprint R2606J

HOW WE DID IT

Robinhood's CEO on Refounding the Business Post-IPO

Vlad Tenev | page 130

When Robinhood's pandemic-fueled growth abruptly reversed after its 2021 IPO, CEO Vlad Tenev realized the company needed to reinvent itself. He overhauled the organization, made difficult decisions on layoffs, ended companywide wellness weeks, brought employees back to the office, and clarified accountability through a new management structure. At the same time, Robinhood shifted from a trading app to a broader financial platform built around three strategic pillars: serving active traders, expanding customers' financial relationships through products like Robinhood Gold, and building a global financial platform. The experience convinced Tenev that adaptability, speed, and organizational clarity matter most when markets change.

HBR Reprint R2606K

POSTMASTER
Send domestic address changes, orders, and inquiries to: *Harvard Business Review,* Subscription Service, P.O. Box 37457, Boone, IA 50037. GST Registration No. 1247384345. Periodical postage paid at Boston, MA, and additional mailing offices.

Printed in the U.S.A. *Harvard Business Review* (ISSN 0017–8012; USPS 0236–520), published every other month for professional managers, is an education program of Harvard Business School, Harvard University; Srikant Datar, dean.

Published by Harvard Business School Publishing Corporation, 60 Harvard Way, Boston, MA 02163.

HARVARD BUSINESS REVIEW

STATEMENT OF OWNERSHIP, MANAGEMENT, AND CIRCULATION

HARVARD BUSINESS REVIEW
Publication No. 0017-8012

ISSUE FREQUENCY Bimonthly

NUMBER OF ISSUES PUBLISHED ANNUALLY 6

ANNUAL SUBSCRIPTION PRICE
$99, $109 (U.S.)
$120, $139 (Canada and Mexico)
$169 (international)

GENERAL BUSINESS OFFICE
60 Harvard Way, Boston, MA 02163; Middlesex County

OWNERS Harvard Business Publishing; 60 Harvard Way, Boston, MA 02163
President and Fellows of Harvard College | Cambridge, MA 02138

VP, DEPUTY PUBLISHER
Erika Heilman
60 Harvard Way, Boston, MA 02163

EDITOR Amy Bernstein
60 Harvard Way, Boston, MA 02163

MANAGING EDITOR Dan McGinn
60 Harvard Way, Boston, MA 02163

BONDHOLDERS, MORTGAGEES, OR OTHER SECURITY HOLDERS None

The purpose, function, and nonprofit status of this organization and the exempt status for federal income tax purposes has not changed during the preceding 12 months.

ISSUE DATE FOR THE DATA BELOW
September 1, 2026 (filed and subject to audit)

AVERAGE NUMBER OF COPIES FOR EACH ISSUE DURING THE PRECEDING 12 MONTHS	
Total number of copies (net press run)	284,844
Paid circulation (by mail and outside the mail):	
Mailed outside-county paid subscriptions stated on PS Form 3541	126,869
Mailed in-county paid subscriptions stated on PS Form 3541	0
Paid distribution outside the mails including sales through dealers and carriers, street vendors, counter sales, and other paid distribution outside USPS	113,301
Paid distribution by other classes of mail through the USPS	0
Total paid distribution	**240,170**
Free or nominal rate distribution (by mail and outside the mail):	
Outside-county copies included on PS Form 3541	0
In-county copies included on PS Form 3541	0
Copies mailed at other classes through the USPS	832
Distribution outside the mail	0
Total free or nominal rate distribution	**832**
Total distribution	241,002
Copies not distributed	43,843
Total	**284,844**
Percent paid and/or requested circulation	99.65%

NUMBER OF COPIES OF SINGLE ISSUE PUBLISHED NEAREST TO FILING DATE	
Total number of copies (net press run)	264,705
Paid circulation (by mail and outside the mail):	
Mailed outside-county paid subscriptions stated on PS Form 3541	123,137
Mailed in-county paid subscriptions stated on PS Form 3541	0
Paid distribution outside the mails including sales through dealers and carriers, street vendors, counter sales, and other paid distribution outside USPS	111,544
Paid distribution by other classes of mail through the USPS	0
Total paid circulation	**234,681**
Free or nominal rate distribution (by mail and outside the mail):	
Outside-county copies included on PS Form 3541	0
In-county copies included on PS Form 3541	0
Copies mailed at other classes through the USPS	727
Distribution outside the mail	0
Total free or nominal rate distribution	**727**
Total distribution	235,408
Copies not distributed	29,297
Total	**264,705**
Percent paid and/or requested circulation	99.69%

I certify that all information furnished on this form is true and complete.

Gregory Daly, ASSOCIATE DIRECTOR, OPERATIONS

Happier teams start with happier managers.

ARTHUR C. BROOKS,
HARVARD PROFESSOR AND
BEST-SELLING AUTHOR

The Leader's Happiness Reset is a free six-week newsletter on finding happiness at work—for yourself and your team—from HBR and Arthur C. Brooks.

The Leader's Happiness Reset

HARVARD BUSINESS REVIEW

SIGN UP TODAY
hbr.org/newsletters

"You've got to accept that failure and rejection are part of the job. Deal with it. Don't take it personally. If you can't, don't become an actor."

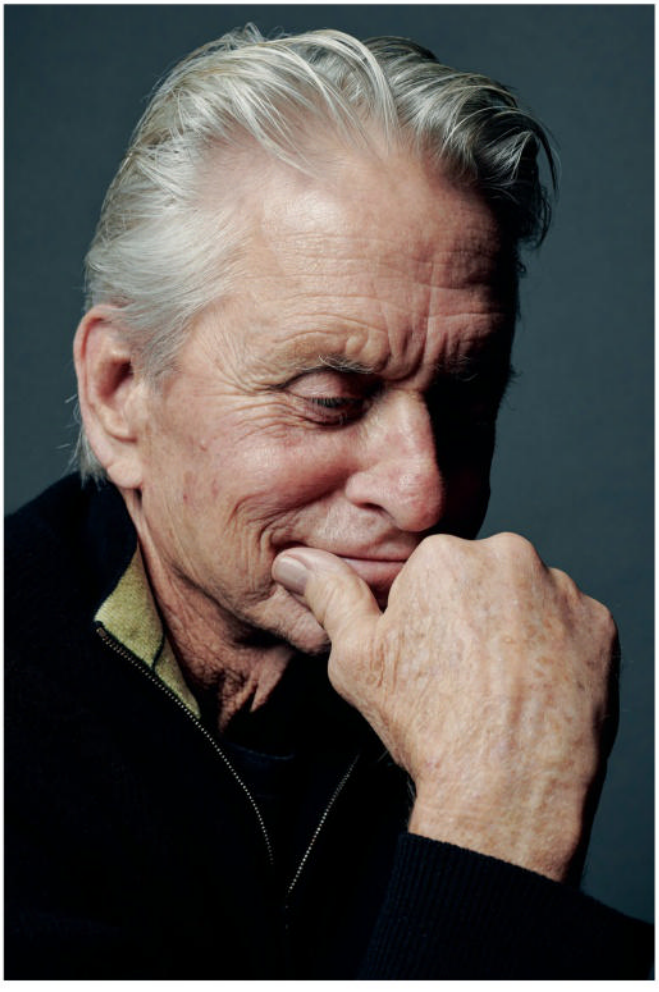

Vincent Tullo/*The New York Times*/Redux

Michael Douglas

The son of two actors—one a mid-20th-century megastar—Douglas never intended to join the family business. But he overcame his hesitancy (and early stage fright) to emerge as one of the film industry's most creative producers, winning his first Academy Award at age 31 for his work on One Flew over the Cuckoo's Nest. *Later he became one of Hollywood's most bankable leading men, with hits like* Romancing the Stone, Fatal Attraction, Wall Street *(for which he earned a best actor Oscar),* The American President, *and* Traffic. *Now 82, he has a new memoir called* One Helluva Ride. **Interview by Alison Beard**

HBR: Early on, did you worry that people might think you were trading on your dad's name?

DOUGLAS: I don't know of any parent who, if their son or daughter wants to get into their business, doesn't try to help. And actually, the track record for second-generation actors isn't very strong. When I auditioned, the last thing I wanted to hear was, "You look just like your father." I wanted to get hired for my own unique qualities. It wasn't until the year of *Fatal Attraction* and *Wall Street* that I felt I had proven myself and stepped out of his shadow.

How did you develop such good creative judgment?

It starts with gut instinct. Does the script elicit emotions or not? That's the initial read. Then I analyze the structure for a first, second, and third act. Another thing is that I'm not looking for my part but at the whole movie. I don't pick a vehicle for me as an actor. It's only about the movie.

As a producer, how do you assemble cohesive project teams?

You start with "No dickheads allowed." If you've been around the business, you know who's trouble, and you don't want them in any department. There are too many other talented people to hire. This is a quality-of-life issue. I want it to be fun, and I don't work well under tension myself, so park your ego at the door. Also, when I'm number one on the call sheet, I set the tone. If a star lollygags, everyone else does, too, and you lose time. My approach is to always ask for a 10-minute warning before I'm expected on set, so I'm always the first ready. Soon the other actors are asking for 10-minute warnings so that we can all get going together. It's an accordion effect.

How do you synthesize feedback from varied stakeholders?

As an actor I normally don't hear enough feedback, so I have to be careful to not be too critical of myself and to ask for others' opinions. If I'm the producer, I talk to the director a lot and hope that we and the studio are on the same wavelength. I welcome feedback from everybody on a rough cut because if you hear the same note more than two or three times, it's generally something you should consider. You also have to evaluate who you're talking to. Is the studio executive an idiot or knowledgeable? Either way, you have to kiss their ass, but in some cases you can work around notes that aren't good.

You've worked with so many top directors. What did they teach you about managing talent?

Every director has an image of how the scene should be and how the characters should act. The good ones hold that back and tell the actors, "Let's see what you've got," and if they like what they see, use it. The not-so-good directors hold on to their vision and want you to comport to it. In terms of how you get people's best, though, there are different ways. Oliver Stone is tough. He tests you and tells you what he thinks. But look at the incredible record he has of getting strong performances. Others, like Adrian Lyne, are more like little kids, getting so excited and encouraging you. Both work. ■

HBR Reprint R2606M

For more from Michael Douglas, go to **HBR.org**.

Look further.
Act sooner.
Create impact that lasts.
Turn foresight and ambition into a radically different future.
KEARNEY
100 Years of Impact
kearney.com

Publicité Et Nous

HERMÈS
PARIS

ECHO OF THE EARTH

TERRE D'HERMÈS

PURE PERFUME